

---

# What's New for Developer's in PostgreSQL 17

Deepak Mahto  
DataCloudGaze Consulting

# About Me

---

I am Deepak Mahto, and I like to call myself your Database Guy.

- Founder at DataCloudGaze Consulting.
- Conduct training on PostgreSQL and Migration at DatabaseGyaan.
- Loves to blog on Databases Learnings.
- I live in Mumbai and have a 3-year-old child.
- Loves to explore street food.



# Agenda

---

- Key upcoming features in PostgreSQL 17 for Developers.
  - Optimizer
  - Utility Command
  - Partitions
  - Merge
  - PL/pgSQL
  - Partitions
  - JSON
  - Trigger
  - Privilege
  - And more..



---

# PostgreSQL 17 - Optimizer Sections

# Optimizer - Null Constraints and Performance.

Improve optimization of IS NOT NULL and IS NULL query restrictions.

```
pg17=# alter table t1 alter column col1 set not null;  
ALTER TABLE
```

```
pg17=# \d t1
```

Column	Type	Collation	Nullable	Default
col1	integer		not null	
col2	integer			
col3	numeric			
col4	text			

# Optimizer - Null Constraints and Performance.

Improve optimization of IS NOT NULL and IS NULL query restrictions.

PG 17 – IS NULL Filter

```
pg17=# explain (analyze, buffers,  
costs off) select count(1) from t1  
where col1 is null;
```

Aggregate (actual  
time=0.003..0.004 rows=1 loops=1)  
-> Result (actual  
time=0.001..0.001 rows=0 loops=1)  
*One-Time Filter: false*  
Planning Time: 0.040 ms  
Execution Time: 0.021 ms  
(5 rows)

PG 17 – IS Not NULL Filter

```
pg17=# explain (analyze, buffers,  
costs off) select count(1) from t1  
where col1 is not null;
```

Aggregate (actual  
time=0.032..0.033 rows=1 loops=1)  
Buffers: shared hit=1  
-> Seq Scan on t1 (actual  
time=0.011..0.020 rows=100  
loops=1) *no Filter Clause is specified*  
Buffers: shared hit=1  
Planning Time: 0.046 ms  
Execution Time: 0.051 ms  
(6 rows)

# Optimizer - enable\_group\_by\_reordering

Allow GROUP BY columns to be internally ordered for better plan.

```
pg17=# set enable_group_by_reordering to off;
SET
pg17=# explain (settings, analyze, costs off) select count(*),c2,c1 from t1 group by c2,c1;
                                         QUERY PLAN
-----
HashAggregate (actual time=325.221..325.876 rows=1000 loops=1)
  Group Key: c2, c1
  Batches: 1  Memory Usage: 3153kB
    -> Seq Scan on t1 (actual time=0.196..167.200 rows=1000000 loops=1)
Settings: enable_group_by_reordering = 'off', max_parallel_workers_per_gather = '0'
Planning Time: 1.499 ms
Execution Time: 328.934 ms
(7 rows)
```

New Flag in PostgreSQL 17, by default it is enable.

# Optimizer - enable\_group\_by\_reordering

Allow GROUP BY columns to be internally ordered for better plan.

```
pg17=# explain (settings, analyze, costs off) select count(*),c2,c1 from t1 group by c2,c1;  
                                         QUERY PLAN  
-----  
GroupAggregate (actual time=0.611..206.106 rows=1000 loops=1)  
  Group Key: c1, c2  
    -> Index Only Scan using t1_c1_c2_idx on t1 (actual time=0.248..121.351 rows=1000000 loops=1)  
        Heap Fetches: 0  
Settings: max_parallel_workers_per_gather = '0'  
Planning Time: 1.357 ms  
Execution Time: 206.461 ms  
(7 rows)
```

The diagram shows a green box highlighting the 'Group Key: c1, c2' part of the query plan. Another green box highlights the 't1\_c1\_c2\_idx' index name. A green arrow points from the 't1\_c1\_c2\_idx' box to the text 'Index on c1,c2' located below the execution statistics.

# Optimizer - Correlated IN Clause Transformation.

Allow correlated IN subqueries to be transformed into optimal joins

```
explain (analyze, costs off)
select * from test_in_correlated1 t1 where t1.col1 in (select t2.col1*99999
from test_in_correlated2 t2 where t1.col2= t2.col2);
```

PG16

QUERY PLAN

```
Seq Scan on test_in_correlated1 t1 (actual time=4546.820..4546.829 rows=0 loops=1)
  Filter: (SubPlan 1)
  Rows Removed by Filter: 10000
    SubPlan 1
      -> Seq Scan on test_in_correlated2 t2 (actual time=0.023..0.451 rows=10 loops=10000)
          Filter: (t1.col2 = col2)
          Rows Removed by Filter: 9990
Planning Time: 0.216 ms
Execution Time: 4546.909 ms
(9 rows)
```

SubPlan is looped 10k times to resolve correlated filter

# Optimizer - Correlated IN Clause Transformation.

Allow correlated IN subqueries to be transformed into optimal joins

```
explain (analyze, costs off)
select * from test_in_correlated1 t1 where t1.col1 in (select t2.col1*99999
from test_in_correlated2 t2 where t1.col2= t2.col2);
```

QUERY PLAN

```
Hash Join (actual time=31.352..31.355 rows=0 loops=1)
  Hash Cond: ((t2.col2 = t1.col2) AND ((t2.col1 * 99999) = t1.col1))
    -> HashAggregate (actual time=12.418..13.889 rows=10000 loops=1)
      Group Key: t2.col2, (t2.col1 * 99999)
      Batches: 1 Memory Usage: 1169kB
      -> Seq Scan on test_in_correlated2 t2 (actual time=0.712..8.530 rows=10000 loops=1)
    -> Hash (actual time=15.098..15.098 rows=10000 loops=1)
      Buckets: 16384 Batches: 1 Memory Usage: 645kB
      -> Seq Scan on test_in_correlated1 t1 (actual time=1.619..13.035 rows=10000 loops=1)
```

Planning Time: 7.713 ms  
Execution Time: 32.389 ms  
(11 rows)

Correlated IN Clause is Transform to Hash Join

# Optimizer - Improve CTE Plans

Use CTE columns statistics and sort order for better plan and estimates.

```
pg16=# explain with x as materialized (select c3 from t1 b)
pg16=# select count(1) from t1 a where c3 in (select * from x);
          QUERY PLAN
```

table t1 and same column c3 is used within CTE and outside PG16

Aggregate (cost=62754.00..62754.01 rows=1 width=8)

CTE x

-> Seq Scan on t1 b (cost=0.00..15406.00 rows=1000000 width=4)

estimated 1000k rows

-> Hash Join (cost=22504.50..46098.00 rows=500000 width=0)

Hash Cond: (a.c3 = x.c3)

-> Seq Scan on t1 a (cost=0.00..15406.00 rows=1000000 width=4)

500k is estimated.

-> Hash (cost=22502.00..22502.00 rows=200 width=4)

-> HashAggregate (cost=22500.00..22502.00 rows=200 width=4)

Group Key: x.c3

-> CTE Scan on x (cost=0.00..20000.00 rows=1000000 width=4)

(10 rows)

c3 column stats  
is not preserve  
outside of CTE.

# Optimizer - Improve CTE Plans

Use CTE columns statistics and sort order for better plan and estimates.

```
pg17=# explain with x as materialized (select c3 from t1 b)
select count(1) from t1 a where c3 in (select * from x);
          QUERY PLAN
```

```
Aggregate  (cost=95190.00..95190.01 rows=1 width=8)
  CTE x
    -> Seq Scan on t1 b  (cost=0.00..15406.00 rows=1000000 width=4)
    -> Hash Semi Join (cost=36407.00..77284.00 rows=1000000 width=0)
      Hash Cond: (a.c3 = x.c3)
        -> Seq Scan on t1 a  (cost=0.00..15406.00 rows=1000000 width=4)
        -> Hash  (cost=20000.00..20000.00 rows=1000000 width=4)
          -> CTE Scan on x  (cost=0.00..20000.00 rows=1000000 width=4)
```

(8 rows)

*c3 stats outside CTE helps to choose better plan*

*c3 column stats is used outside CTE*

---

# PostgreSQL 17 - PL/pgSQL

# PL/pgSQL - New Array Declaration

Allow plpgsql %TYPE and %ROWTYPE specifications to represent arrays of non array type.

```
pg17=# do language plpgsql
$$
<<block1>>
declare
var1 testplpgsql17.col1%type[];
var2 testplpgsql17%rowtype[];
begin
SELECT array_agg(col1),array_agg(tbl.*)
into var1, var2
FROM (select * from testplpgsql17 limit 2) as tbl ;
raise notice 'PL\pgsql PG 17 - type[] - %',var1;
raise notice 'PL\pgsql PG 17 - rowtype[] - %',var2;
end block1;
$$;
NOTICE: PL\pgsql PG 17 - type[] - {1,2}
NOTICE: PL\pgsql PG 17 - rowtype[] - {"(1,1,9.9)","(2,2,19.8)"}
DO
```

Declare Array using %type[] and %rowtype[] for non array type.

---

# PostgreSQL 17 - Utility Command

# Utility Command - Explain (Memory)

## Allow EXPLAIN to report optimizer memory usage

```
--> Seq Scan on pgbench_accounts_3499 (actual time=0.010..0.010 rows=0 loops=1)
      Filter: (bid = 1)
--> Seq Scan on pgbench_accounts_3500 (actual time=0.004..0.004 rows=0 loops=1)
      Filter: (bid = 1)
Planning:
  Memory: used=34007kB allocated=41819kB
Planning Time: 623.421 ms
Execution Time: 313.460 ms
(8730 rows)
```

With large set of partitions to scan planning operation memory need is increases.

3.5k partition scan

```
explain (analyze, costs off, memory) select * from pgbench_accounts where bid =1;
```

Completing running Execution plan for many partitions scans , check Planning Memory usage.

```
explain (analyze, costs off, memory) select * from pgbench_accounts where aid =1;
QUERY PLAN
```

```
Seq Scan on pgbench_accounts_1 pgbench_accounts (actual time=0.029..0.040 rows=1 loops=1)
  Filter: (aid = 1)
  Rows Removed by Filter: 57
Planning:
  Memory: used=56kB allocated=68kB
Planning Time: 1.081 ms
Execution Time: 0.075 ms
(7 rows)
```

Planning memory reduces with lower partitions on scans.

Completing running Execution plan for few partitions scans, check Planning Memory usage.

# Utility Command - Explain (serialize)

Allow Explain command to show overhead of convert and de-toasting and sending it over the wire.

Table "public.testtoast"					
Column	Type	Collation	Nullable	Default	
col1	bigint		not null	nextval('testtoast_col1_seq)::regclass	
col2	text				
table_schema	table_name	table_toast	table	index	total
public	testtoast	26 MB	96 kB	0 bytes	26 MB

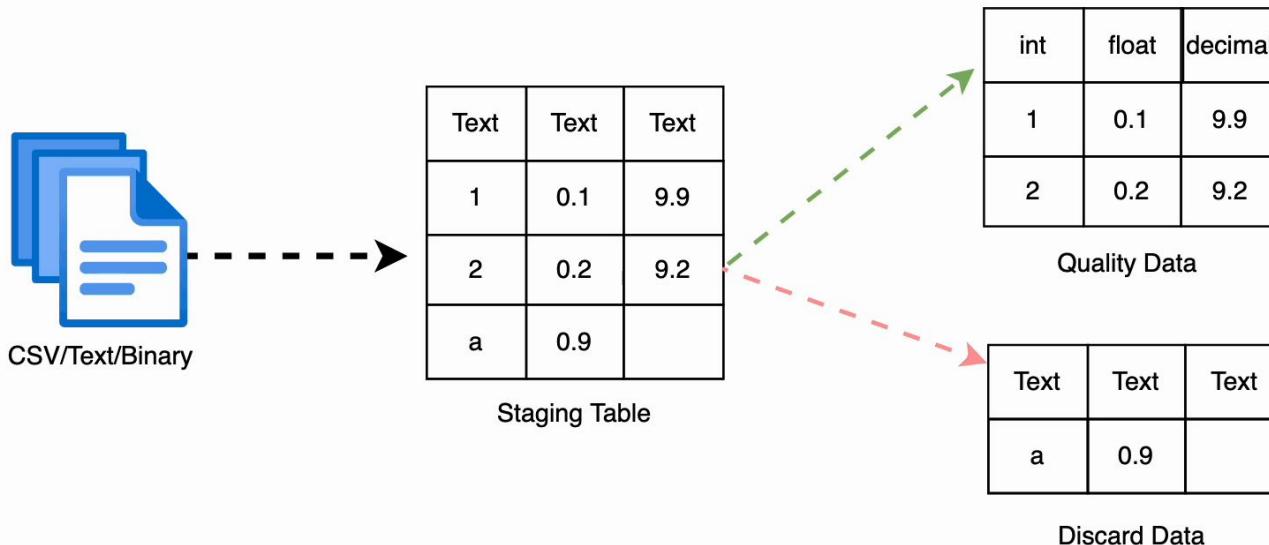
```
Running - explain (analyze, buffers) select * from testtoast
          QUERY PLAN
-----  
Seq Scan on testtoast  (cost=0.00..14.00 rows=600 width=72) (actual time=1.297..2.838 rows=600 loops=1)
  Buffers: shared hit=8
Planning:
  Buffers: shared hit=8
Planning Time: 0.229 ms
Execution Time: 2.928 ms
(6 rows)
```

With our serialize, we can't measure overhead of serializing toasted columns over the wire.

```
Running - explain (analyze, buffers,serialize) select * from testtoast
          QUERY PLAN
-----  
Seq Scan on testtoast  (cost=0.00..14.00 rows=600 width=72) (actual time=0.025..0.436 rows=600 loops=1)
  Buffers: shared hit=8
Planning Time: 0.101 ms
Serialization: time=1475.352 ms output=2225399kB format=text
  Buffers: shared hit=1829 read=3084
Execution Time: 1476.359 ms
(6 rows)
```

# Utility Command - Copy

New option to ignore Bad record and Track it with Verbosity option.



# Utility Command - Copy

---

- Handling data type incompatibilities with ON\_ERROR options
- Improved logging with LOG\_Verbosity
- Observability via pg\_stat\_progress\_copy with tuples\_skipped

# Utility Command - Copy command is more user friendly

Load fails due to data type mismatch or bad data.

```
pg17=# \copy copy17 from 'test.csv' (format csv);  
ERROR: date/time field value out of range: "2024-01-01 12:00:61"  
CONTEXT: COPY copy17, line 4, column col7: "2024-01-01 12:00:61"
```

```
pg17=# \copy copy17 from 'test.csv' (ON_ERROR ignore ,  
LOG_VERTOSITY default , format csv);  
NOTICE: 5 rows were skipped due to data type incompatibility  
COPY 3
```

```
pg17=# \copy copy17 from 'test.csv' (ON_ERROR ignore ,  
LOG_VERTOSITY verbose , format csv);  
NOTICE: skipping row due to data type incompatibility at line 4 for column col7: "2024-01-01 12:00:61"  
NOTICE: skipping row due to data type incompatibility at line 5 for column col1: "a"  
NOTICE: skipping row due to data type incompatibility at line 6 for column col1: "99999999999999999999999999999999"  
NOTICE: skipping row due to data type incompatibility at line 7 for column col5: "adhoc"  
NOTICE: skipping row due to data type incompatibility at line 8 for column col6: "20241301"  
NOTICE: 5 rows were skipped due to data type incompatibility  
COPY 3
```

New ON\_ERROR ignore option and verbosity enable skips bad data.

Enable get additional information on skips records

# Utility Command - tuples\_skipped

```
1 | select command, tuples_processed , tuples_excluded ,  
2 | tuples_skipped from pg_stat_progress_copy;
```

```
-[ RECORD 1 ]-----  
command           | COPY FROM  
tuples_processed | 3  
tuples_excluded | 0  
tuples_skipped | 5
```

---

# PostgreSQL 17 - New JSON Functions.

# SQL/JSON Function - JSON\_TABLE

JSON\_TABLE() to convert JSON data to a table representation

```
pg17=# SELECT jt.* FROM
my_films,
JSON_TABLE ( js, '$.favorites[*]'      Allow to Query Json as Table
            COLUMNS (
              id FOR ORDINALITY,
              kind text PATH '$.kind',
              NESTED PATH '$.films[*]' COLUMNS (
                title text FORMAT JSON PATH '$.title' OMIT QUOTES,
                director text PATH '$.director' KEEP QUOTES))) AS jt;
id | kind    |       title        |   director
---+-----+-----+-----+
 1 | comedy  | Nuvvu Naaku Nachav | "K. Vijaya Bhaskar"
 1 | comedy  | Bangalore Days     | "Anjali Menon"
 2 | horror   | Pizza             | "Karthik Subbaraj"
 3 | thriller | Kalki            | "Prasanth Varma"
 4 | drama    | Bahubali          | "S. S. Rajamouli"
 4 | drama    | Salaar            | "Prashanth Neel"
(6 rows)
```

# SQL/JSON Constructor Functions

```
pg17=# select json_scalar(random(1,10)), json_scalar(now()::timestamp without time zone);
          json_scalar |      json_scalar
-----+-----+
 2           | "2024-08-30T00:27:22.058111"
-- 

pg17=# select JSON_SERIALIZE('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}' RETURNING text);
          json_serialize
-----+
{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}
(1 row)

pg17=# select JSON_SERIALIZE('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}' RETURNING bytea);
          json_serialize
-----+
\x7b226b696e64223a20226472616d61222c2266696c6d73223a205b7b227469746c65223a20224261687562616c69222c226469726563746f72223a2022532e2
0532e2052616a616d6f756c69227d5d7d

pg17=# select JSONC('{"kind": , "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}');
ERROR: invalid input syntax for type json
DETAIL: Expected JSON value, but found ", ".
CONTEXT: JSON data, line 1: {"kind": ,...
pg17=# select JSONC('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}');
          json
-----+
{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}
```

# SQL/JSON query functions

```
pg17=# SELECT JSON_EXISTS('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}', '$.films[0]');
json_exists
-----
t
(1 row)

pg17=# SELECT JSON_EXISTS('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}', '$.films[1]');
json_exists
-----
f
(1 row)

pg17=# SELECT JSON_QUERY('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}', '$.films[*].title');
json_query
-----
"Bahubali"
(1 row)

pg17=# SELECT JSON_VALUE('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}', '$.kind');
json_value
-----
drama
(1 row)

pg17=# SELECT JSON_VALUE('{"kind": "drama", "films": [{"title": "Bahubali", "director": "S. S. Rajamouli"}]}', '$.films[0].title');
json_value
-----
Bahubali
(1 row)
```

---

# PostgreSQL 17 - Partitions.

# Partition - Identity Column support

```
CREATE TABLE pg17_partition (
    id      int8 generated always AS IDENTITY,
    created_at timestamp without time zone NOT NULL,
    addinfo  TEXT
) PARTITION BY RANGE ( created_at );
```

Identity space is shared across Partitions.

```
CREATE TABLE pg17_partition_old PARTITION OF pg17_partition FOR VALUES FROM (MINVALUE) TO ('2022-01-01');
CREATE TABLE pg17_partition_2022 PARTITION OF pg17_partition FOR VALUES FROM ('2022-01-01') TO ('2023-01-01');
CREATE TABLE pg17_partition_2023 PARTITION OF pg17_partition FOR VALUES FROM ('2023-01-01') TO ('2024-01-01');
CREATE TABLE pg17_partition_2024 PARTITION OF pg17_partition FOR VALUES FROM ('2024-01-01') TO ('2025-01-01');
CREATE TABLE pg17_partition_2025 PARTITION OF pg17_partition FOR VALUES FROM ('2025-01-01') TO ('2026-01-01');
```

```
INSERT INTO pg17_partition (created_at, addinfo)
SELECT (now() - '5 years'::INTERVAL * random())::timestamp without time zone, 'Add info - ' || i FROM generate_series(1,10) i;
```

```
pg17=# select * from pg17_partition;
```

id	created_at	addinfo
3	2021-07-08 20:22:36.767674	Add info - 3
9	2021-08-25 10:33:17.049274	Add info - 9
10	2020-07-11 21:02:40.761274	Add info - 10
1	2022-09-29 19:10:33.225274	Add info - 1
7	2022-01-16 08:46:41.116474	Add info - 7
8	2022-04-13 02:33:59.327674	Add info - 8
2	2023-08-11 17:00:42.278074	Add info - 2
5	2023-10-14 23:00:36.748474	Add info - 5
4	2024-05-14 20:11:12.911674	Add info - 4
6	2024-04-14 01:56:31.545274	Add info - 6

(10 rows)

Based on  
Insertion  
order  
Identity  
Space is  
Shared  
Across  
Partitions.

# Partition - Split existing Partition Support

```
pg17=# ALTER TABLE pg17_partition
          SPLIT PARTITION pg17_partition
          (PARTITION pg17_partition_2024_01 FOR VALUES FROM ('2024-01-01') TO ('2024-04-01'),
           PARTITION pg17_partition_2024_02 FOR VALUES FROM ('2024-04-01') TO ('2024-07-01'),
           PARTITION pg17_partition_2024_03 FOR VALUES FROM ('2024-07-01') TO ('2024-10-01'),
           PARTITION pg17_partition_2024_04 FOR VALUES FROM ('2024-10-01') TO ('2025-01-01'));
```

```
ALTER TABLE
```

```
pg17=# \dt pg17_partition_2024*
          List of relations
```

Schema	Name	Type	Owner
public	pg17_partition_2024_01	table	postgres
public	pg17_partition_2024_02	table	postgres
public	pg17_partition_2024_03	table	postgres
public	pg17_partition_2024_04	table	postgres
(4 rows)			

Reverted in 17 Beta 3

Split existing partition to further small subsets.

Acquire access exclusive locks  
Hash partition is not supported.

# Partition - Merge existing Partition Support

```
pg17=# ALTER TABLE pg17_partition
      MERGE PARTITION
      pg17_partition_2022
      pg17_partition_2023
      INTO pg17_partition_old;
ALTER TABLE
```

```
pg17=# \dt pg17_partition*
```

List of relations

Schema	Name	Type	Owner
public	pg17_partition	partitioned table	postgres
public	pg17_partition_2023	table	postgres
public	pg17_partition_2024_01	table	postgres
public	pg17_partition_2024_02	table	postgres
public	pg17_partition_2024_03	table	postgres
public	pg17_partition_2024_04	table	postgres
public	pg17_partition_2025	table	postgres
public	pg17_partition_old	table	postgres

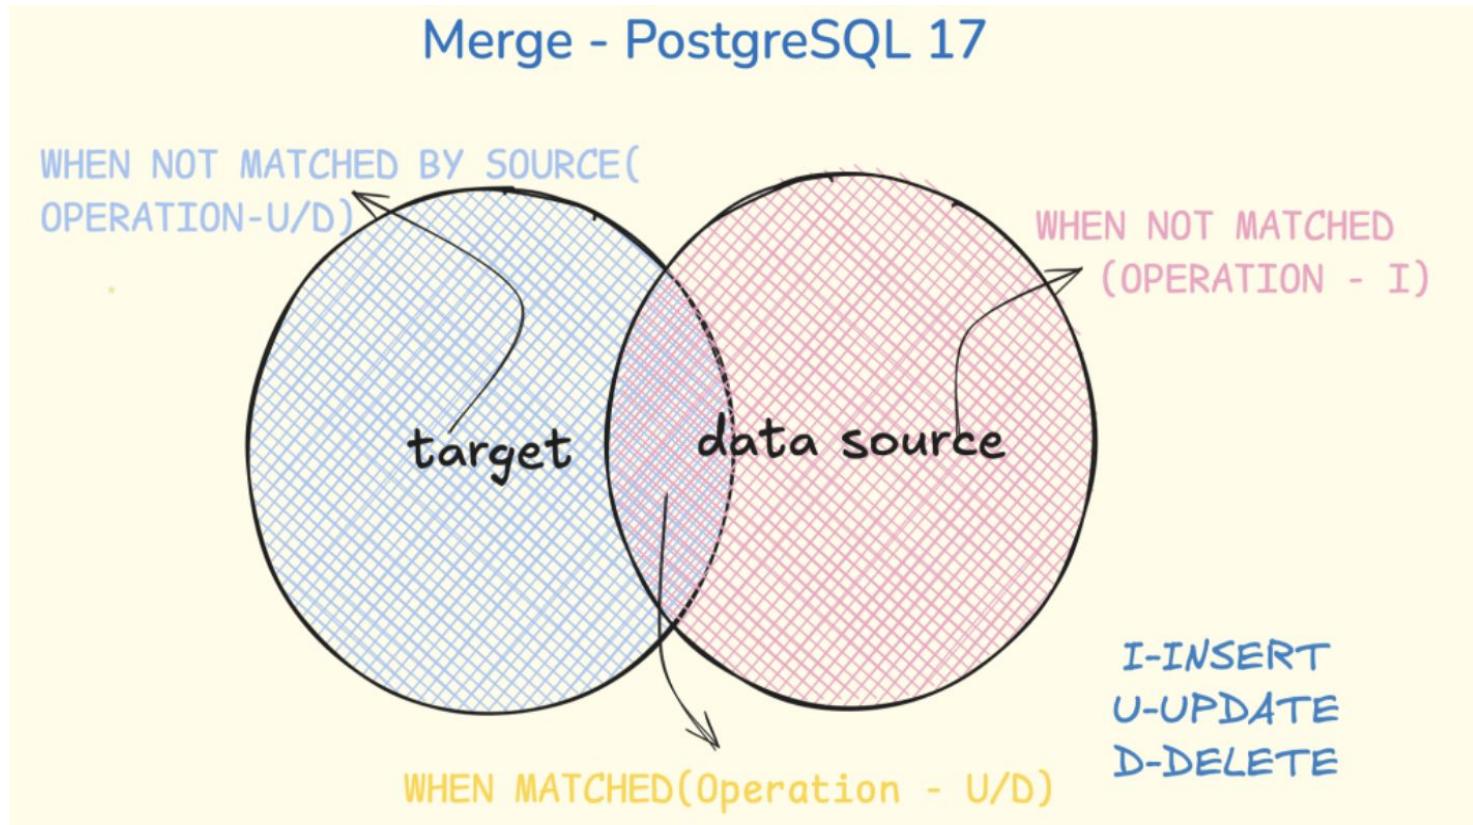
Reverted in 17 Beta 3

Merge existing Partition

---

# PostgreSQL 17 - Merge

# Merge Command in PostgreSQL 17



# Merge Command in PostgreSQL 17

## Fun Fact: Internal Join Strategies in PostgreSQL 17 MERGE

<code>when_clause</code>	<code>join_method</code>
<i>all (matched, not_matched, not_mached_by_source)</i>	<i>full join</i>
<i>only matched</i>	<i>inner join</i>
<i>only not matched</i>	<i>right join</i>
<i>only not_mached_by_source</i>	<i>left join</i>

# Merge Command in PostgreSQL 17

New merge\_action() Function: Highlights operations performed on the target table, now with added support for the RETURNING clause.

merge\_action() highlights operation perform on target

merge_action	product_id	product_name	quantity	last_updated	addinfo
UPDATE	1	Laptop	45	2024-08-09 11:02:27.739527	sale
UPDATE	2	Smartphone	105	2024-08-09 11:02:27.739527	sale+restock
UPDATE	3	Tablet	31	2024-08-09 11:02:27.739527	restock
DELETE	5	DVD Player	5	2024-08-07 10:00:00	
INSERT	6	Smartwatch	25	2024-08-09 11:02:27.739527	new
INSERT	7	GoPro	25	2024-08-09 11:02:27.739527	new

(6 rows)

---

# PostgreSQL 17 - psql command line

# Allow psql's \watch to stop after a min number of rows returned

```
pg17=# select generate_series(1,random(1,10));  
generate_series  
-----  
1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
(10 rows)
```

```
pg17=# \watch i=1 m=3
```

---

# PostgreSQL 17 - Server Configuration

# Terminating connection due to transaction\_timeout

Set timeout for Transactions that can help to overcome vacuum and table bloat challenges.

## List of configuration parameters

Parameter	Value
idle_in_transaction_session_timeout	0
idle_session_timeout	0
statement_timeout	0
transaction_timeout	0

Added in PostgreSQL 17

# Terminating connection due to transaction\_timeout

```
pg17=# set transaction_timeout=3000;  
SET  
pg17=# begin;select;select pg_sleep(2) as "sleep for 2 sec";select;select pg_sleep(2) as  
"sleep for further 2 seconds";  
BEGIN  
-- Trans started  
(1 row)  
  
sleep for 2 sec  added time of 2 seconds in transaction.  
-----  
  
(1 row)  
--  
(1 row)  
                  started sleep for 2 more sec  
  
FATAL: terminating connection due to transaction timeout  
server closed the connection unexpectedly  
This probably means the server terminated abnormally  
before or while processing the request.
```

Trans started

added time of 2 seconds in transaction.

started sleep for 2 more sec

FATAL: terminating connection due to transaction timeout  
server closed the connection unexpectedly  
This probably means the server terminated abnormally  
before or while processing the request.

# Configurable SLRU cache

---

- Postgres 17 adds configurable SLRU cache sizes to boost performance in large workloads.
- Splitting SLRU caches into multiple banks with locks enhances scalability.
- Configuration changes require a restart, key for workloads facing contention

List of configuration parameters

Parameter	Value
-----------	-------

Parameter	Value
commit_timestamp_buffers	256kB
multixact_member_buffers	256kB
multixact_offset_buffers	128kB
notify_buffers	128kB
Serializable_buffers	256kB
subtransaction_buffers	256kB
transaction_buffers	256kB

(7 rows)

---

# PostgreSQL 17 - Privileges

# Privileges - maintain / pg\_maintain at user level.

Allow granting the right to perform maintenance operations

```
pg17=> set role myuser1;
SET
pg17=> select count(1) from t1;
      count
-----
      1000000
(1 row)

pg17=> vacuum t1;
WARNING: permission denied to vacuum "t1", skipping it
VACUUM
pg17=> analyze t1;
WARNING: permission denied to analyze "t1", skipping it
ANALYZE
pg17=> set role postgres;
SET
pg17=# grant maintain on t1 to myuser1;
GRANT
pg17=# set role myuser1;
SET
pg17=> vacuum analyze t1;
VACUUM
```

myuser1 can select on t1.

myuser1 cannot analyze or vacuum.

new "maintain" privilege  
grant myuser1 to perform maintenance related activity.

---

# PostgreSQL 17 - Builtin Collation

# New Builtin Collation

---

- Built-in collation provider with UTF-8 encoding
- Provides sorting semantics similar to 'C' collation
- Guaranteed immutable for consistent sort results across systems

```
pg17=# \d0 pg_c_utf8|ucs_basic
|
Schema | Name | Provider | (
-----+-----+-----+
pg_catalog | pg_c_utf8 | builtin | 
pg_catalog | ucs_basic | builtin | 
(2 rows)
```

---

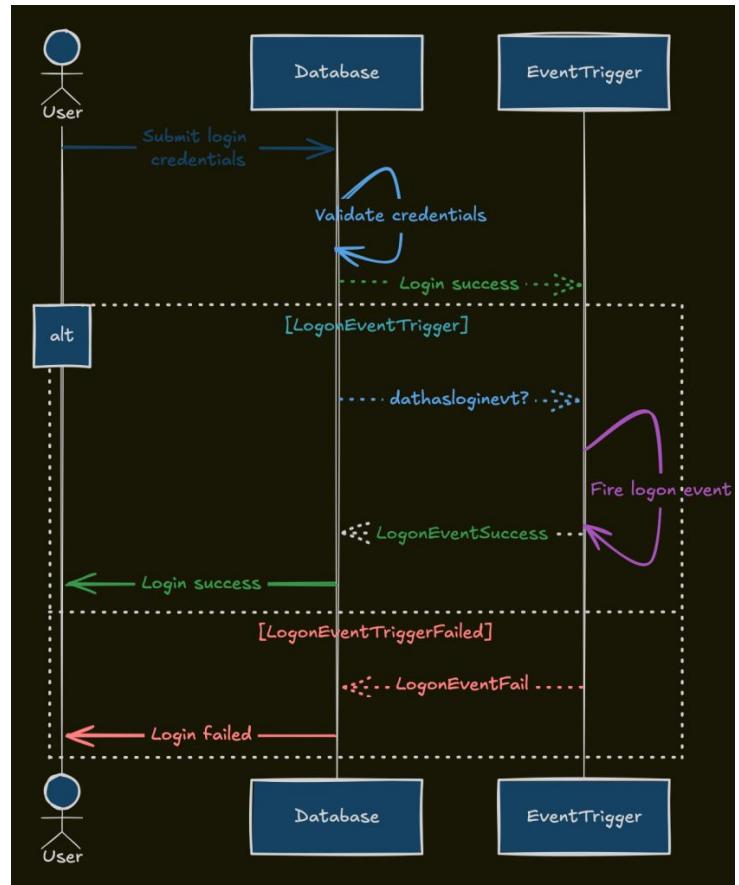
# PostgreSQL 17 - Event Trigger on Logon

# Event Triggers - Login

Implement login related usecases using login Event Trigger

```
CREATE OR REPLACE FUNCTION mandate_app_info()
RETURNS event_trigger AS $$  
BEGIN  
    IF nullif(current_setting('application_name'), '') is null then
        RAISE EXCEPTION 'Login not allowed for missing application_name'  
USING HINT = 'Please set application_name context';
    END IF;
END;  
$$ LANGUAGE plpgsql;
```

```
CREATE EVENT TRIGGER on_login_trigger ON login
EXECUTE PROCEDURE mandate_app_info();
ALTER EVENT TRIGGER on_login_trigger ENABLE ALWAYS;
```



# Event Triggers - Login in action

```
asmitamahto@dcg pghyd2024 % docker exec -it pg17b3 psql -h localhost -U postgres -d 'dbname=pg17 application_name='
psql: error: connection to server at "localhost" (::1), port 5432 failed: FATAL:  Login not allowed for missing application_name
```

HINT: Please set application\_name context

CONTEXT: PL/pgSQL function mandate\_app\_info() line 4 at RAISE

```
asmitamahto@dcg pghyd2024 % docker exec -it pg17b3 psql -h localhost -U postgres -d 'dbname=pg17 application_name=pgsql'
```

```
psql (17beta3 (Debian 17~beta3-1.pgdg120+1))
```

Type "help" for help.

```
pg17=# \dy
```

List of event triggers					
Name	Event	Owner	Enabled	Function	Tags
on_login_trigger	login	postgres	always	mandate_app_info	

(1 row)

```
pg17=# \dconfig event_triggers
```

List of configuration parameters	
Parameter	Value
event_triggers	on

(1 row)

Connection allowed

configuration to control logon trigger

# Getting Started and Test your self.

---

Running PostgreSQL 17 Beta within Docker.

```
docker run -d --name=pg17b3 -e POSTGRES_PASSWORD=*****
postgres:17beta3 -p 5432:5435
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
f660743d3f2a	postgres:17beta3	"docker-entrypoint.s..."	6 days ago	Up 46 hours	5432/tcp	pg17b3

```
docker exec -it pg17b3 psql -h localhost -U postgres -d postgres
```

# Thank you!

---



<https://www.linkedin.com/in/mahtodeepak/>



<https://x.com/mahtodeepak05>



<https://databaserookies.wordpress.com/>

DATA CLOUD GAZE

<https://www.datacloudgaze.com/>

DATA CLOUD GAZE

<https://www.databasegyaan.com>



Code : PGHYD30 , 30% off