


What's New in PostgreSQL 18



Vinoth Kanna RS

Founding Partner, Mydbops LLP

Mydbops MyWebinar 47

About Me




Vinoth Kanna R S

- ❑ Making Database Management Simple at Scale
- ❑ Leading remote database team at Mydbops
- ❑ I enjoy working with various database technologies and scaling backend systems.
- ❑ Tech Speaker on Open Source Events
- ❑ Databases for living (13 years)



**Your Trusted
Open Source Database
Management Partner**

With 9+ Years of Expertise



Mydbops by the Numbers



9+ years

Of Expertise



800 +

Happy Clients



10 B +

DB Transactions
Handled per Day



6000 +

Servers
Monitored



3000 +

Tickets Handled
per Day

Database Technologies



Agenda

- ❑ Asynchronous IO
- ❑ Vacuuming
- ❑ Persistent statistics on Upgrade
- ❑ UUUIDV7
- ❑ Virtual Columns
- ❑ Index Skip Scans
- ❑ Non Overlapping range Constraints
- ❑ Other Features
- ❑ Cloud Readiness
- ❑ Q&A

Asynchronous IO

Asynchronous IO: The Foundations (PgSQL 17)

- ❑ Streaming and Vectored I/O
- ❑ 1 preadv() (128kB) call can replace 16 read() (8kB) calls
- ❑ Uses posix_fadvise() to prefetch to OS page cache
- ❑ ReadBuffer is replaced by Streaming API (ReadStream)
 - ❑ pg_prewarm - loads a full relation (table, index) into shared buffers
 - ❑ Faster ANALYZE
 - ❑ Faster sequential scans, OLAP workloads

Asynchronous IO: The Present (PgSQL 18)

- ❑ Introduced variable “io_method” to control, configure and optimise
- ❑ Extended ReadSteam usage for
 - ❑ Bitmap heap scans
 - ❑ Vacuum
- ❑ “Effective_io_concurrency” directly controls the prefetch
- ❑ Stats are available via table **pg_aio** to monitor the performance

Asynchronous IO: io_method

Method	How It Works	Best For
<code>sync</code>	Disables AIO.	Baseline testing and troubleshooting.
<code>worker (Default)</code>	Uses a pool of background workers to handle I/O. Portable and robust.	The recommended setting for most environments.
<code>io_uring</code>	A high-performance, Linux-only method using a shared buffer with the kernel for minimal overhead.	Modern Linux systems where Postgres is compiled with liburing support for maximum speed.

Asynchronous IO: io_method

```
vinoth=# select version();

PostgreSQL 18.0 (Homebrew), compiled by Apple clang version 17.0.0 (clang-1700.0.13.3), 64-bit
(1 row)

vinoth=# show io_method;
          io_method
-----
 worker
(1 row)

vinoth=# show io_workers;
      io_workers
-----
 3
(1 row)
```

Asynchronous IO: io_method

```
vinoth=# select datid, datname, pid, usename, wait_event, state, left(query,10), backend_type from pg_stat_activity;
   datid | datname | pid | usename |      wait_event      | state |      left      |           backend_type
-----+-----+-----+-----+-----+-----+-----+-----+
-----+
16393 | vinoth | 92569 | vinoth |          | active | select dat | client backend
      |         | 92204 |          | AutovacuumMain |       |          | autovacuum launcher
      |         | 92190 |          | IoWorkerMain   |       |          | io worker
      |         | 92194 |          | IoWorkerMain   |       |          | io worker
      |         | 92196 |          | IoWorkerMain   |       |          | io worker
      |         | 92200 |          | CheckpointerMain |       |          | checkpointer
      |         | 92201 |          | BgwriterMain   |       |          | background writer
      |         | 92203 |          | WalWriterMain  |       |          | walwriter
(8 rows)
```

Asynchronous IO: The Future (PgSQL 19+)

- ❑ AIO for writes (Checkpointers, BG writers, Backend Writes Etc)
- ❑ AIO integration with Network IO
- ❑ Expose other OS methods like IOCP for windows, Just like IO_URING For Linux
- ❑ POSIX aio_read/aio_write for solaris

FAQ ?

- ❑ Will AIO in PgSQL 18 make the writes faster ?
 - ❑ No, Current AIO implementation aimed to speed up the reads, Implementation of AIO for writes is in progress, Will be shipped in the later releases.
- ❑ What value i can set for io_method to get best performance ?
 - ❑ OS other than linux, only supports worker model, Tune the threads according to server capacity, CPU. Watch on memory as you increase worker threads.
 - ❑ In case of Linux, IO_Uring would be recommended, Benchmark claims it's suited for Random IO, Watch on Connections and Open files limits.

Vacuuming

Vacuum Flow



Vacuum: (PgSQL 18)

- ❑ Uses **read_stems** in all lazy vacuum phases (1, 2 and 3)
- ❑ Supported in B-Tree, GIST, SPGIST indexes
- ❑ New columns total vacuum, auto vacuum, analyze, auto analyze added on pg_stat_*tables

Vacuum: (PgSQL 19+)

- ❑ Parallel vacuum in Phase 1, 3
- ❑ Decoupling vacuum for table and index
- ❑ Running specific vacuum phases independently


Persistent statistics on Upgrade

Persistent statistics on Upgrade

- ❑ On execution of pg_upgrade, statistics will be automatically copied
 - ❑ We have tested this during the upgrade between PostgreSQL 17 to PostgreSQL 18
- ❑ We will be able to dump the statistics manually from old version to restore in the new version
 - ❑ You should be using the pg_dump and pg_restore client tools that comes with PgSQL 18
 - ❑ pg_dump --statistics-only -Fc <DB> > stats.dump
 - ❑ pg_restore --statistics-only -d <DB> stats.dump
- ❑ Execute vacuumdb --missing-stats-only
 - ❑ compute only missing optimizer statistics

UUIDV7

UUIDV4 (PgSQL 17 and Older)




UUIDV4 (PgSQL 17 and Older)

- ❑ Issues when using random generated UUIDV4
 - ❑ Not possible to sort, Usually need other timestamp columns
 - ❑ Diverse index locality due to random ordering
 - ❑ Table bloating / Fragmentation due to unsorted nature of the primary key
- ❑ Blog reference on bloating between random UUID4 and sequentially generated UUID's
 - ❑ <https://www.enterprisedb.com/blog/sequential-uuid-generators>

UUIDV4 (PgSQL 17 and Older)

- ❑ How sequential UUID was generated on the older versions ?
 - ❑ Extensions
 - ❑ <https://github.com/dverite/postgres-uuidv7-sql/blob/main/README.md>
 - ❑ Custom Function
 - ❑ <https://instagram-engineering.com/sharding-ids-at-instagram-1cf5a71e5a5c>
 - ❑ Application logic

UUIDv7 (PgSQL 18+)



UUIDV7 (PgSQL 18+)

```
vinoth=# SELECT uuidv7();
          uuidv7
-----
019a38ce-4774-7b52-a4d3-1155e5d7e058
(1 row)

vinoth=# select uuid_extract_timestamp(uuidv7());
          uuid_extract_timestamp
-----
2025-10-31 11:17:24.123+05:30
(1 row)

vinoth=# select now();
           now
-----
2025-10-31 11:17:31.807555+05:30
(1 row)
```

UUIDV7 (PgSQL 18+)

```
vinoth=# CREATE TABLE test (
    id uuid DEFAULT uuidv7() PRIMARY KEY,
    name text
);
INSERT INTO test (name) VALUES ('foo');
INSERT INTO test (id, name) VALUES (uuidv7(INTERVAL '-1 day'), 'old3');
INSERT INTO test (id, name) VALUES (uuidv7(INTERVAL '-7 day'), 'old2');
INSERT INTO test (id, name) VALUES (uuidv7(INTERVAL '-24 day'), 'old1');

vinoth=# select uuid_extract_timestamp(id), id, name from test;
      uuid_extract_timestamp      |          id          |      name
-----+-----+-----+
2025-10-31 11:17:57.729+05:30 | 019a38ce-f1a1-7deb-8617-2e17342e38b5 | bar
2025-10-30 11:17:57.73+05:30  | 019a33a8-95a2-7089-a576-5a9d5d00a6d7 | old3
2025-10-24 11:17:57.73+05:30  | 019a14c2-6da2-72c7-9eeb-dc05308ebb91 | old2
2025-10-07 11:17:57.73+05:30  | 0199bd36-51a2-7473-ae01-3dc84c5b0fb0 | old1
```

UUIDV7 (PgSQL 18+)

```
vinoth=# explain select uuid_extract_timestamp(id), id, name from test where uuid_extract_timestamp(id) < '2025-10-29';
Seq Scan on test  (cost=0.00..26.94 rows=357 width=56)
  Filter: (uuid_extract_timestamp(id) < '2025-10-29 00:00:00+05:30'::timestamp with time zone)
(2 rows)

vinoth=# create index idx_uuid_ts on test (uuid_extract_timestamp(id));
CREATE INDEX

vinoth=# explain select uuid_extract_timestamp(id), id, name from test where uuid_extract_timestamp(id) < '2025-10-29';
Seq Scan on test  (cost=0.00..1.08 rows=2 width=56)
  Filter: (uuid_extract_timestamp(id) < '2025-10-29 00:00:00+05:30'::timestamp with time zone)
(2 rows)
```

Index Skip Scans

Index Skip Scans

- ❑ Making use of a non-leading column of an index
- ❑ Help to reduce the number of indexes
- ❑ Skip scan works by generating a dynamic equality constraint internally, that matches every possible value in an index column
- ❑ Eg:

Index: (is_deleted, status)

Query: select * from table where status = 'new';

Virtual Generated Columns

Virtual Generated Columns

- ❑ PostgreSQL 18 makes VIRTUAL the new default for generated columns.
 - ❑ STORED (The old way): Computed on write, takes up disk space. Adding one required a full table rewrite and a disruptive lock.
 - ❑ VIRTUAL (The new default): Computed on read, takes up zero storage space.
- ❑ Adding a virtual column is now an instantaneous, metadata-only operation. No more waiting hours for ALTER TABLE on a huge table.
- ❑ Generated virtual column values are calculated on the fly during the execution of the query

Virtual Generated Columns

```
vinoth=# CREATE TABLE users (
    first_name TEXT,
    last_name TEXT,
    full_name TEXT GENERATED ALWAYS AS (first_name || ' ' || last_name) VIRTUAL
);
CREATE TABLE

vinoth=# insert into users values ('vinoth','mydbops');
INSERT 0 1

vinoth=# select * from users;
first_name | last_name |   full_name
-----+-----+-----
vinoth     | mydbops   | vinoth mydbops
(1 row)
```

Extension of RETURNING SQL Clause

Extension of RETURNING SQL Clause

- ❑ OLD and NEW support for RETURNING clauses in INSERT, UPDATE, DELETE, and MERGE commands.
- ❑ Easier to implement the auditing functionality without needing for triggers

```
UPDATE products SET price = price * 1.10
    WHERE price <= 99.99
    RETURNING name, old.price AS old_price, new.price AS new_price,
            new.price - old.price AS price_change;

DELETE FROM products
    WHERE obsolescence_date = 'today'
    RETURNING *;
```

Data Checksums

Data Checksum

- ❑ Data checksum are on by default in PostgreSQL 18
 - ❑ If you're upgrading from instance with checksum turned off, Ensure to initialize DB on PostgreSQL 18 by disabling checksum.
 - ❑ `initdb --no-data-checksums </datadir>`
- ❑ Best practise is to have the checksum turned on, Logical restore and replication can be used to upgrade with checksum turned on

Other Features

- ❑ Add vacuumdb option --missing-stats-only to compute only missing optimizer statistics
 - ❑ To be run post upgrade
- ❑ Pg_overexplain
 - ❑ Understand the internal heuristics behind index, optimisations selection
- ❑ Oauth support

Public Cloud Support

Public Cloud Support

- ❑ AWS
 - ❑ Preview, Available only in Ohio Region
 - ❑ No option in parameter group to tune io_method
- ❑ GCP
 - ❑ Available
 - ❑ io_uring method not available for io_method, Only sync, worker is available
- ❑ Azure
 - ❑ Preview, Available only in the East Asia region
 - ❑ io_uring method not available for io_method, Only sync, worker is available

Key Takeaways

Key Takeaways

- ❑ Linux VM based instances will be best to test all the features
- ❑ Validate compatibility for upgrade with pg_upgrade --check
- ❑ If you want to explore all the new features, Linux based VM would be the best option
- ❑ When tuning for AIO, Benchmark with real time data size, queries, server configuration
- ❑ Benchmarks from the synthetic benchmarks from sysbench, pgbench should be only taken for reference

Connect with us !



Reach us at: info@mydbops.com

Thank You