

Data processing more than billion rows per second

~unleash full capability of GPU and NVME~

HeteroDB, Inc
Chief Architect & CEO
KaiGai Kohei <kaigai@heterodb.com>

Self-Introduction

about myself

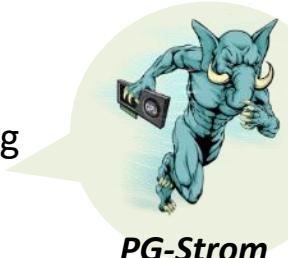


- KaiGai Kohei (海外浩平)
- Chief Architect & CEO of HeteroDB
- Contributor of PostgreSQL (2006-)
- Primary Developer of PG-Strom (2012-)
- Interested in: Big-data, GPU, NVME/PMEM, ...

about our company

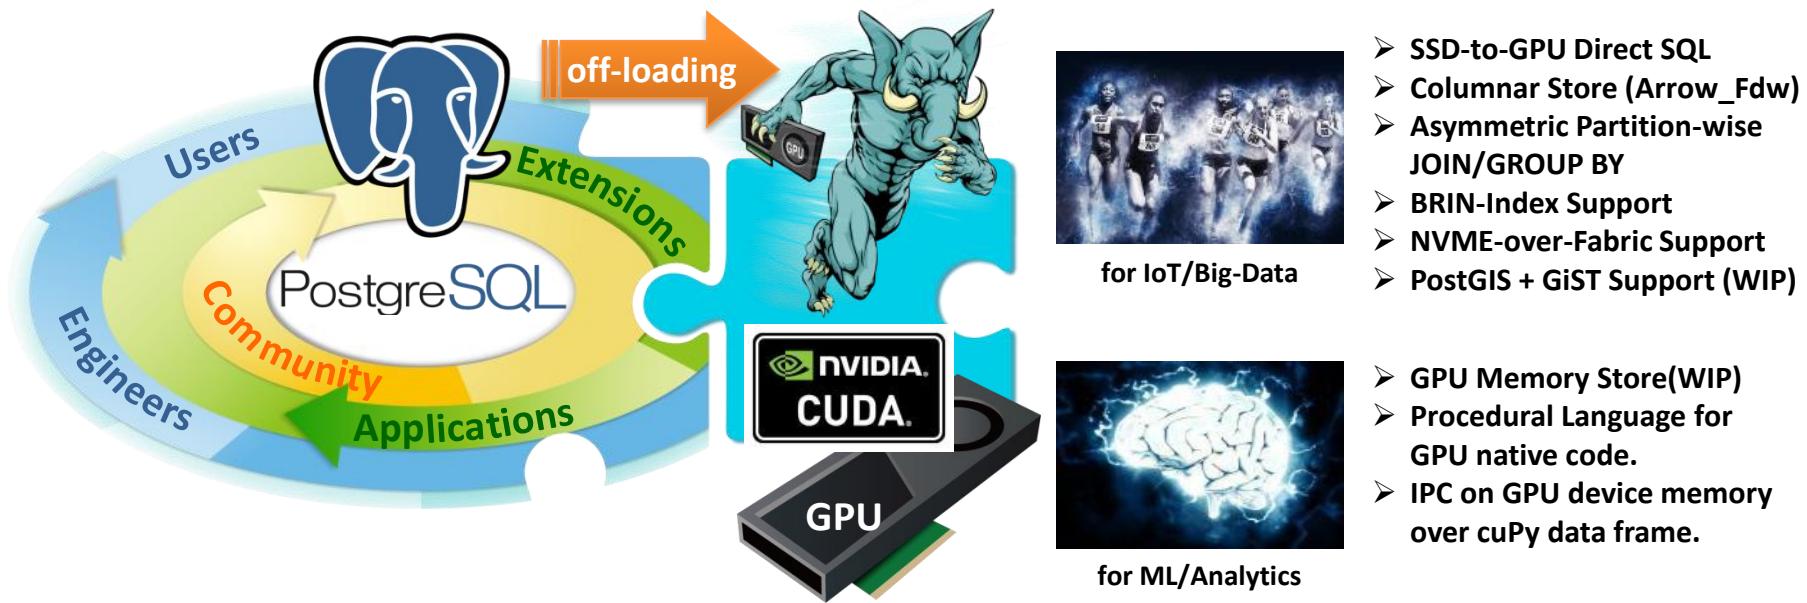


- Established: 4th-Jul-2017
- Location: Shinagawa, Tokyo, Japan
- Businesses:
 - ✓ Sales & development of high-performance data-processing software on top of heterogeneous architecture.
 - ✓ Technology consulting service on GPU&DB area.



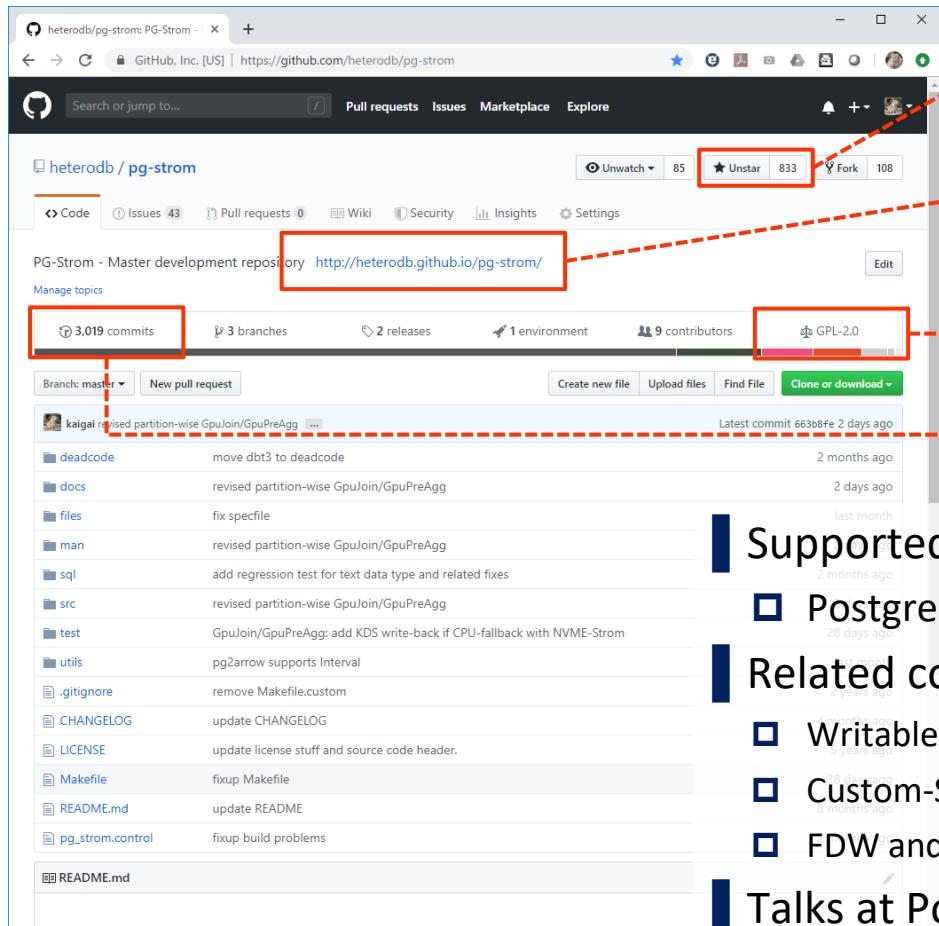
What is PG-Strom?

PG-Strom is an extension of PostgreSQL for terabytes scale data-processing and in-database analytics, by utilization of GPU and NVME-SSD.



- Transparent GPU acceleration for analytics and reporting workloads
- Binary code generation by JIT from SQL statements
- PCIe-bus level optimization by SSD-to-GPU Direct SQL technology
- Columnar-storage for efficient I/O and vector processors

PG-Strom as an open source project



Many stars 😊

Official documentation in English / Japanese
<http://heterodb.github.io/pg-strom/>

Distributed under GPL-v2.0

Has been developed since 2012

Supported PostgreSQL versions

- PostgreSQL v12, v11, and v10. WIP for v13 support

Related contributions to PostgreSQL

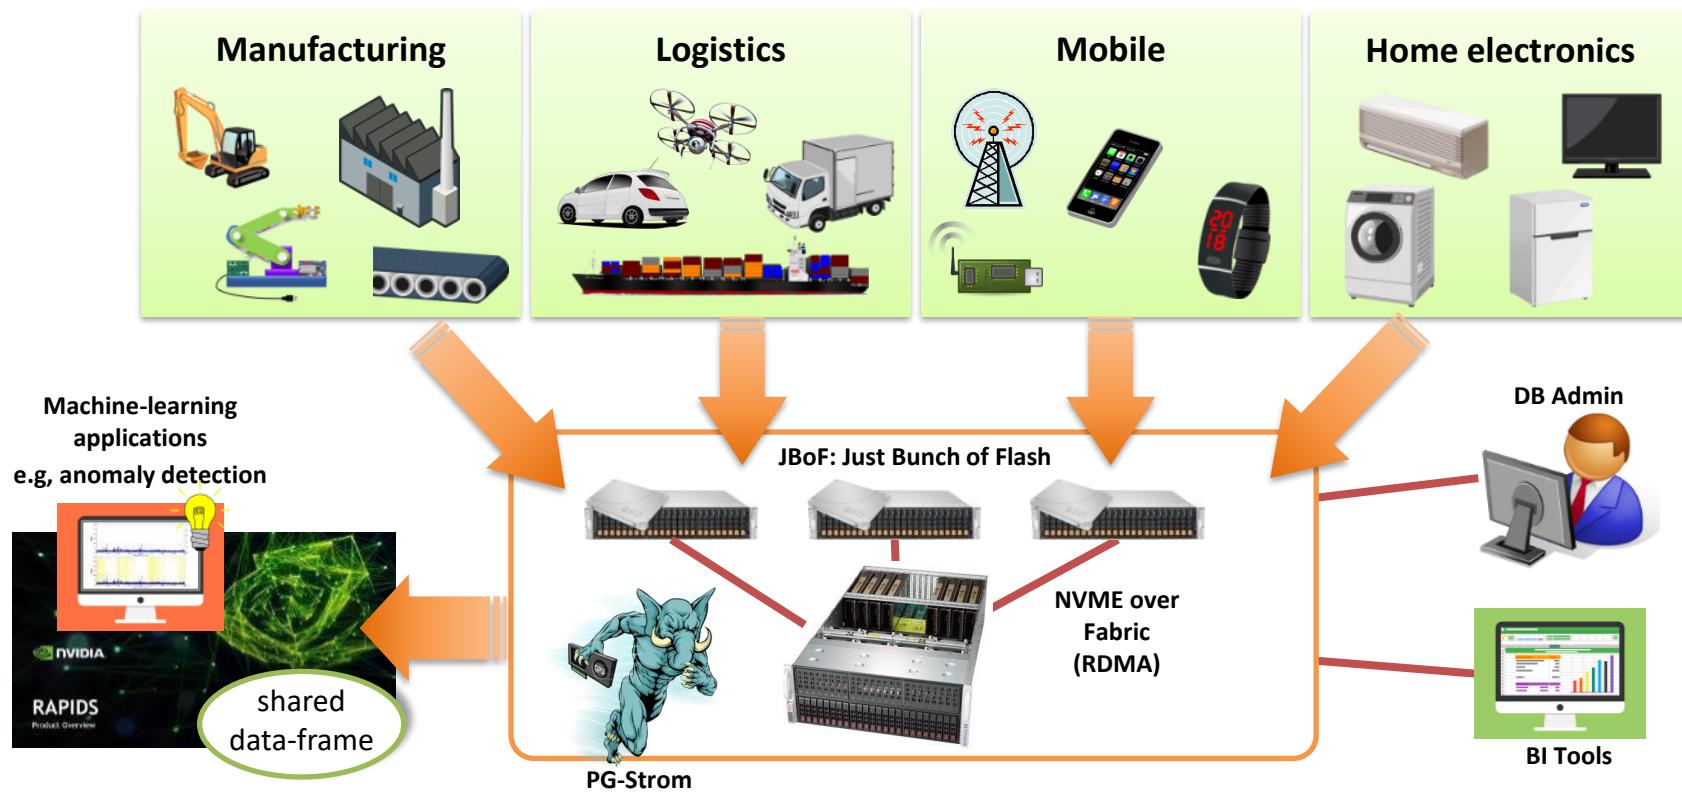
- Writable FDW support (9.3)
- Custom-Scan interface (9.5)
- FDW and Custom-Scan JOIN pushdown (9.5)

Talks at PostgreSQL community

- PGconf.EU 2012 (Prague), 2015 (Vienna), 2018 (Lisbon)
- PGconf.ASIA 2018 (Tokyo), 2019 (Bali, Indonesia)
- PGconf.SV 2016 (San Francisco)
- PGconf.China 2015 (Beijing)

Target Workloads?

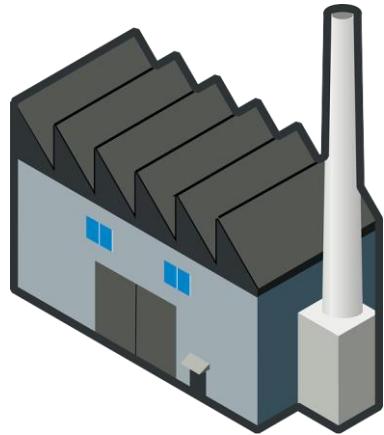
Target: IoT/M2M Log-data processing



- Any kind of devices generate various form of log-data.
- People want/try to find out insight from the data.
- Data importing and processing must be rapid.
- System administration should be simple and easy.

How much data size we expect? (1/3)

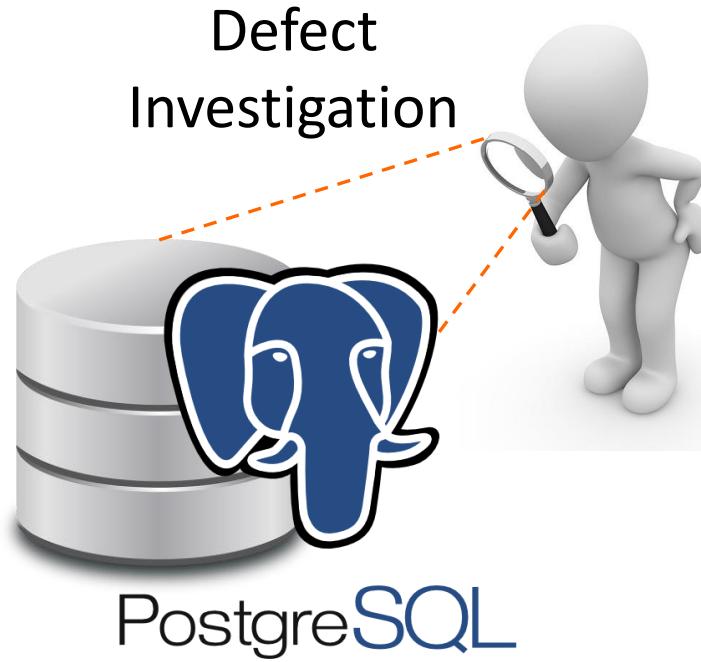
Case: A semiconductor factory managed up to 100TB with PostgreSQL



Total **20TB**
Per Year



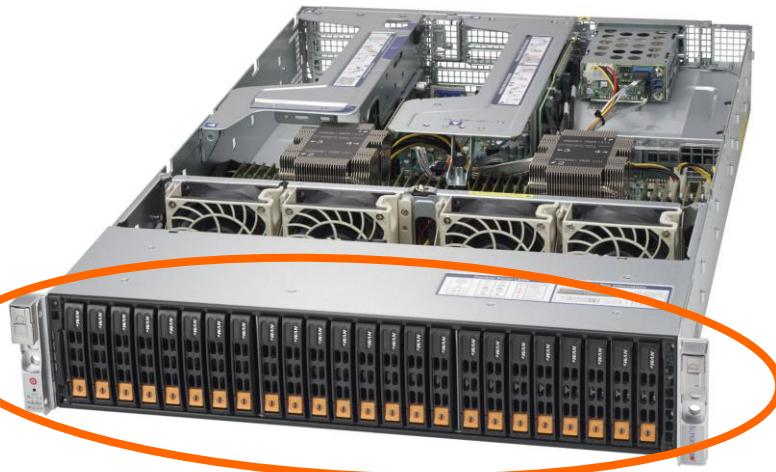
Kept for
5 years



- This is a factory of an international top-brand company.
- Most of users don't have bigger data than them.

How much data size we expect? (2/3)

Nowadays, 2U server is capable for 100TB capacity



8.0TB x 24 = 192TB

How much is it?

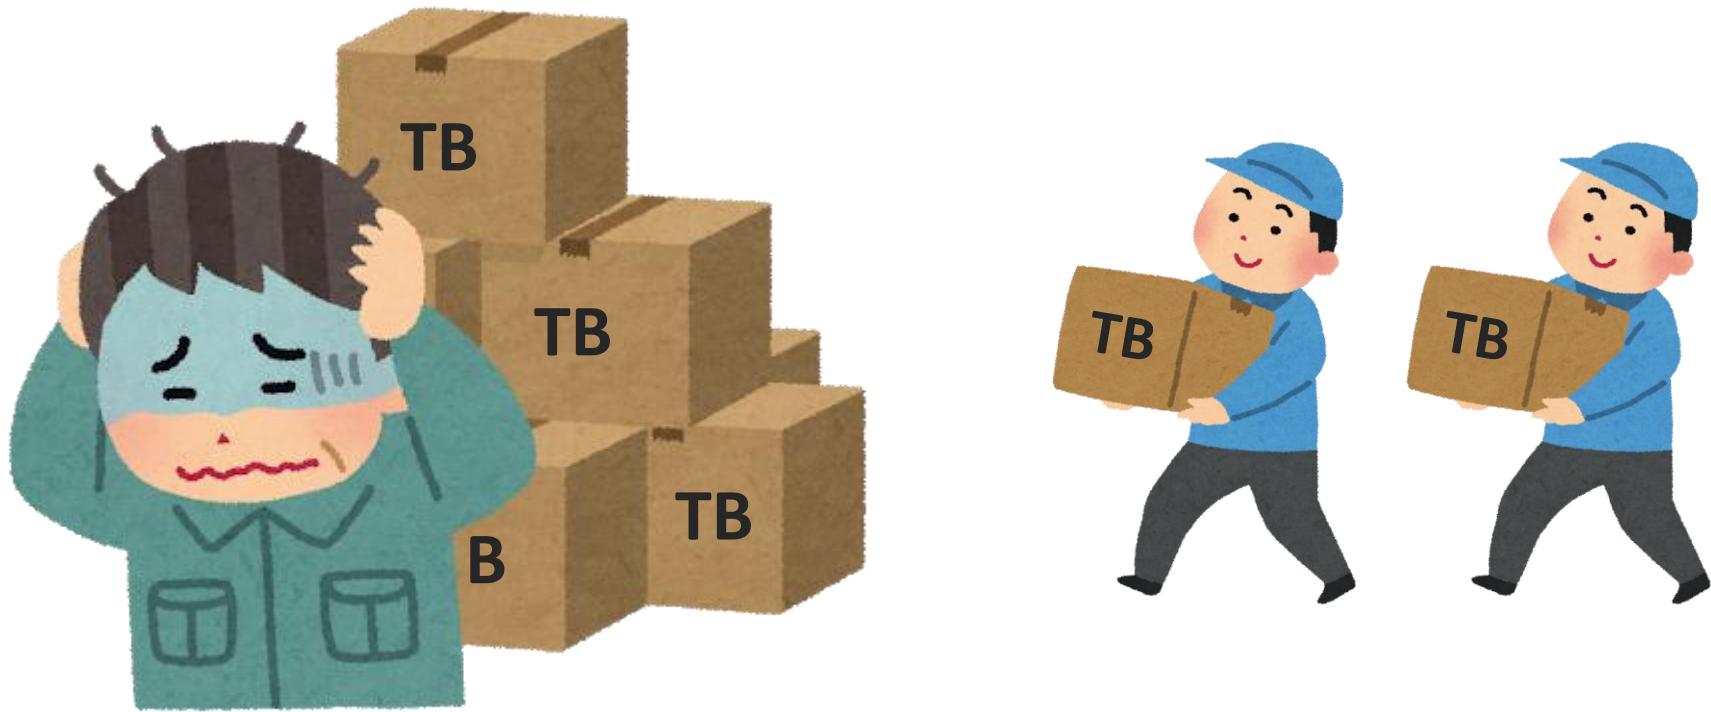
60,932USD

by thinkmate.com
(31st-Aug-2019)

| model | Supermicro 2029U-TN24R4T | Qty |
|-------|------------------------------------|-----|
| CPU | Intel Xeon Gold 6226 (12C, 2.7GHz) | 2 |
| RAM | 32GB RDIMM (DDR4-2933, ECC) | 12 |
| GPU | NVIDIA Tesla P40 (3840C, 24GB) | 2 |
| HDD | Seagate 1.0TB SATA (7.2krpm) | 1 |
| NVME | Intel DC P4510 (8.0TB, U.2) | 24 |
| N/W | built-in 10GBase-T | 4 |

How much data size we expect? (3/3)

On the other hands, it is not small enough.



Efficient Storage

- SSD-to-GPU Direct SQL



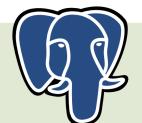
Efficient Data Structure

- Arrow_Fdw



Efficient Data Deployment

- Table Partitioning



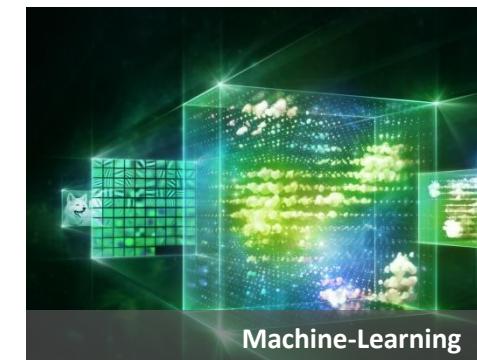
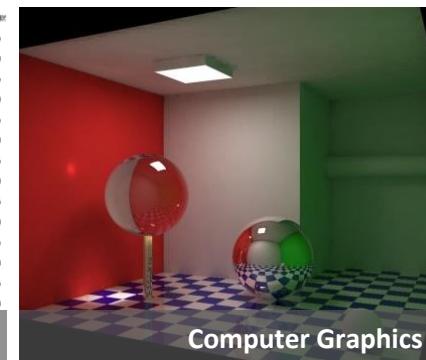
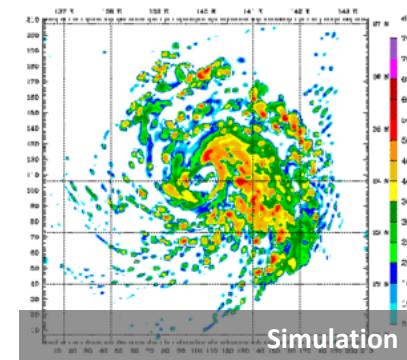
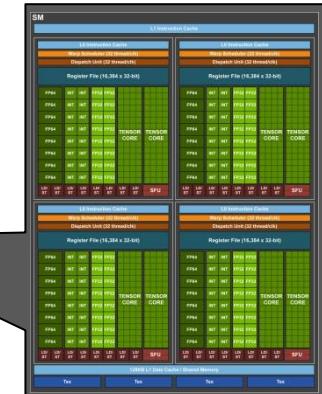
PG-Strom: SSD-to-GPU Direct SQL

GPU's characteristics - mostly as a computing accelerator

Over 10years history in HPC, then massive popularization in Machine-Learning

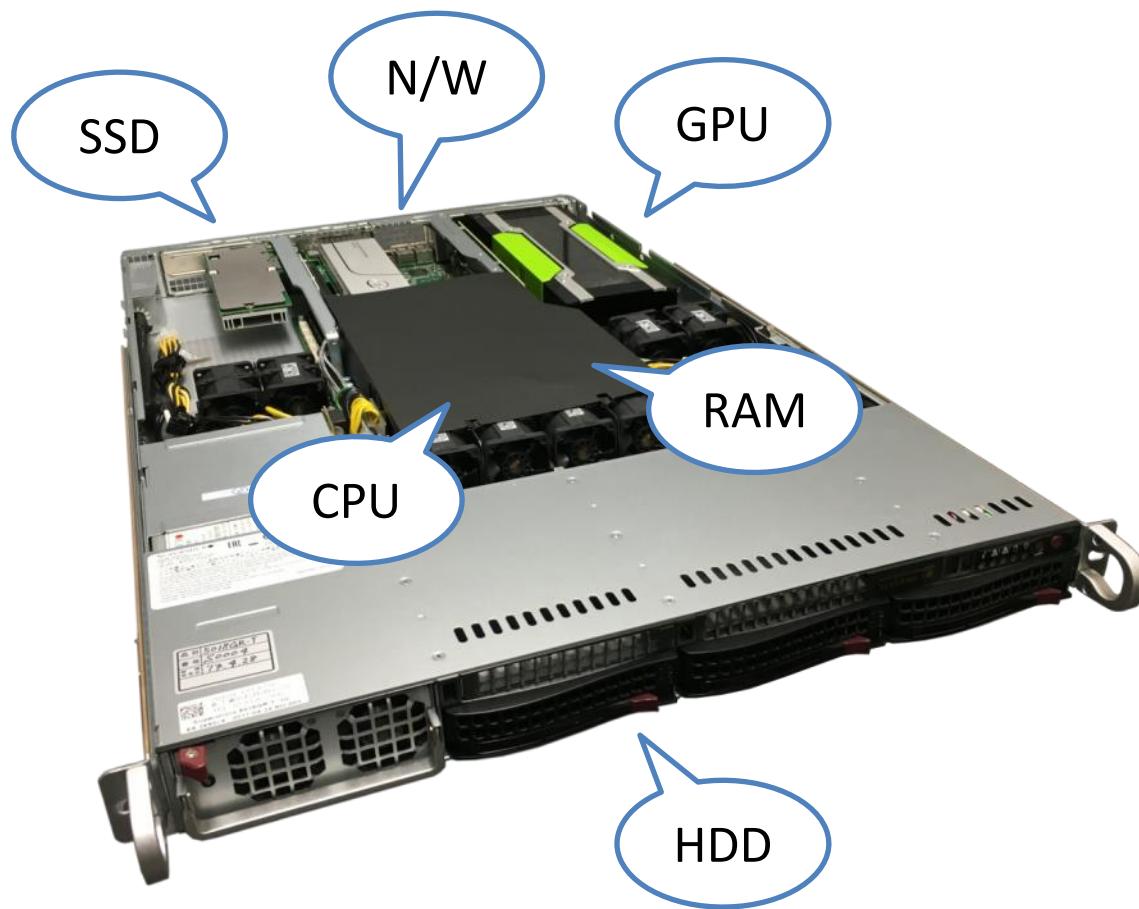


NVIDIA Tesla V100



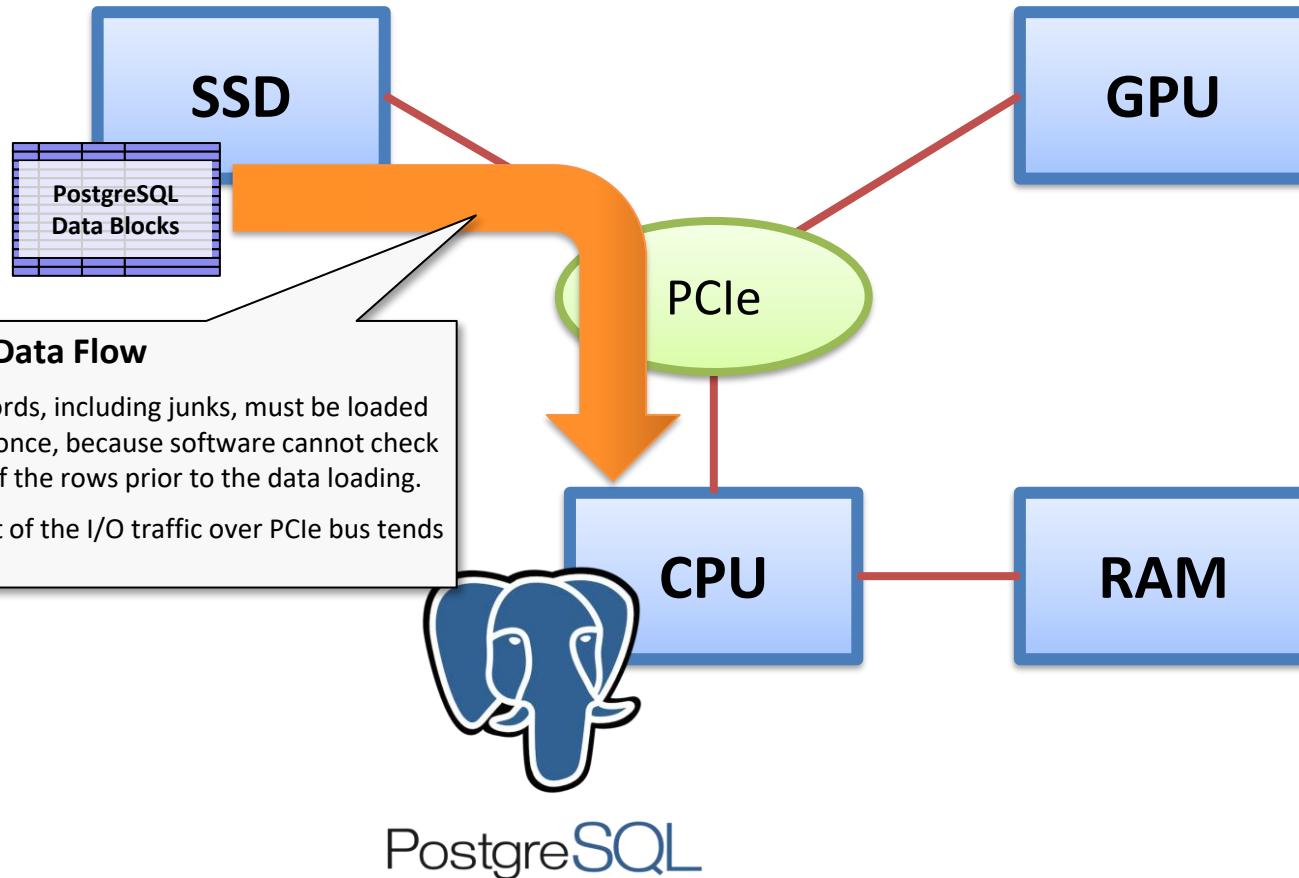
How I/O workloads are accelerated by GPU that is a **computing accelerator**?

A usual composition of x86_64 server

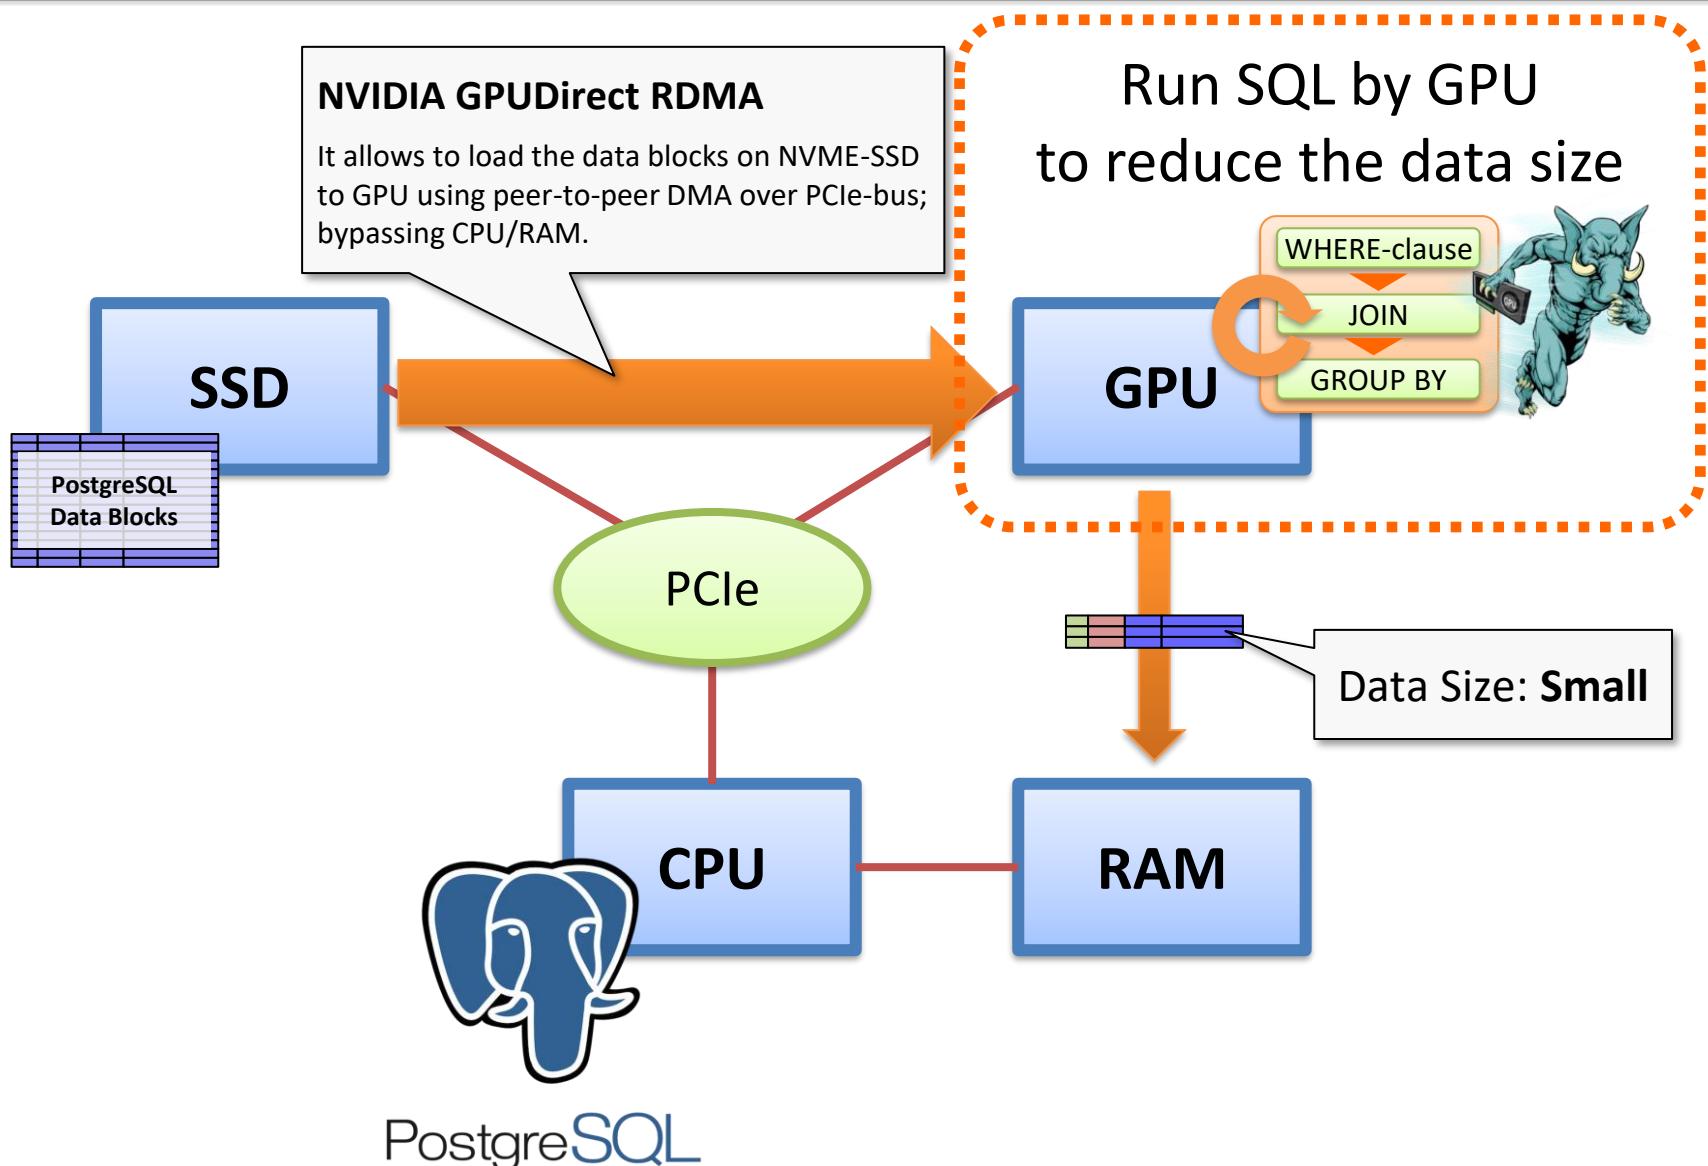


Data flow to process a massive amount of data

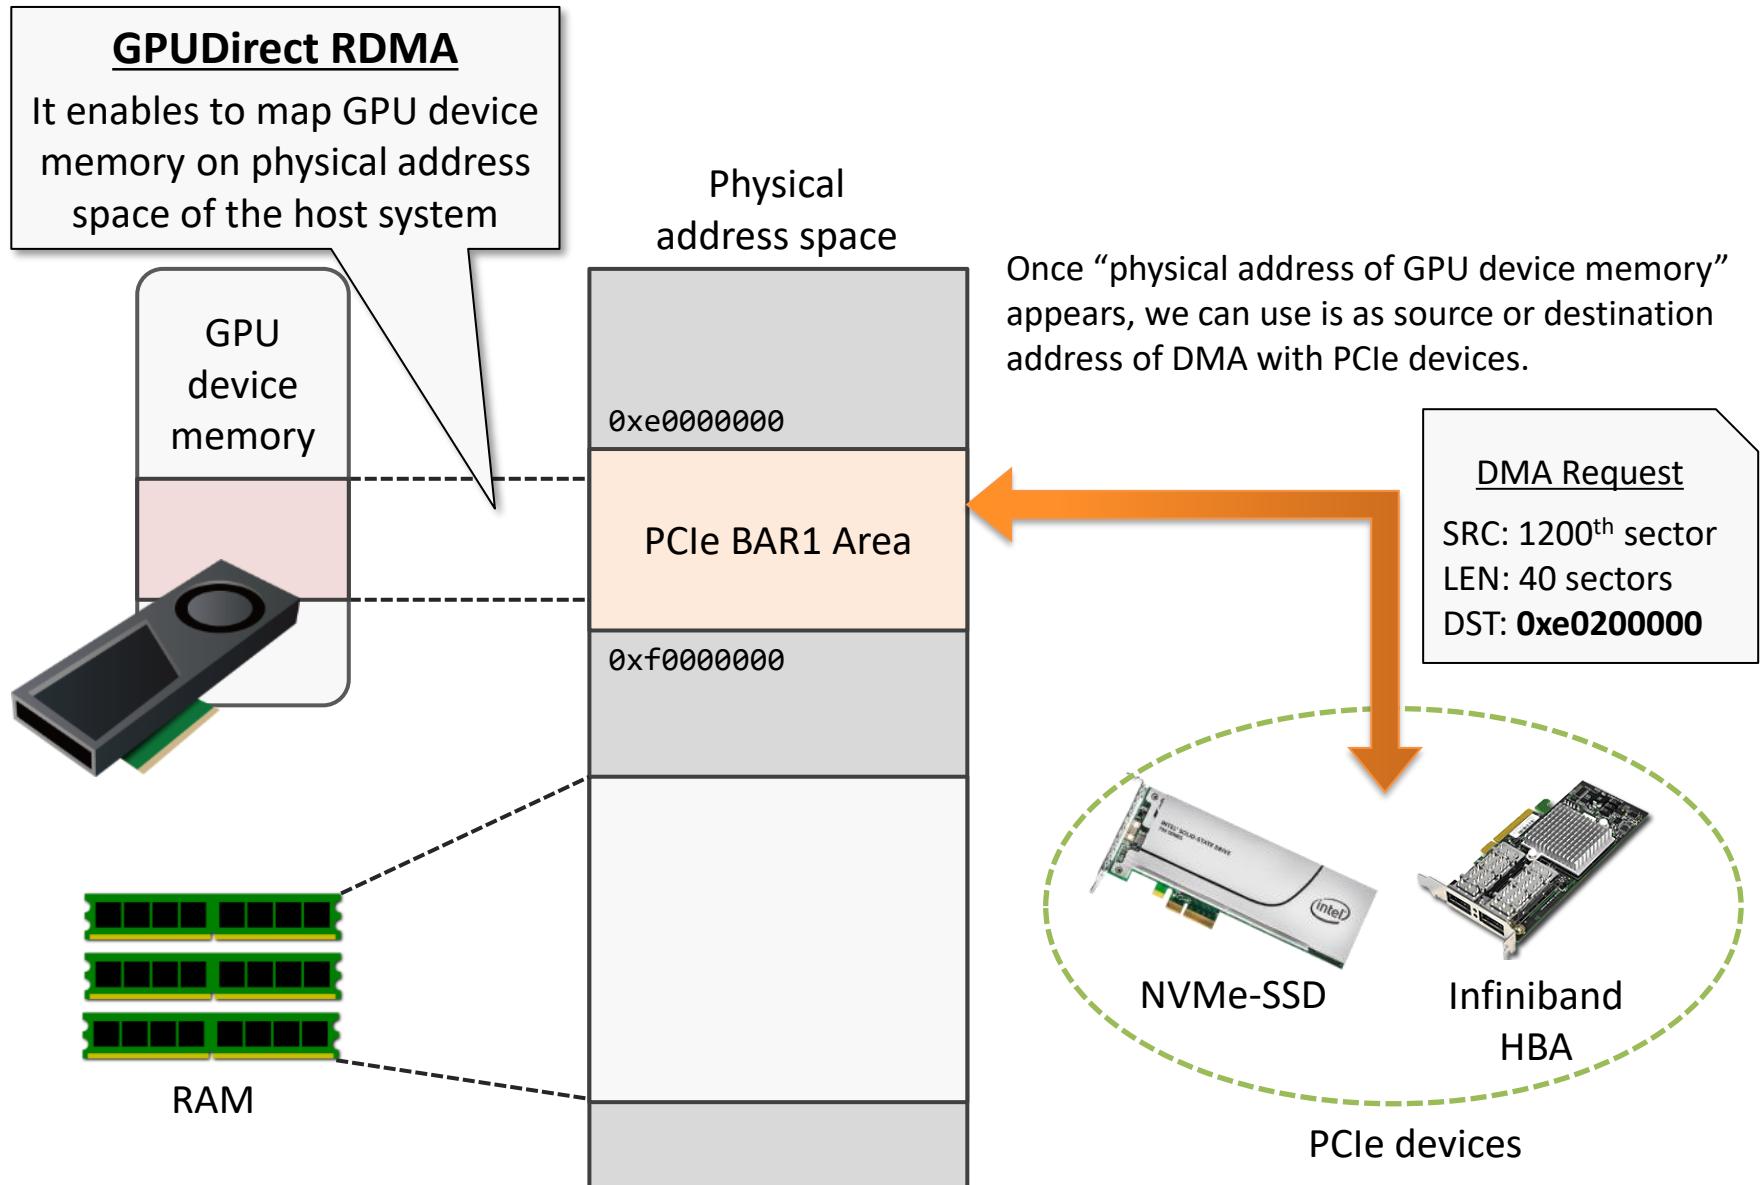
Unless records are not loaded to CPU/RAM once, over the PCIe bus, software cannot check its necessity even if they are “junk” records.



SSD-to-GPU Direct SQL (1/4) – Overview



Background - GPUDirect RDMA by NVIDIA



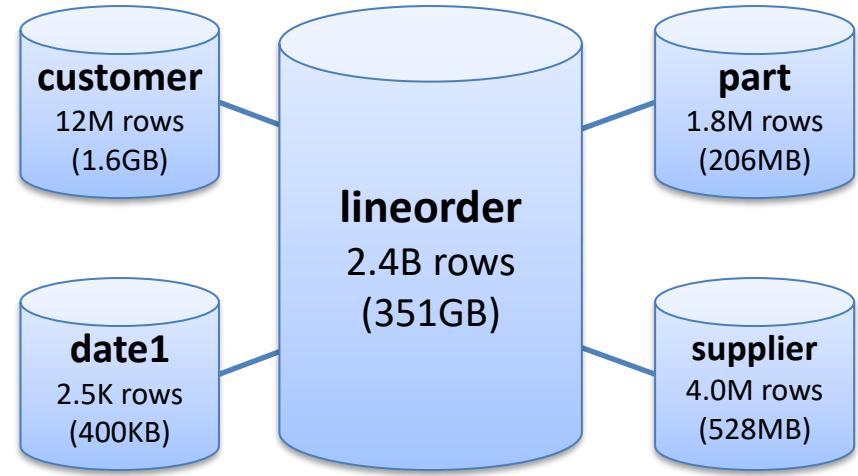
SSD-to-GPU Direct SQL (2/4) – System configuration and workloads

Summarizing queries for typical Star-Schema structure on simple 1U server



Supermicro SYS-1019GP-TT

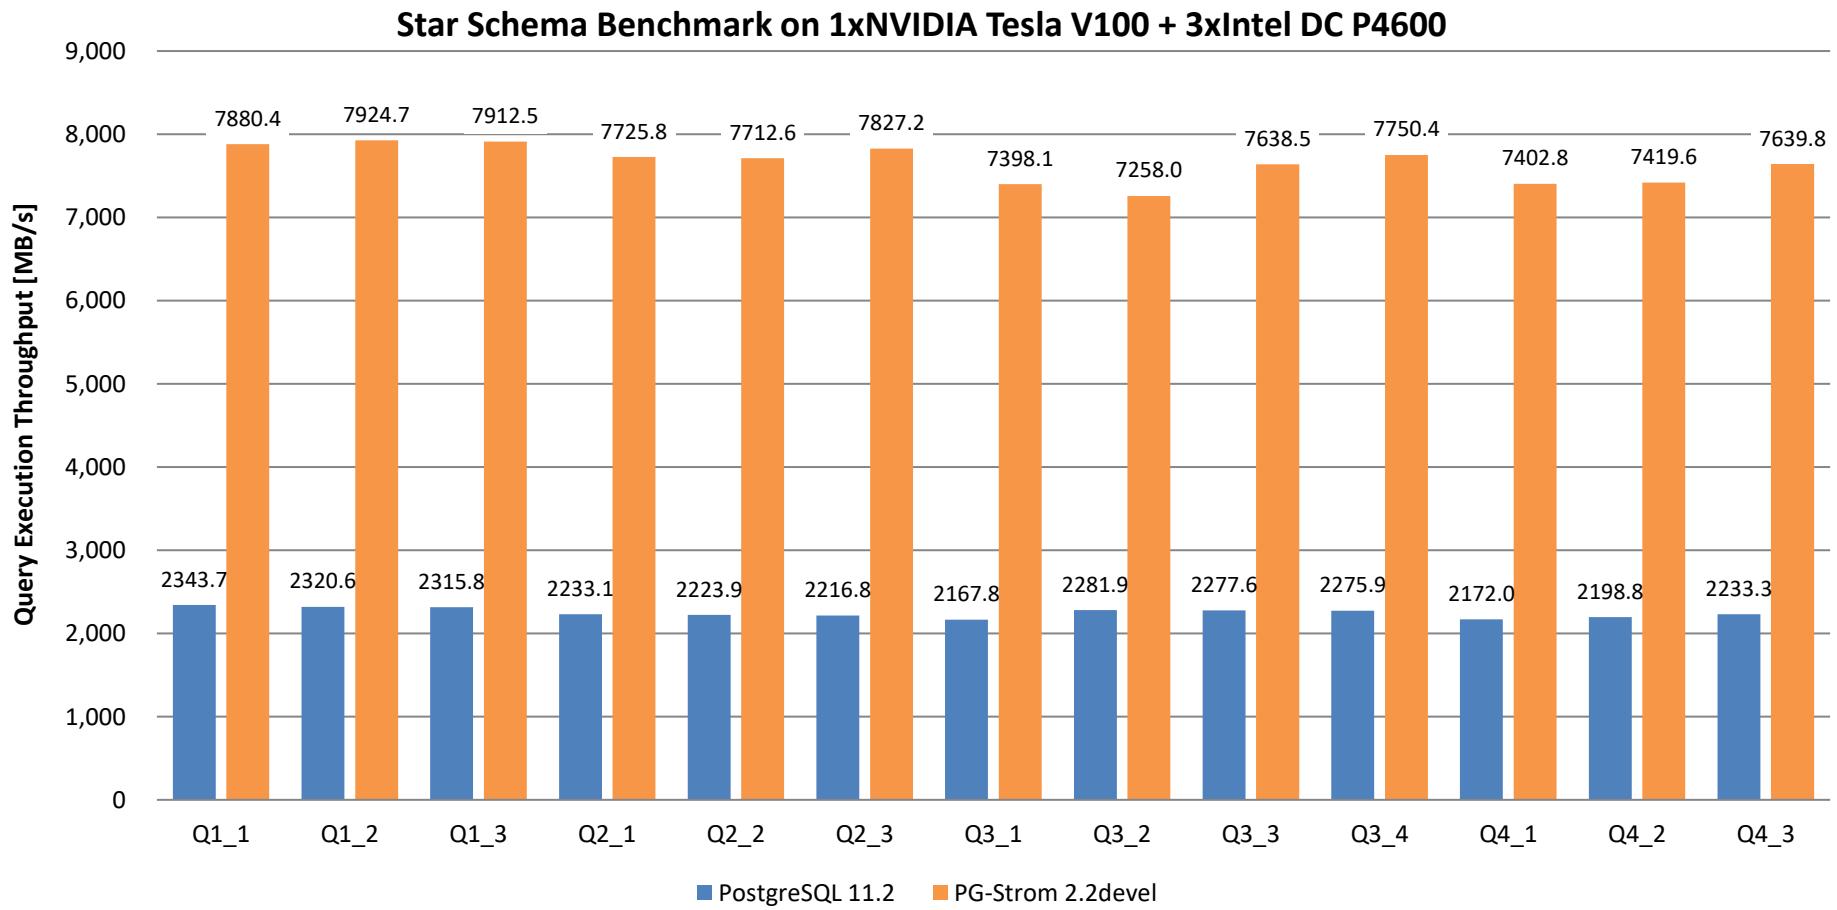
| | |
|---------|---|
| CPU | Xeon Gold 6126T (2.6GHz, 12C) x1 |
| RAM | 192GB (32GB DDR4-2666 x 6) |
| GPU | NVIDIA Tesla V100 (5120C, 16GB) x1 |
| SSD | Intel SSD DC P4600 (HHHL; 2.0TB) x3 (striping configuration by md-raid0) |
| HDD | 2.0TB(SATA; 72krpm) x6 |
| Network | 10Gb Ethernet 2ports |
| OS | Red Hat Enterprise Linux 7.6 CUDA 10.1 + NVIDIA Driver 418.40.04 |
| DB | PostgreSQL v11.2 PG-Strom v2.2devel |



■ Query Example (Q2_3)

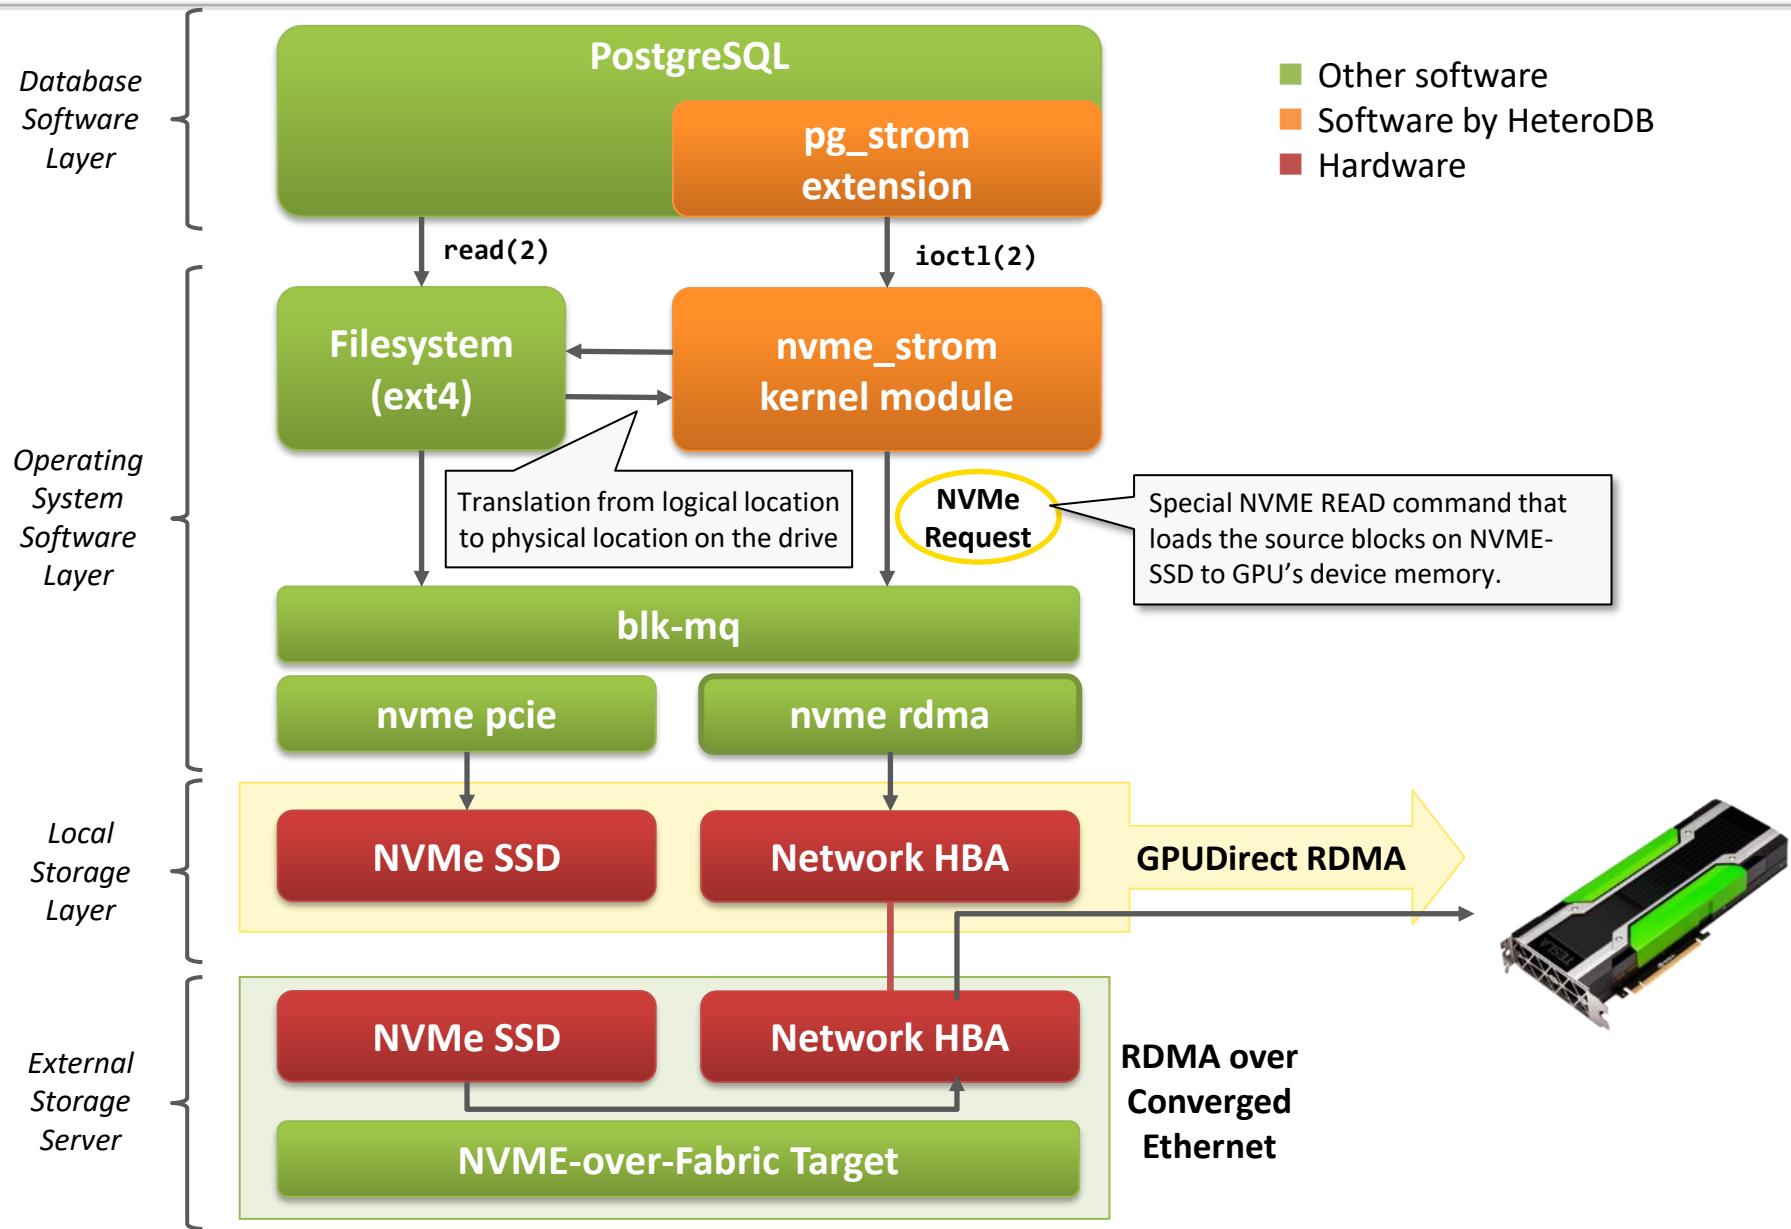
```
SELECT sum(lo_revenue), d_year, p_brand1
  FROM lineorder, date1, part, supplier
 WHERE lo_orderdate = d_datekey
   AND lo_partkey = p_partkey
   AND lo_suppkey = s_suppkey
   AND p_category = 'MFGR#12'
   AND s_region = 'AMERICA'
 GROUP BY d_year, p_brand1
 ORDER BY d_year, p_brand1;
```

SSD-to-GPU Direct SQL (3/4) – Benchmark results



- Query Execution Throughput = (353GB; DB-size) / (Query response time [sec])
- SSD-to-GPU Direct SQL runs the workloads close to the hardware limitation (8.5GB/s)
- about x3 times faster than filesystem and CPU based mechanism

SSD-to-GPU Direct SQL (4/4) – Software Stack



PG-Strom: Arrow_Fdw (columnar store)

Background: Apache Arrow (1/2)

Characteristics

- Column-oriented data format designed for analytics workloads
- A common data format for inter-application exchange
- Various primitive data types like integer, floating-point, date/time and so on

| | session_id | timestamp | source_ip |
|-----|------------|-----------------|----------------|
| w 1 | 1331246660 | 3/8/2012 2:44PM | 99.155.155.225 |
| w 2 | 1331246351 | 3/8/2012 2:38PM | 65.87.165.114 |
| w 3 | 1331244570 | 3/8/2012 2:09PM | 71.10.106.181 |
| w 4 | 1331261196 | 3/8/2012 6:46PM | 76.102.156.138 |

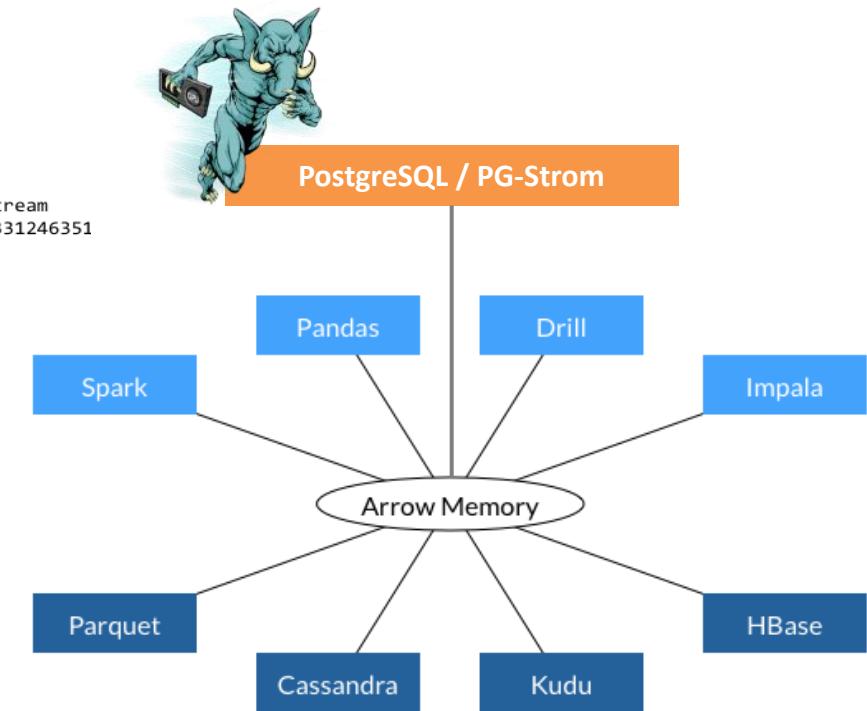
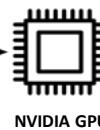
Traditional Memory Buffer

| | |
|------|-----------------|
| w 1 | 1331246660 |
| w 2 | 3/8/2012 2:44PM |
| w 3 | 99.155.155.225 |
| w 4 | 1331246351 |
| w 5 | 3/8/2012 2:38PM |
| w 6 | 65.87.165.114 |
| w 7 | 1331244570 |
| w 8 | 3/8/2012 2:09PM |
| w 9 | 71.10.106.181 |
| w 10 | 1331261196 |
| w 11 | 3/8/2012 6:46PM |
| w 12 | 76.102.156.138 |

Arrow Memory Buffer

| | |
|------------|-----------------|
| session_id | 1331246660 |
| timestamp | 1331246351 |
| source_ip | 1331244570 |
| session_id | 1331261196 |
| timestamp | 3/8/2012 2:44PM |
| source_ip | 3/8/2012 2:38PM |
| session_id | 3/8/2012 2:09PM |
| timestamp | 3/8/2012 6:46PM |
| source_ip | 99.155.155.225 |
| session_id | 65.87.165.114 |
| timestamp | 71.10.106.181 |
| source_ip | 1331261196 |
| session_id | 3/8/2012 6:46PM |
| timestamp | 76.102.156.138 |

SELECT * FROM clickstream
WHERE session_id = 1331246351



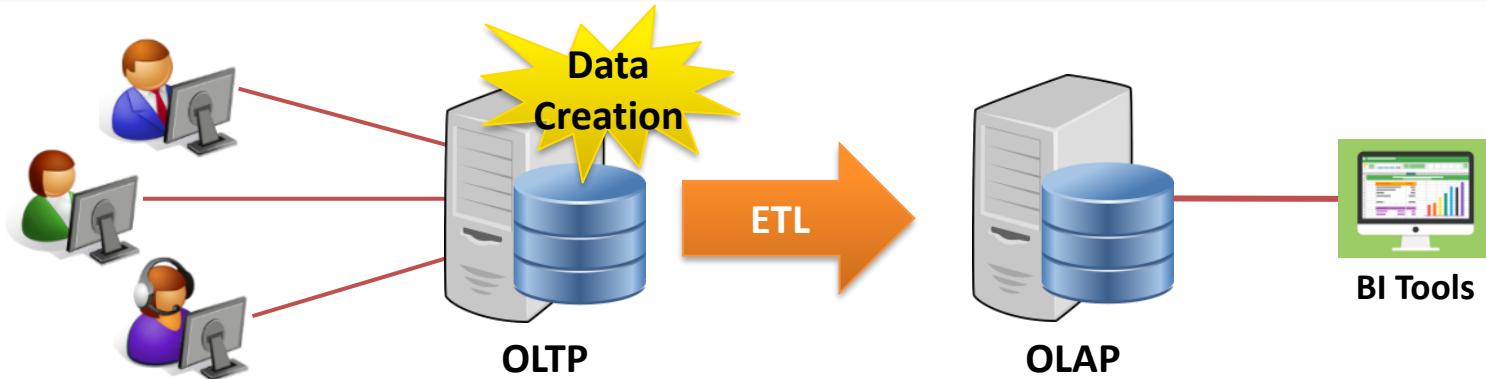
Background: Apache Arrow (2/2)

Most of data types are convertible between Apache Arrow and PostgreSQL.

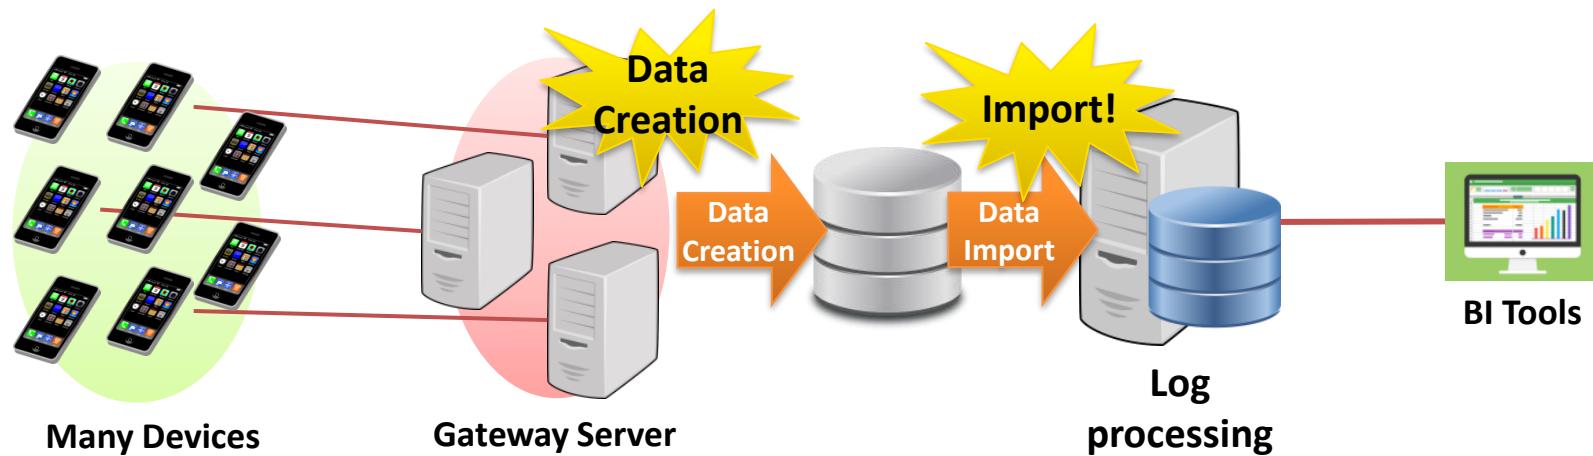
| Apache Arrow Data Types | PostgreSQL Data Types | extra description |
|-------------------------|------------------------|---|
| Int | int2, int4, int8 | |
| FloatingPoint | float2, float4, float8 | float2 is an enhancement by PG-Strom |
| Binary | bytea | |
| Utf8 | text | |
| Bool | bool | |
| Decimal | numeric | Decimal = fixed-point 128bit value |
| Date | date | adjusted to unitsz = Day |
| Time | time | adjusted to unitsz = MicroSecond |
| Timestamp | timestamp | adjusted to unitsz = MicroSecond |
| Interval | interval | |
| List | array types | Only 1-dimensional array is supportable |
| Struct | composite types | |
| Union | ----- | |
| FixedSizeBinary | char(n) | |
| FixedSizeList | ----- | |
| Map | ----- | |

Log-data characteristics (1/2) – WHERE is it generated on?

Traditional OLTP&OLAP – Data is generated **inside** of database system

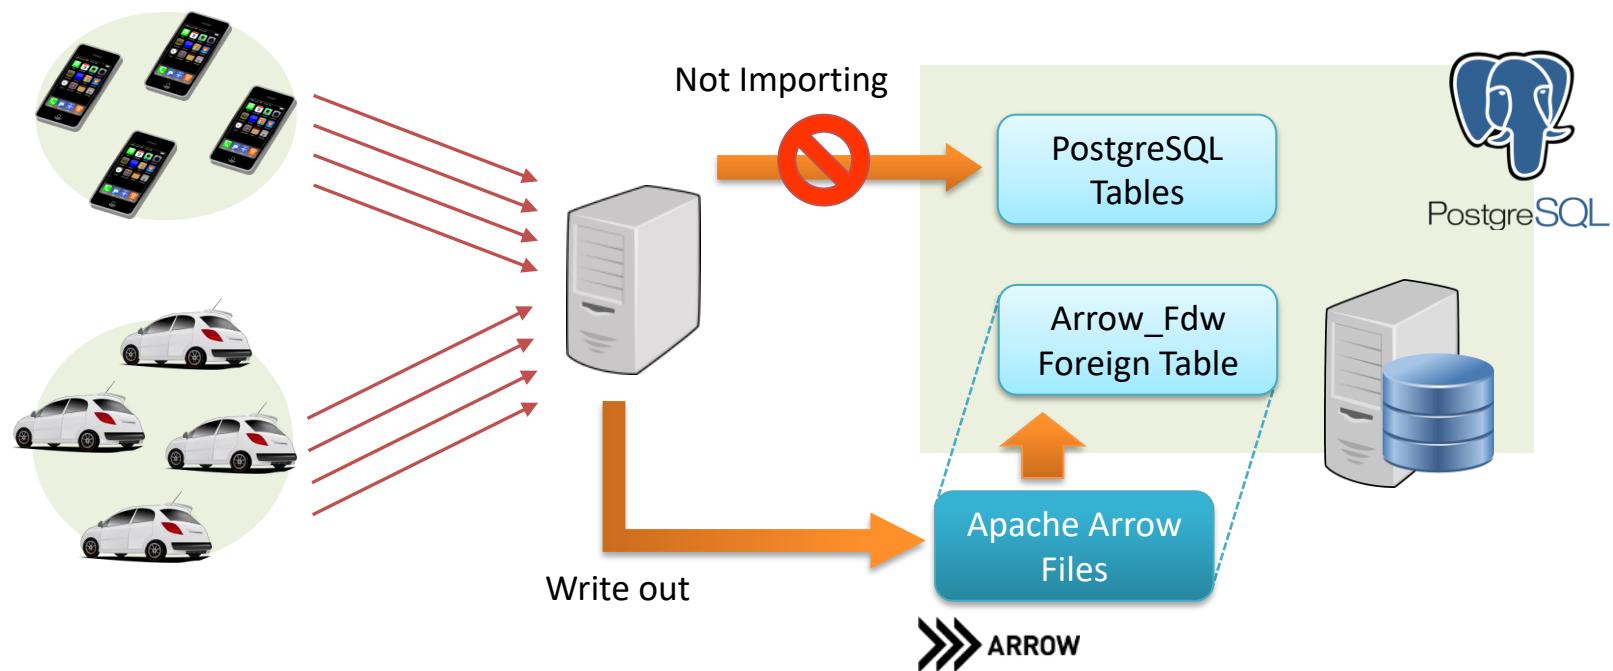


IoT/M2M use case – Data is generated **outside** of database system



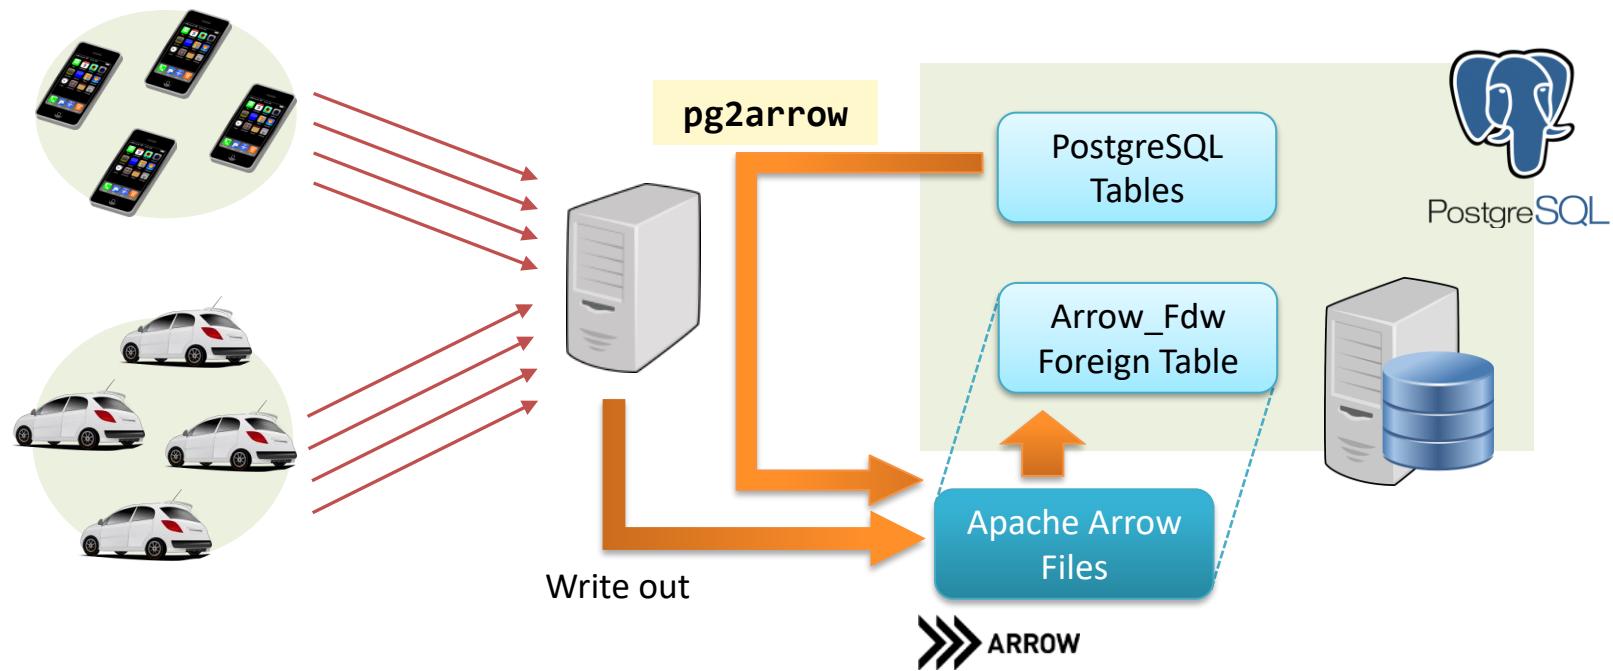
Data Importing becomes a heavy time-consuming operations for big-data processing.

Arrow_Fdw - maps Apache Arrow files as a foreign table



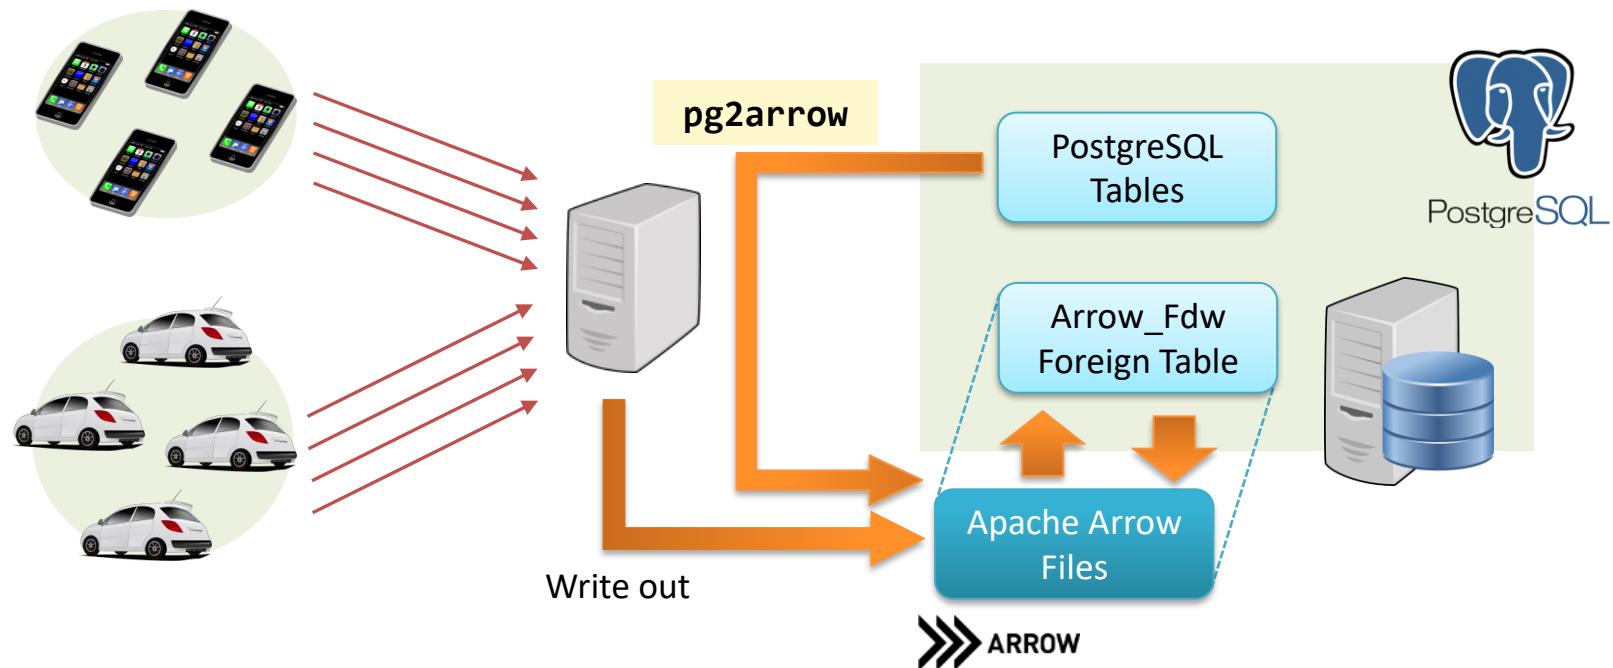
- Arrow_Fdw enables to map Apache Arrow files as a PostgreSQL foreign table.
- Not Importing, so Apache Arrow files are immediately visible.

Arrow_Fdw - maps Apache Arrow files as a foreign table



- Arrow_Fdw enables to map Apache Arrow files as a PostgreSQL foreign table.
- Not Importing, so Apache Arrow files are immediately visible.
- **pg2arrow** generates Apache Arrow file from query results of PostgreSQL.

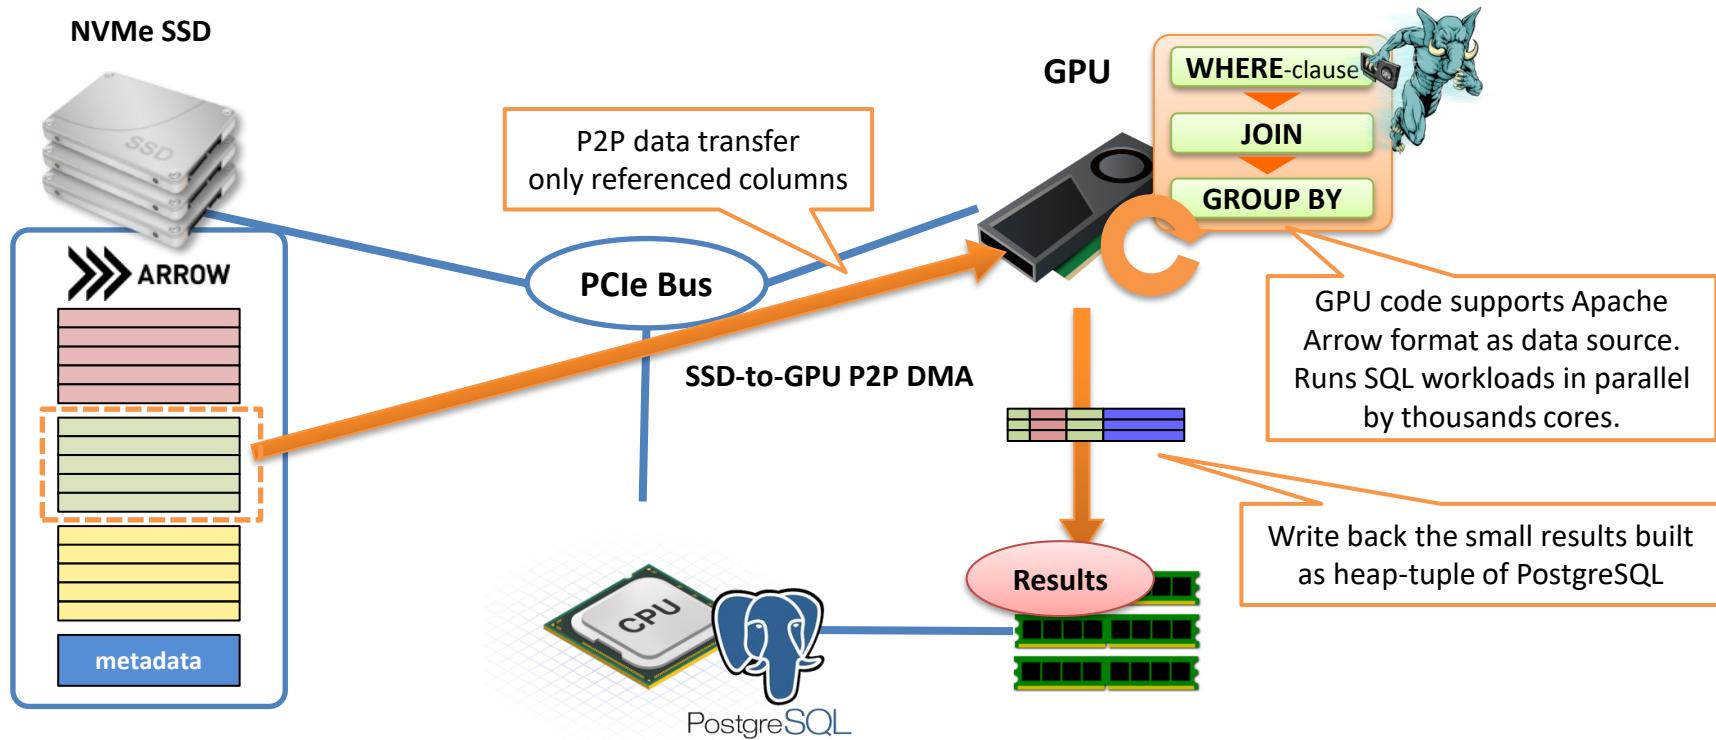
Arrow_Fdw - maps Apache Arrow files as a foreign table



- Arrow_Fdw enables to map Apache Arrow files as a PostgreSQL foreign table.
- Not Importing, so Apache Arrow files are immediately visible.
- pg2arrow generates Apache Arrow file from query results of PostgreSQL.
- Arrow_Fdw can be writable, but only batched-INSERT is supported.

SSD-to-GPU Direct SQL on Arrow_Fdw (1/3)

It transfers **ONLY Referenced Columns** over SSD-to-GPU Direct SQL mechanism

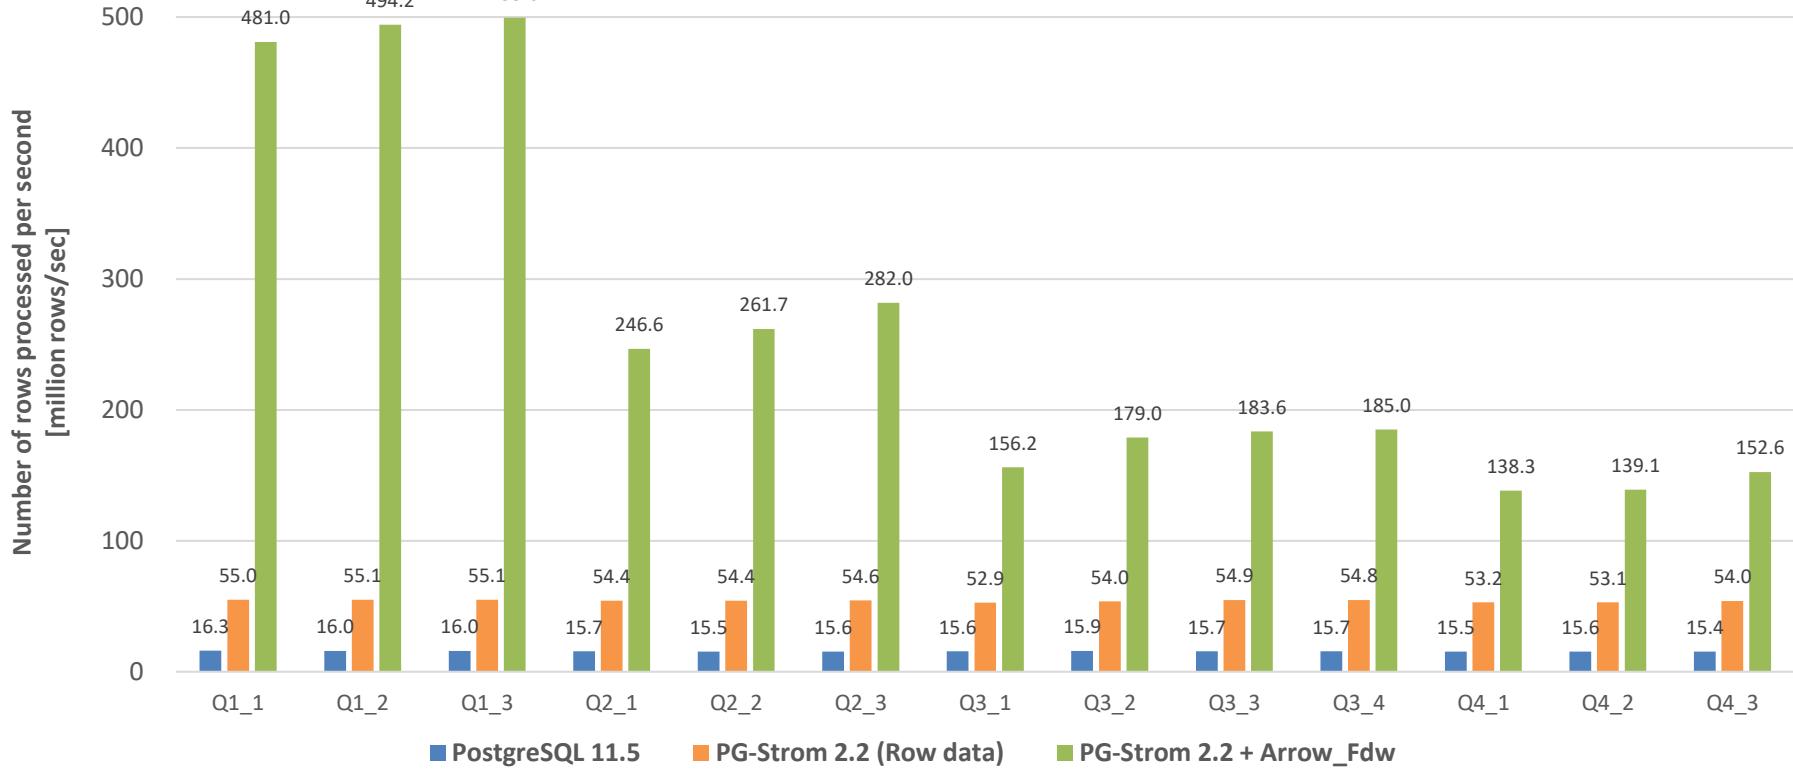


Why Apache Arrow is beneficial?

- ❑ Less amount of I/O to be loaded; only referenced columns
- ❑ Higher utilization of GPU core; by vector processing and wide memory bus
- ❑ Read-only structure; No MVCC checks are required on run-time

SSD-to-GPU Direct SQL on Arrow_Fdw (2/3) - Benchmark Results

Star Schema Benchmark ($s=1000$, DBsize: 879GB) on Tesla V100 x1 + DC P4600 x3



- Query Execution Throughput = $(6.0B \text{ rows}) / (\text{Query Response Time [sec]})$
- Combined usage of SSD-to-GPU Direct SQL and Columnar-store pulled out the extreme performance; 130-500M rows per second
- Server configuration is identical to what we show on p.17. (1U, 1CPU, 1GPU, 3SSD)

SSD-to-GPU Direct SQL on Arrow_Fdw (3/3) - Benchmark Validation

Almost equivalent physical data transfer, but no need to load unreferenced columns

Foreign table "public.flineorder"

| Column | Type | Size | |
|--------------------|---------------|---------|-------------------------|
| lo_orderkey | numeric | 89.42GB | |
| lo_linenumber | integer | 22.37GB | |
| lo_custkey | numeric | 89.42GB | |
| lo_partkey | integer | 22.37GB | <-- ★Referenced by Q2_1 |
| lo_suppkey | numeric | 89.42GB | <-- ★Referenced by Q2_1 |
| lo_orderdate | integer | 22.37GB | <-- ★Referenced by Q2_1 |
| lo_orderpriority | character(15) | 83.82GB | |
| lo_shipppriority | character(1) | 5.7GB | |
| lo_quantity | integer | 22.37GB | |
| lo_extendedprice | integer | 22.37GB | |
| lo_ordertotalprice | integer | 22.37GB | |
| lo_discount | integer | 22.37GB | |
| lo_revenue | integer | 22.37GB | <-- ★Referenced by Q2_1 |
| lo_supplycost | integer | 22.37GB | |
| lo_tax | integer | 22.37GB | |
| lo_commit_date | character(8) | 44.71GB | |
| lo_shipmode | character(10) | 55.88GB | |

FDW options: (file '/opt/nvme/lineorder_s401.arrow') ... file size = 310GB

Q2_1 actuall reads only 157GB of 681GB (23.0%) from the NVME-SSD.

Execution time of Q2_1 is 24.3s, so 157GB / 24.3s = 6.46GB/s

→ Reasonable performance for 3x Intel DC P4600 on single CPU configuration.

Table Partitioning

Log-data characteristics (2/2) – INSERT-only

Transactional Data

**INSERT
UPDATE
DELETE**



Log Data

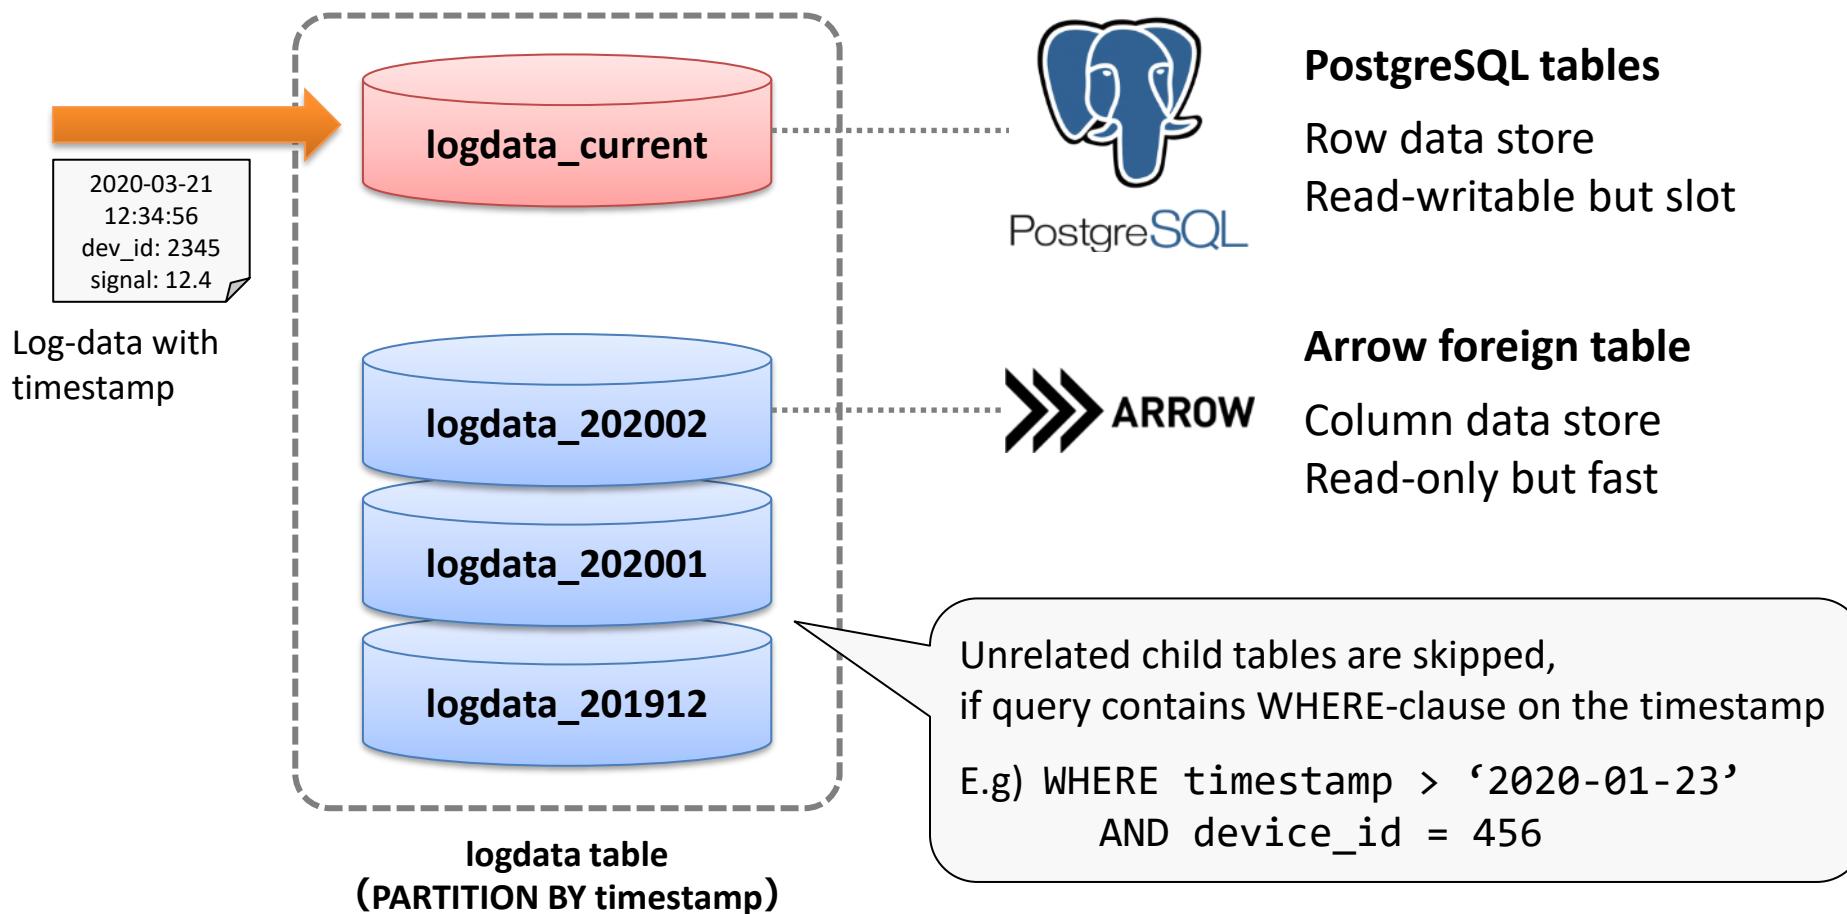
**INSERT
UPDATE
DELETE**



- ✓ MVCC visibility check is (relatively) not significant.
- ✓ Rows with old timestamp will never inserted.

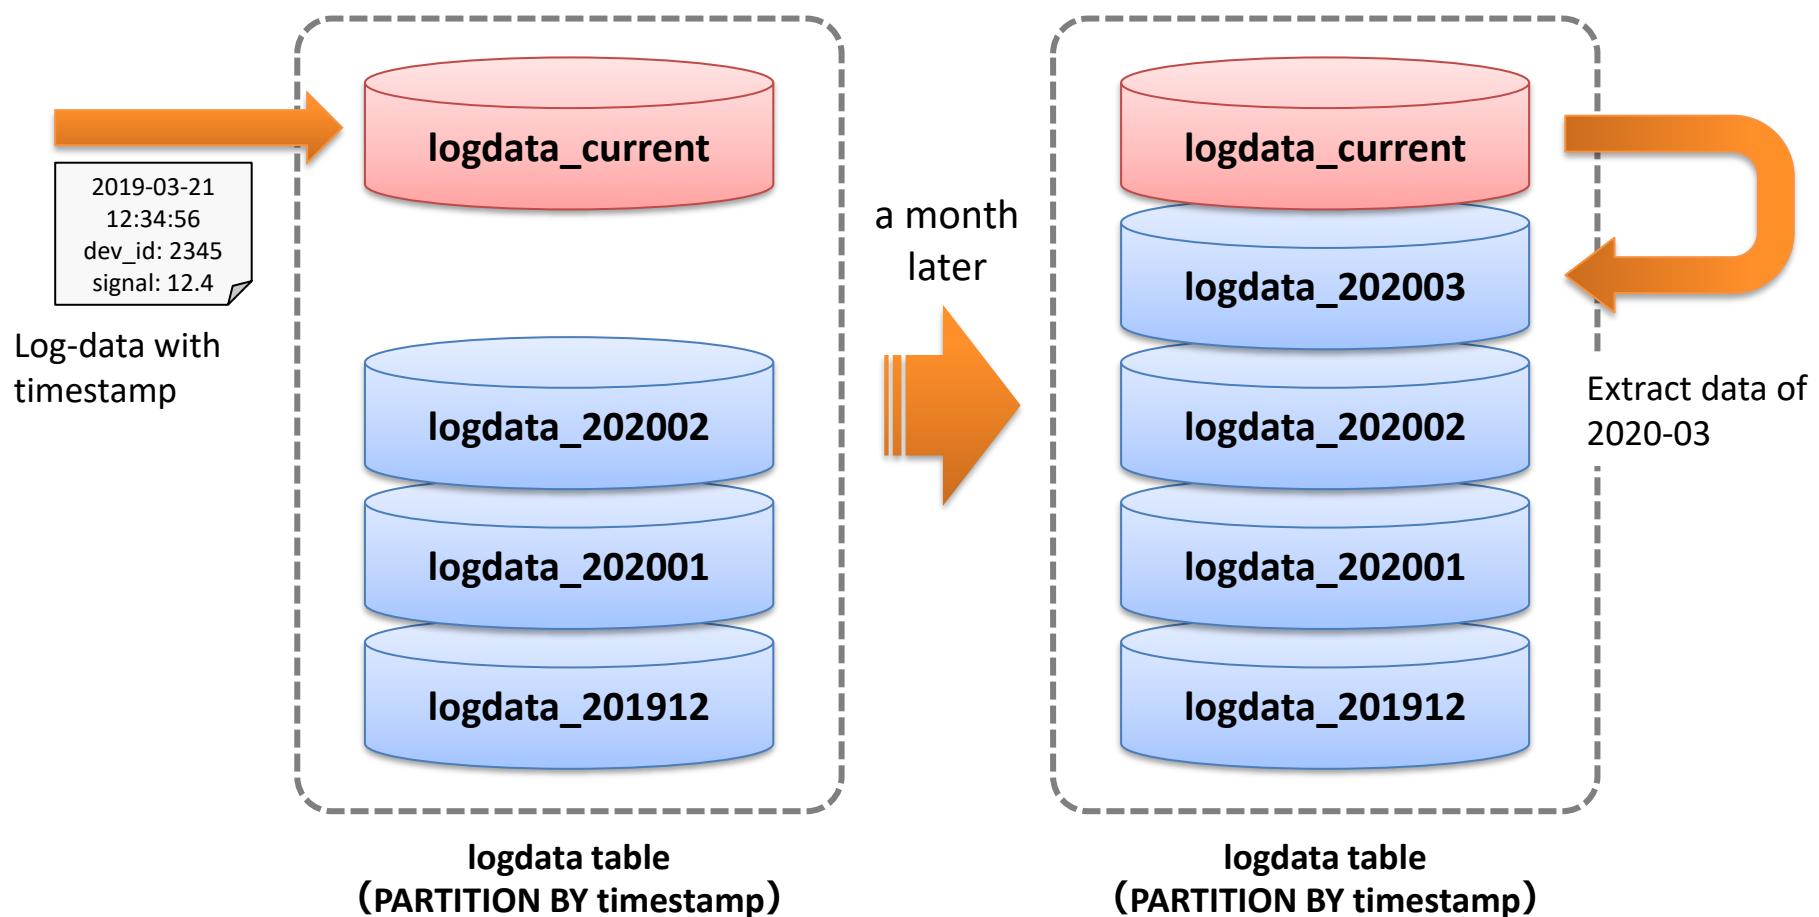
Configuration of PostgreSQL partition for log-data (1/2)

- Mixture of PostgreSQL table and Arrow foreign table in partition declaration
- Log-data should have timestamp, and never updated
- Old data can be moved to Arrow foreign table for more efficient I/O



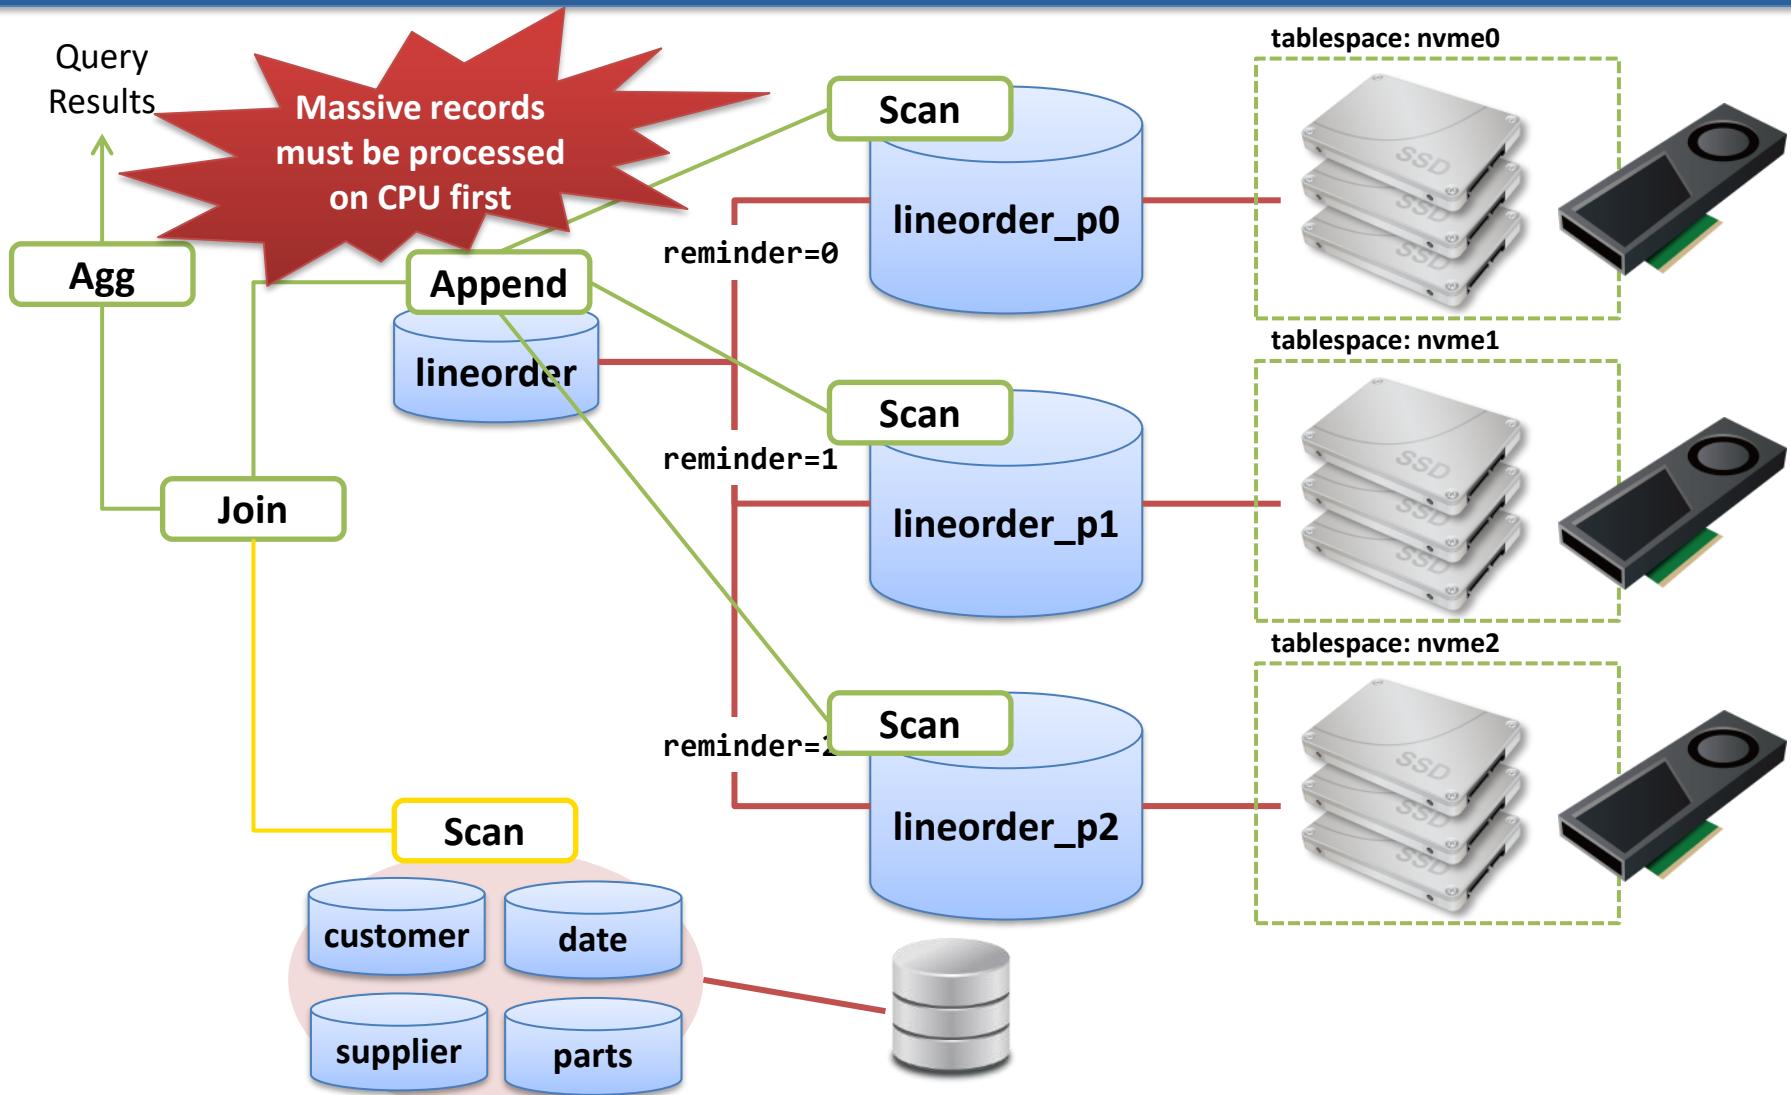
Configuration of PostgreSQL partition for log-data (2/2)

- Mixture of PostgreSQL table and Arrow foreign table in partition declaration
- Log-data should have timestamp, and never updated
- Old data can be moved to Arrow foreign table for more efficient I/O



Asymmetric Partition-wise JOIN (1/4)

Records from partition-leafs must be backed to CPU and processed once!



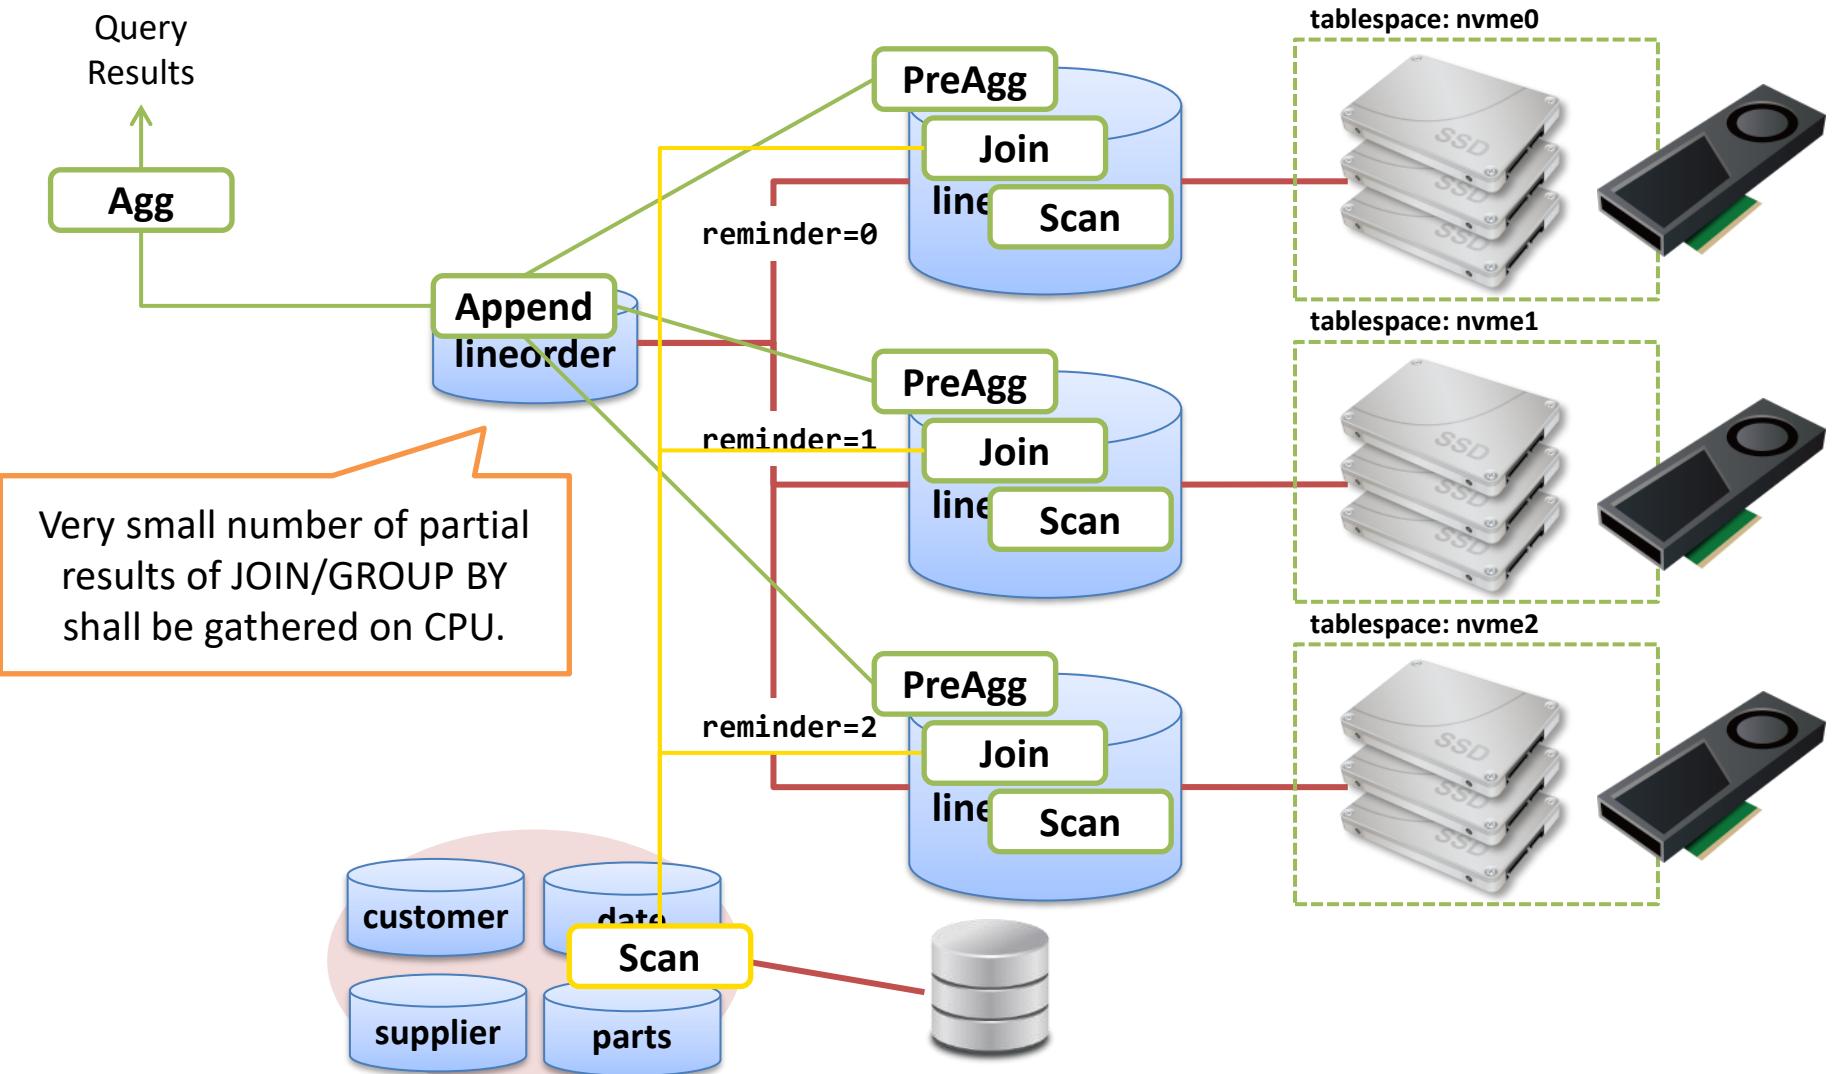
Asymmetric Partition-wise JOIN (2/4)

```
postgres=# explain select * from ptable p, t1 where p.a = t1.aid;  
                      QUERY PLAN
```

```
-----  
Hash Join  (cost=2.12..24658.62 rows=49950 width=49)  
  Hash Cond: (p.a = t1.aid)  
    -> Append  (cost=0.00..20407.00 rows=1000000 width=12)  
      -> Seq Scan on ptable_p0 p  (cost=0.00..5134.63 rows=333263 width=12)  
      -> Seq Scan on ptable_p1 p_1  (cost=0.00..5137.97 rows=333497 width=12)  
      -> Seq Scan on ptable_p2 p_2  (cost=0.00..5134.40 rows=333240 width=12)  
    -> Hash  (cost=1.50..1.50 rows=50 width=37)  
      -> Seq Scan on t1  (cost=0.00..1.50 rows=50 width=37)  
(8 rows)
```

Asymmetric Partition-wise JOIN (3/4)

Push down JOIN/GROUP BY, even if smaller half is not a partitioned table.



Asymmetric Partition-wise JOIN (4/4)

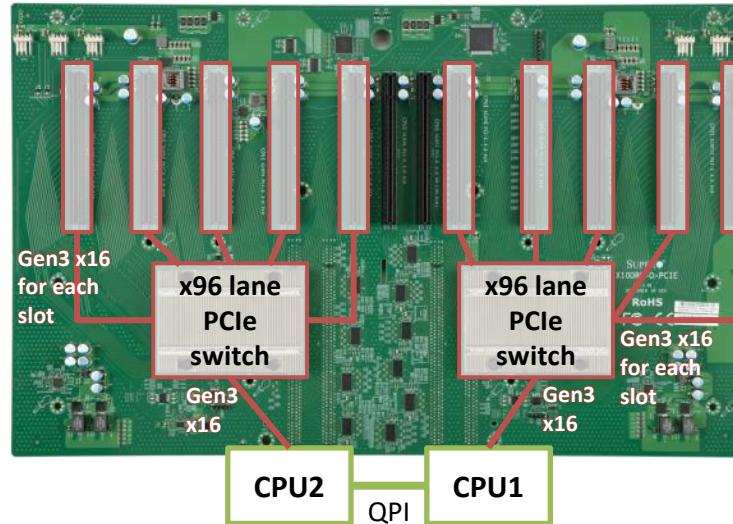
```
postgres=# set enable_partitionwise_join = on;
SET
postgres=# explain select * from ptable p, t1 where p.a = t1.aid;
                                         QUERY PLAN
-----
Append  (cost=2.12..19912.62 rows=49950 width=49)
 -> Hash Join  (cost=2.12..6552.96 rows=16647 width=49)
      Hash Cond: (p.a = t1.aid)
      -> Seq Scan on ptable_p0 p  (cost=0.00..5134.63 rows=333263 width=12)
      -> Hash  (cost=1.50..1.50 rows=50 width=37)
            -> Seq Scan on t1  (cost=0.00..1.50 rows=50 width=37)
 -> Hash Join  (cost=2.12..6557.29 rows=16658 width=49)
      Hash Cond: (p_1.a = t1.aid)
      -> Seq Scan on ptable_p1 p_1  (cost=0.00..5137.97 rows=333497 width=12)
      -> Hash  (cost=1.50..1.50 rows=50 width=37)
            -> Seq Scan on t1  (cost=0.00..1.50 rows=50 width=37)
 -> Hash Join  (cost=2.12..6552.62 rows=16645 width=49)
      Hash Cond: (p_2.a = t1.aid)
      -> Seq Scan on ptable_p2 p_2  (cost=0.00..5134.40 rows=333240 width=12)
      -> Hash  (cost=1.50..1.50 rows=50 width=37)
            -> Seq Scan on t1  (cost=0.00..1.50 rows=50 width=37)
(16 rows)
```

PCIe-bus level optimization (1/3)

HPC Server – optimization for GPUDirect RDMA



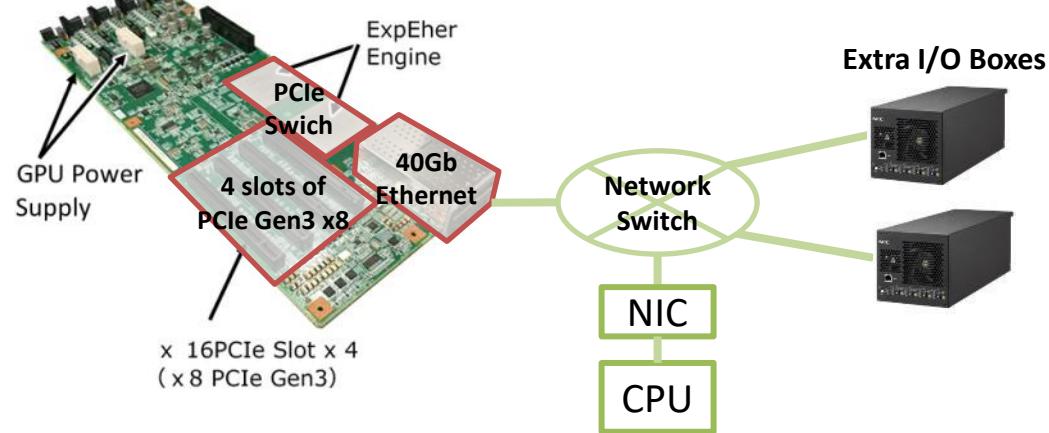
Supermicro
SYS-4029TRT2



I/O Expansion Box

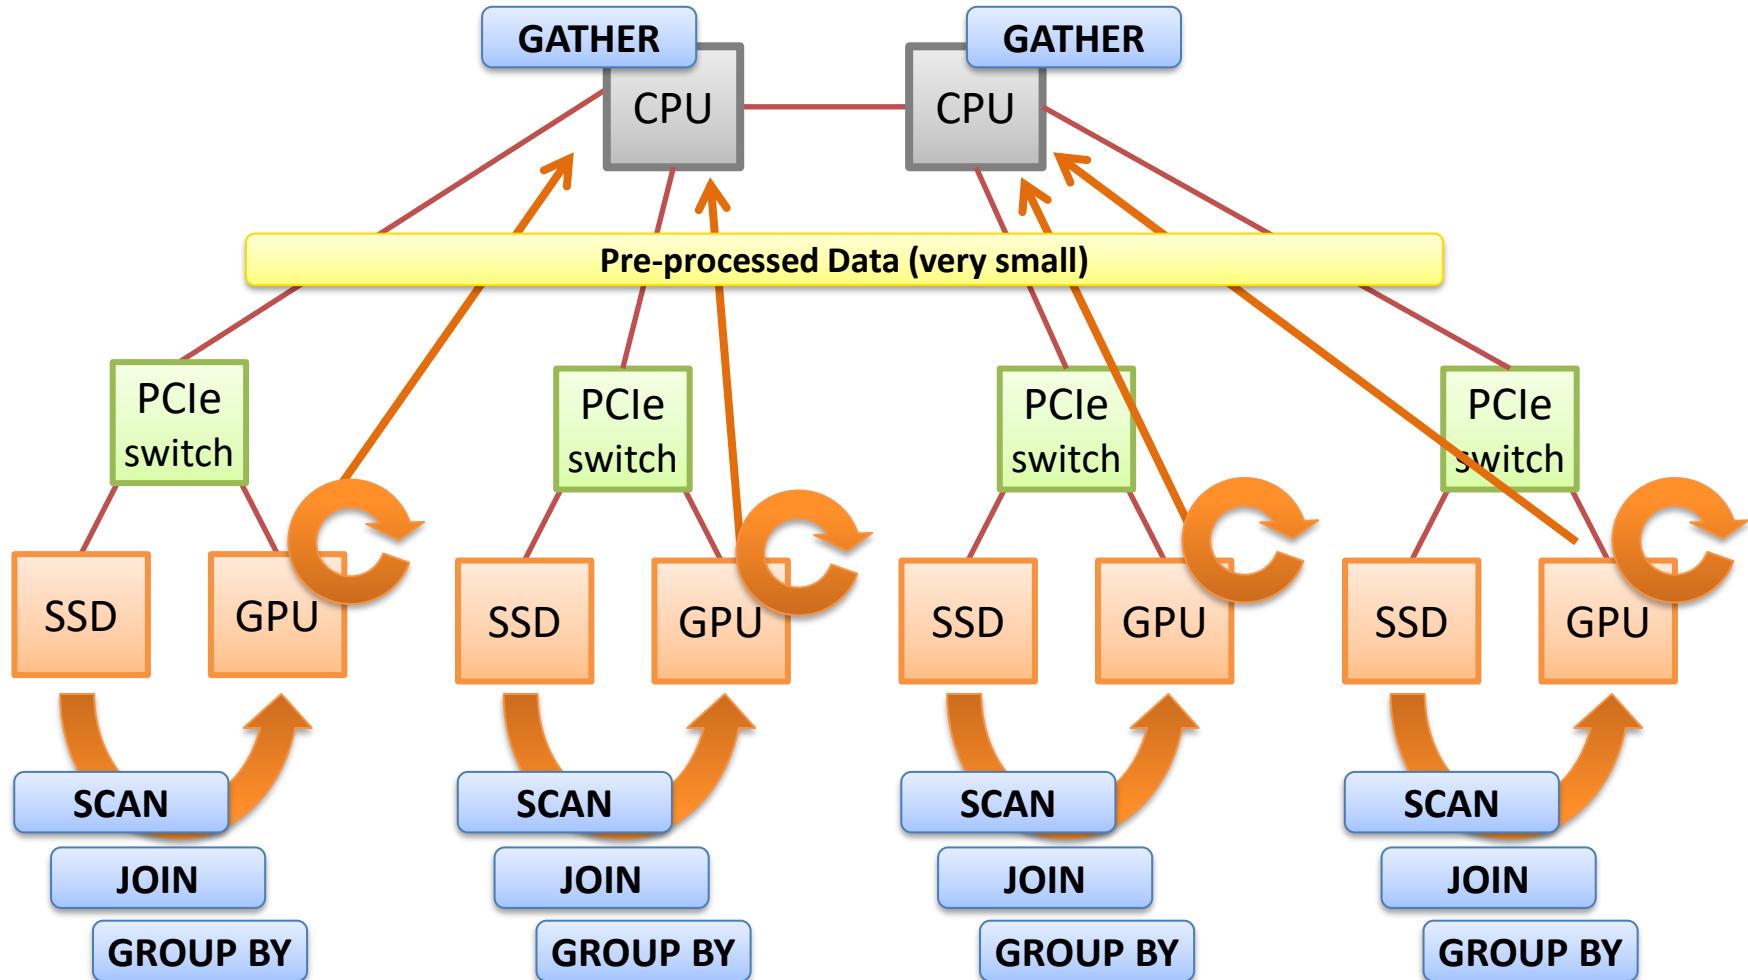


NEC ExpEther 40G
(4slots edition)



PCIe-bus level optimization (2/3)

By P2P DMA over PCIe-switch, major data traffic bypass CPU



PCIe-bus level optimization (3/3)

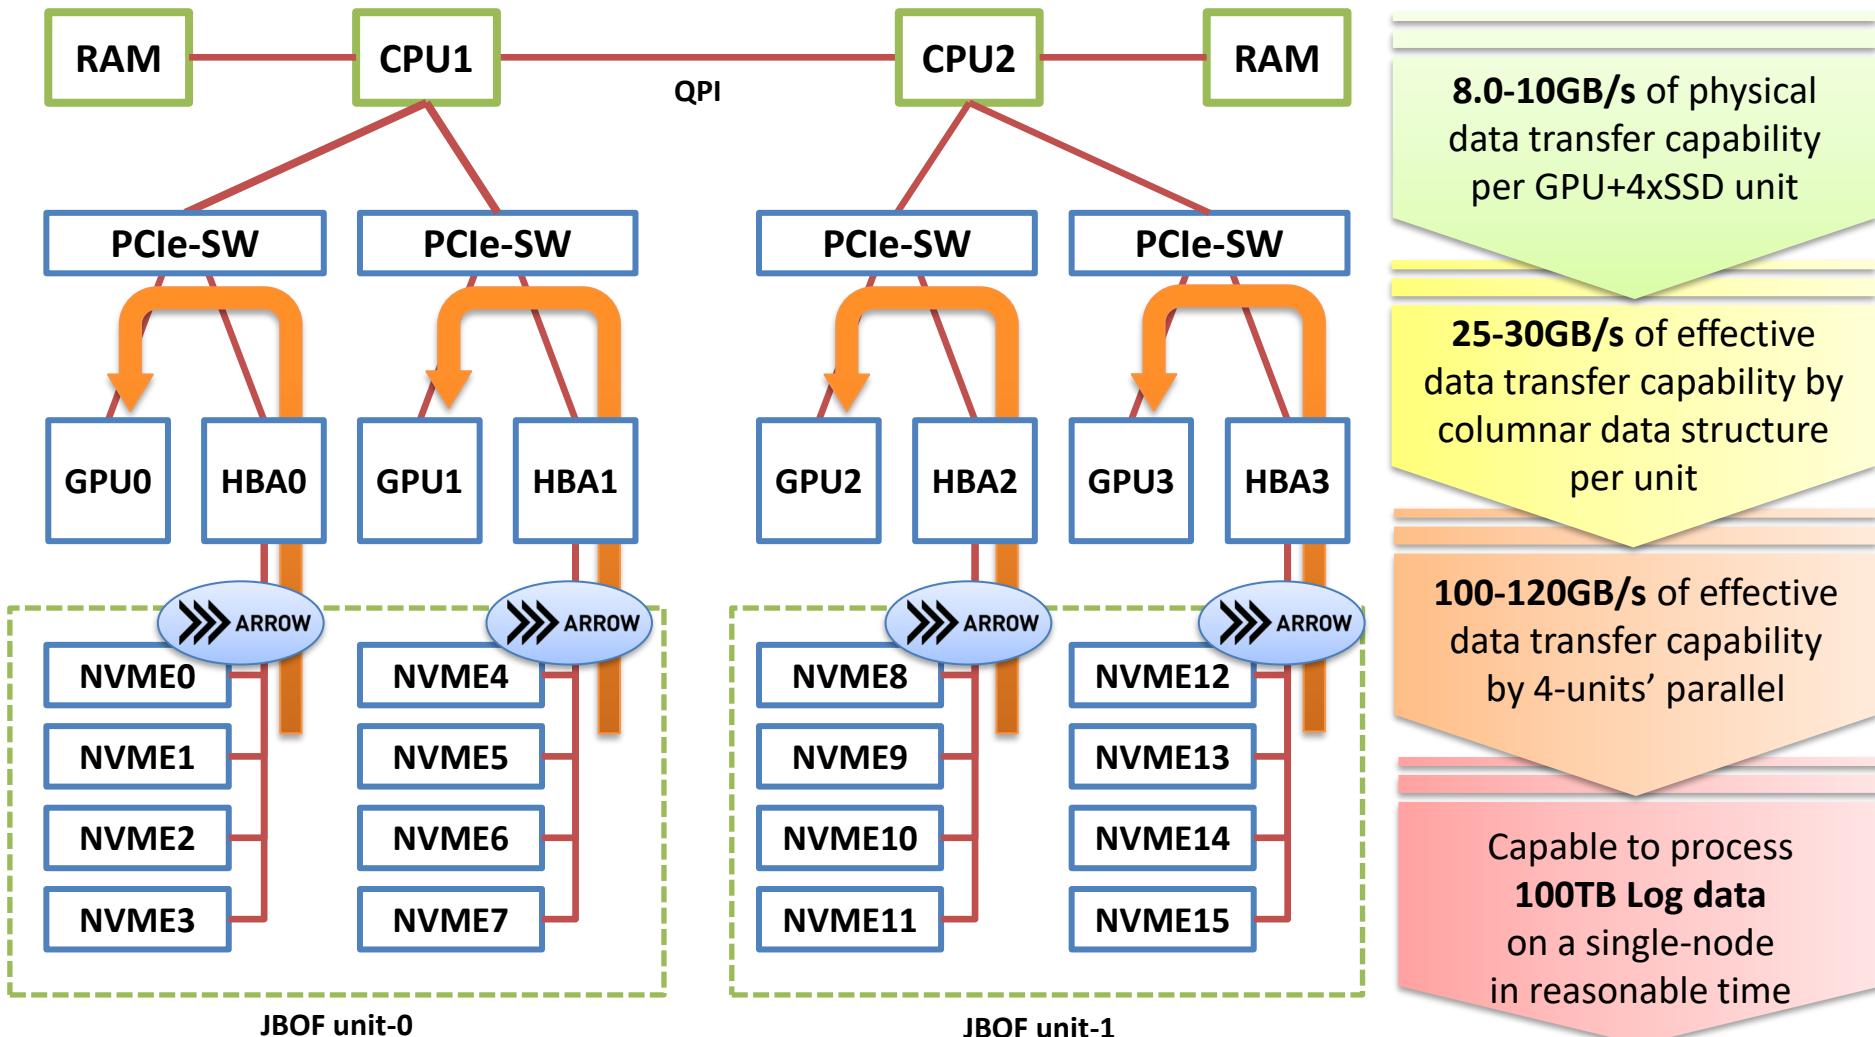
```
$ pg_ctl restart
:
LOG:    -  PCIe[0000:80]
LOG:        -  PCIe(0000:80:02.0)
LOG:            -  PCIe(0000:83:00.0)
LOG:                -  PCIe(0000:84:00.0)
LOG:                    -  PCIe(0000:85:00.0) nvme0 (INTEL SSDPEDKE020T7)
LOG:                -  PCIe(0000:84:01.0)
LOG:                    -  PCIe(0000:86:00.0) GPU0 (Tesla P40)
LOG:                -  PCIe(0000:84:02.0)
LOG:                    -  PCIe(0000:87:00.0) nvme1 (INTEL SSDPEDKE020T7)
LOG:        -  PCIe(0000:80:03.0)
LOG:            -  PCIe(0000:c0:00.0)
LOG:                -  PCIe(0000:c1:00.0)
LOG:                    -  PCIe(0000:c2:00.0) nvme2 (INTEL SSDPEDKE020T7)
LOG:                -  PCIe(0000:c1:01.0)
LOG:                    -  PCIe(0000:c3:00.0) GPU1 (Tesla P40)
LOG:                -  PCIe(0000:c1:02.0)
LOG:                    -  PCIe(0000:c4:00.0) nvme3 (INTEL SSDPEDKE020T7)
LOG:        -  PCIe(0000:80:03.2)
LOG:            -  PCIe(0000:e0:00.0)
LOG:                -  PCIe(0000:e1:00.0)
LOG:                    -  PCIe(0000:e2:00.0) nvme4 (INTEL SSDPEDKE020T7)
LOG:                -  PCIe(0000:e1:01.0)
LOG:                    -  PCIe(0000:e3:00.0) GPU2 (Tesla P40)
LOG:                -  PCIe(0000:e1:02.0)
LOG:                    -  PCIe(0000:e4:00.0) nvme5 (INTEL SSDPEDKE020T7)
LOG: GPU<->SSD Distance Matrix
LOG:      GPU0      GPU1      GPU2
LOG:  nvme0  (  3)      7      7
LOG:  nvme5      7      7  (  3)
LOG:  nvme4      7      7  (  3)
LOG:  nvme2      7  (  3)      7
LOG:  nvme1  (  3)      7      7
LOG:  nvme3      7  (  3)      7
```

PG-Strom chooses the closest GPU
for the tables to be scanned,
according to the PCI-E device distance.

Large-scale Benchmark

Large-scale Benchmark (1/4)

All of SSD-to-GPU Direct SQL, Apache Arrow and PCI-E bus optimization are used



Large-scale Benchmark (2/4)

Build HPC 4U Server + 4 of GPU + 16 of NVME-SSD configuration



Intel SSD
DC P4510 (1.0TB) x16



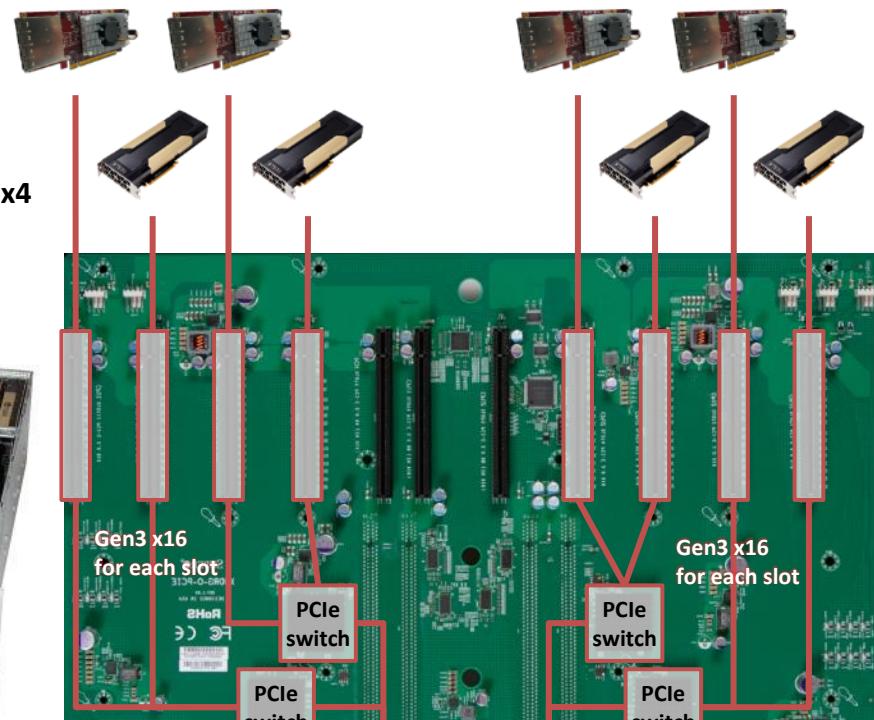
SerialCables
dual PCI-ENC8G-08A
(U.2 NVME JBOF; 8slots) x2



Supermicro
SYS-4029GP-TRT

via SFF-8644 based PCIe x4 cables
(3.2GB/s x 16 = max 51.2GB/s)

SerialCables
PCI-AD-x16HE x4



SPECIAL THANKS



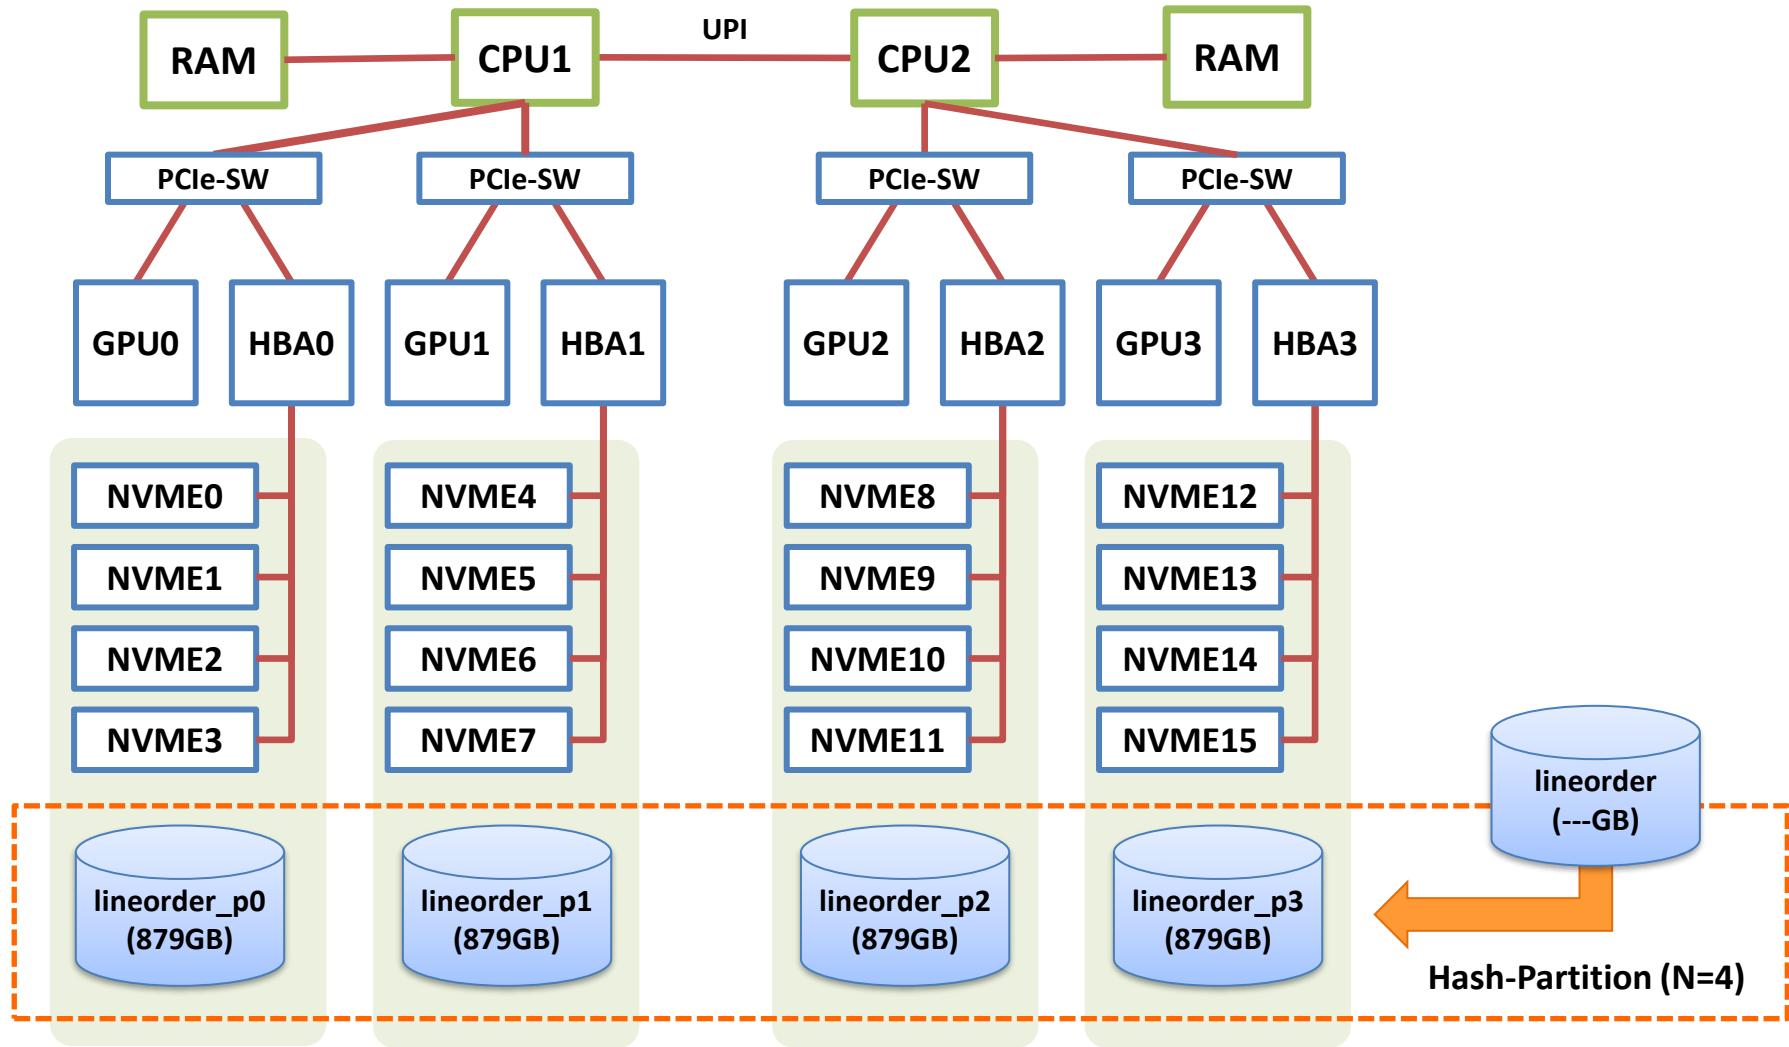
Large-scale Benchmark (3/4)

Build HPC 4U Server + 4 of GPU + 16 of NVME-SSD configuration



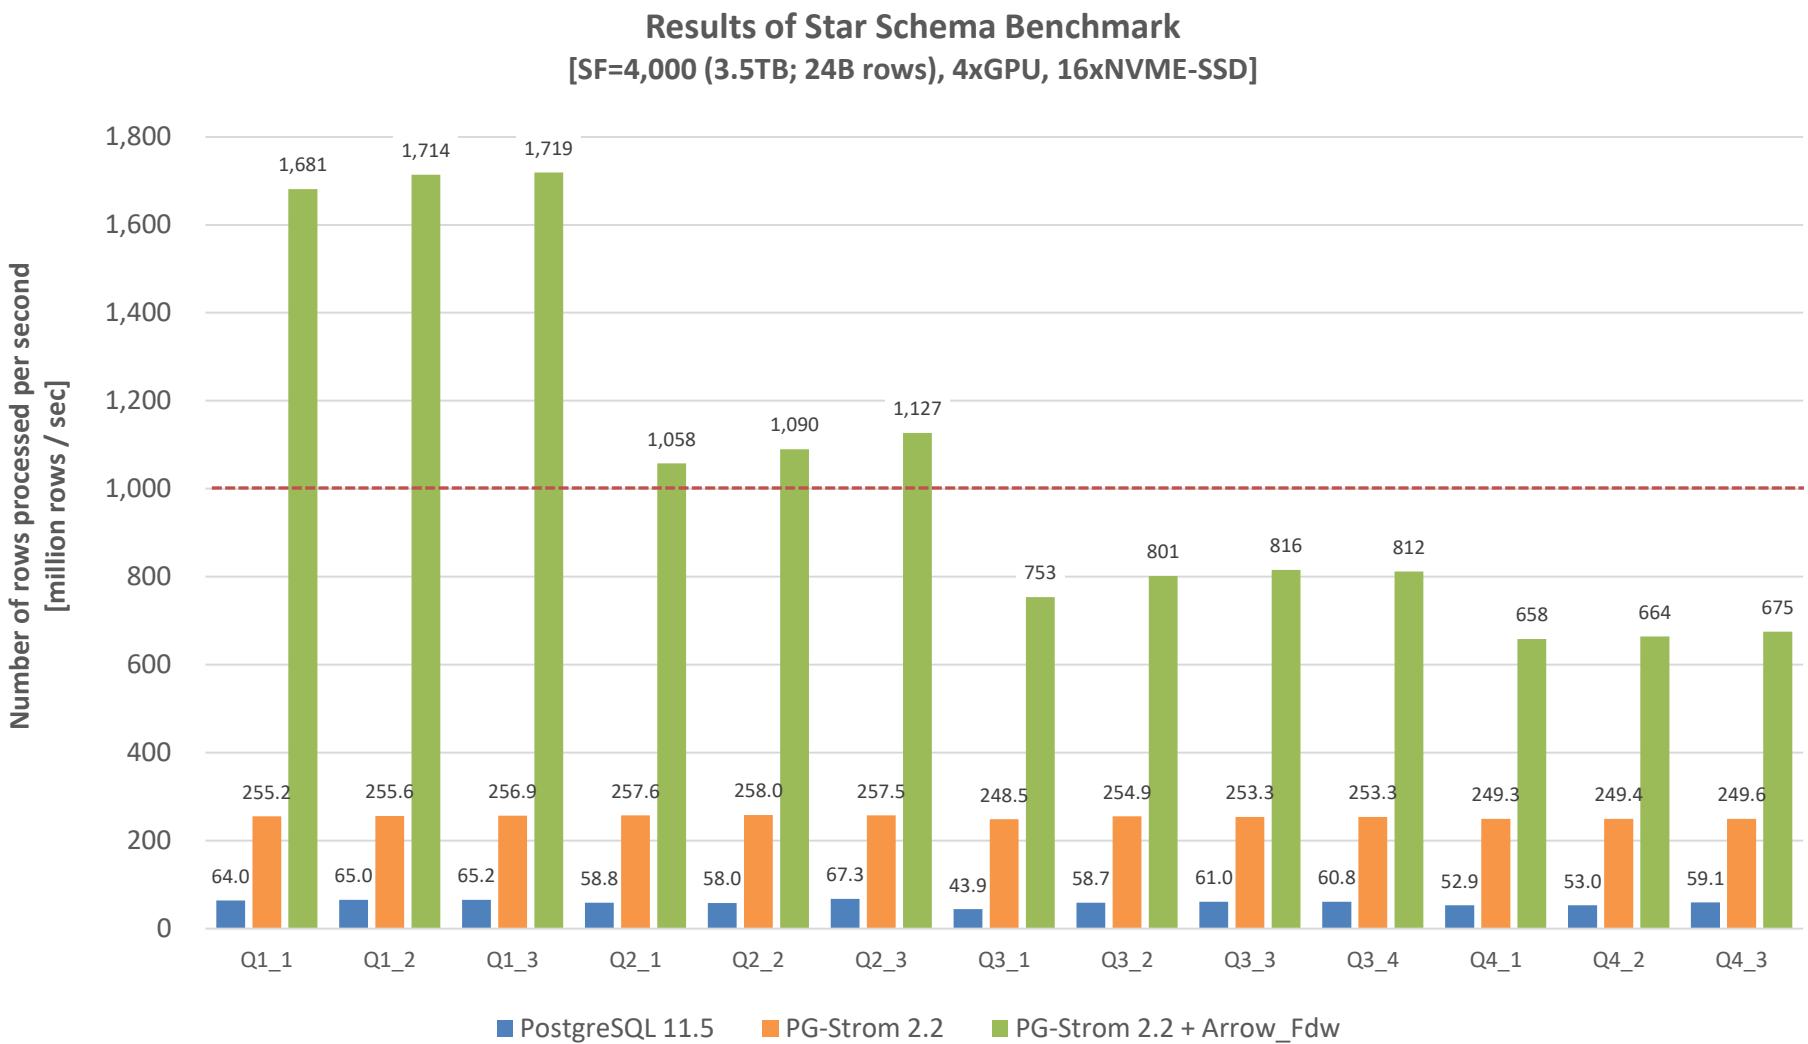
Large-scale Benchmark (4/4)

Distributed 3.5TB test data into 4 partitions; 879GB for each



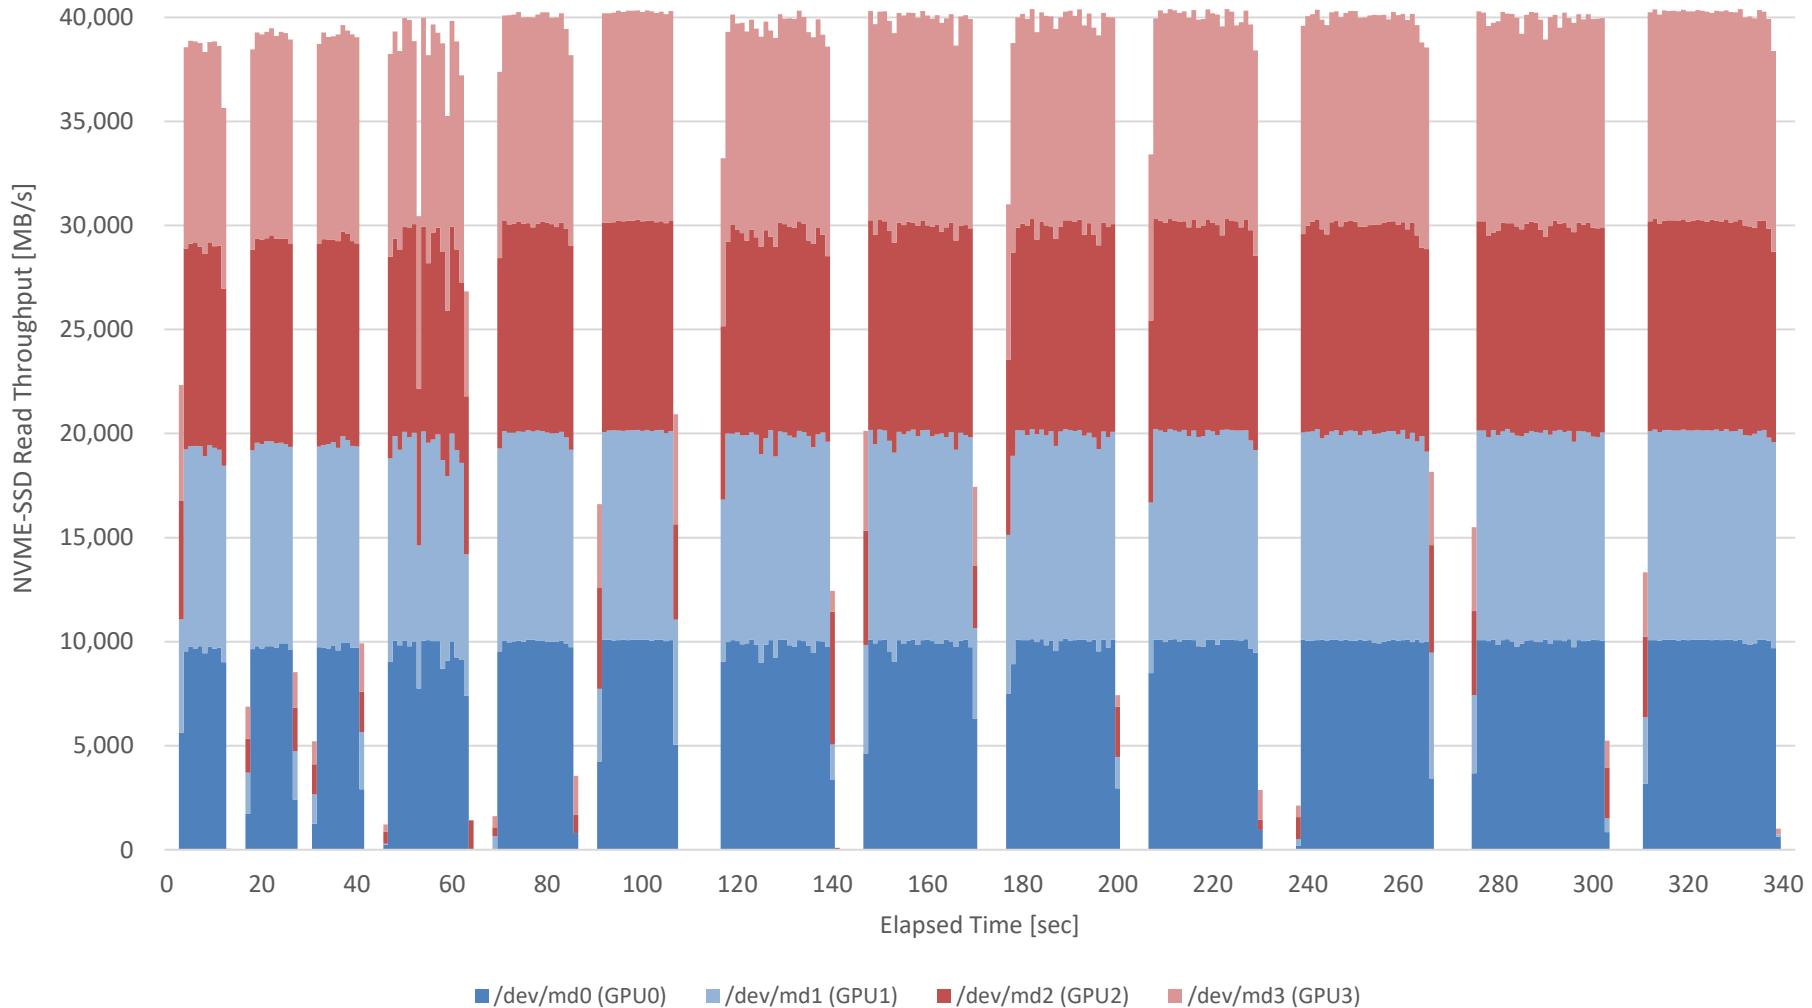
Benchmark Results (1/2)

Billion rows per second at a single-node PostgreSQL!!



Benchmark Results (2/2)

40GB/s Read from NVME-SSDs during SQL execution (95% of H/W limit)



Future Direction

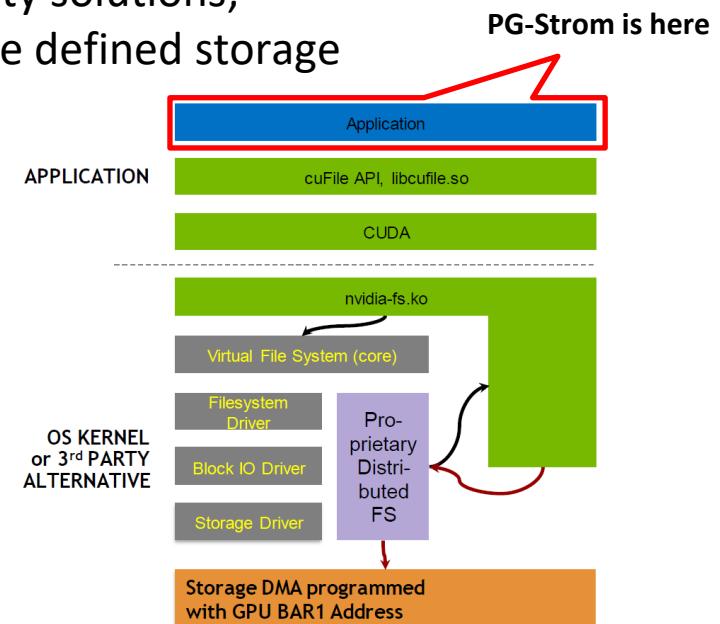
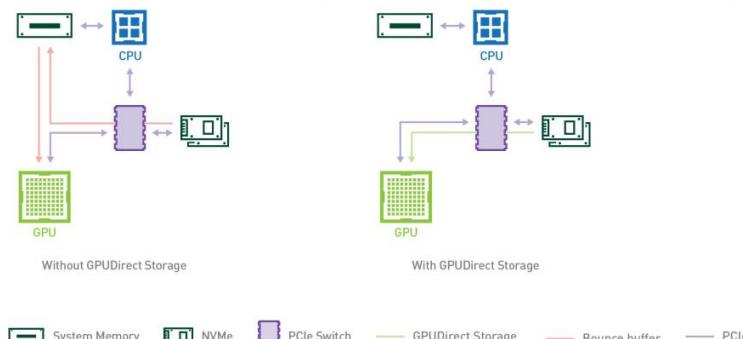
NVIDIA GPUDirect Storage (cuFile API) support

NVIDIA GPUDirect Storage?

- ❑ API & Driver stack for direct read from NVME to GPU
 - ✓ Features are almost equivalent to NVME-Strom
 - ✓ Ubuntu 18.04 & RHEL8/CentOS8 shall be released at the initial release
- ❑ NVIDIA will officially release the software in 2020.

Why beneficial?

- ❑ Linux kernel driver that is tested/evaluated by NVIDIA's QA process.
- ❑ Involvement of broader ecosystem with third-party solutions;
 - like distributed-filesystem, block storage, software defined storage
- ❑ Broader OS support; Ubuntu 18.04.



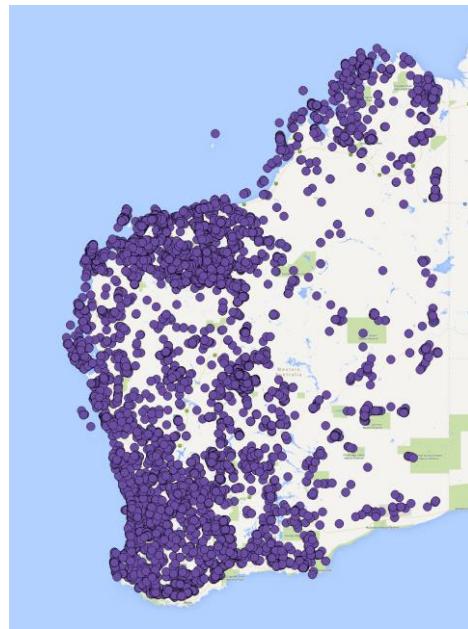
Ref: **GPUDIRECT STORAGE: A DIRECT GPU-STORAGE DATA PATH**

<https://on-demand.gputechconf.com/supercomputing/2019/pdf/sc1922-gpudirect-storage-transfer-data-directly-to-gpu-memory-alleviating-io-bottlenecks.pdf>

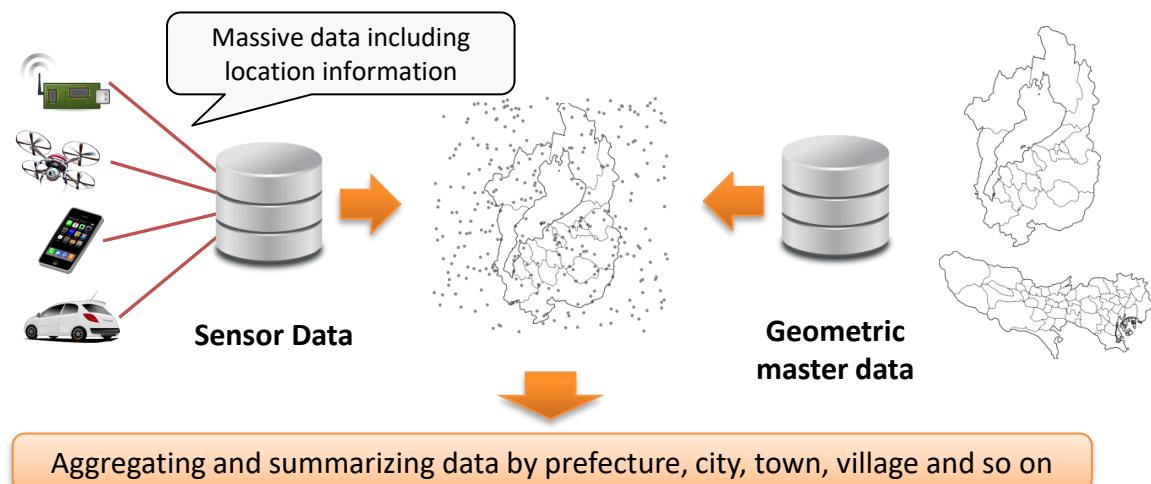
PostGIS Functions & Operator support (1/2)

PostGIS

- An extension of handle geographic objects and relational operations on them.
 - ✓ Geographic objects: Point, LineString, Polygon, and etc
 - ✓ Geographic relational operations: contains, crosses, touch, distance, and etc...
- Wise design for fast search
 - ✓ Bounding-box, GiST-index, CPU Parallel
- v1.0 is released at 2005, then its development has continued over 15 years.



【IoT/M2Mデータに対しての想定利用シーン】



© GAIA RESOURCES

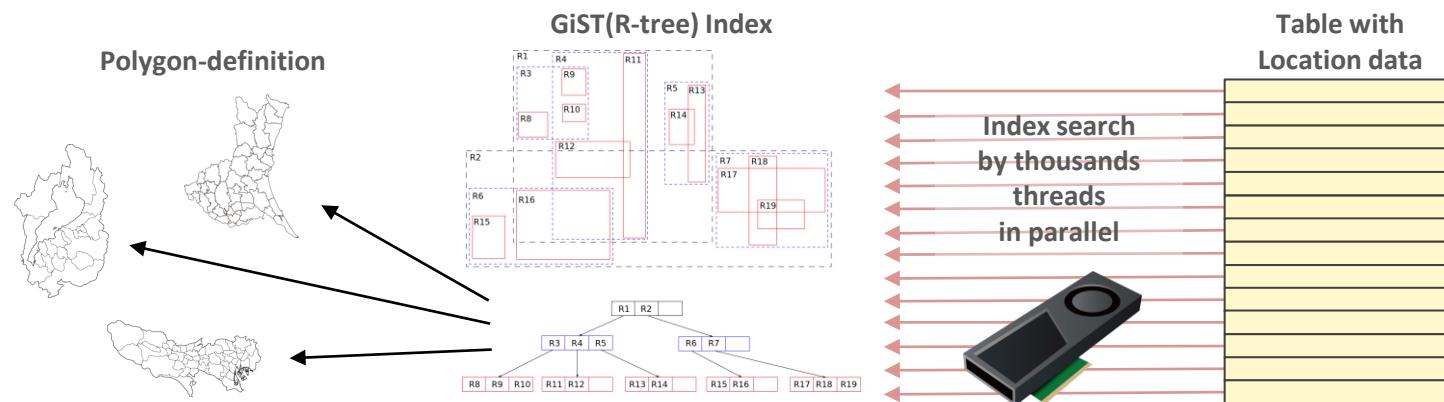
PostGIS Functions & Operator support (2/2)

Current status

- ❑ `geometry st_makewkt(float8, float8, ...)`
- ❑ `float8 st_distance(geometry, geometry)`
- ❑ `bool st_dwithin(geometry, geometry, float8)`
- ❑ `bool st_contains(geometry, geometry)`
- ❑ `bool st_crosses(geometry, geometry)`

Next: GiST (R-tree) Index support

- ❑ Index-based nested-loop up to million polygon x billion points.
- ❑ Runs index-search by GPU's thousands threads in parallel.
- ➔ Because of GiST internal structure, GPU-parallel is likely efficient.



Resources



Repository

- ❑ <https://github.com/heterodb/pg-strom>
- ❑ <https://heterodb.github.io/swdc/>

Document

- ❑ <http://heterodb.github.io/pg-strom/>

Contact

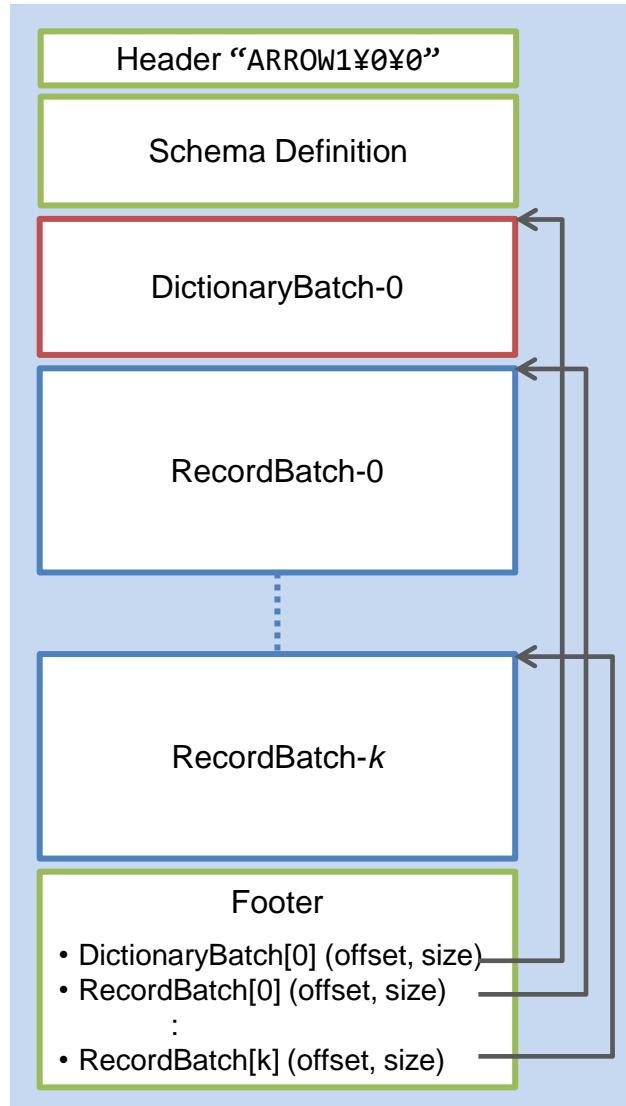
- ❑ kaigai@heterodb.com
- ❑ Tw: [@kkaigai / @heterodb](#)



Backup Slides

Apache Arrow Internal (1/3)

Apache Arrow file



Apache Arrow Internal

□ Header

- Fixed byte string: "ARROW1¥0¥0"

□ Schema Definition

- # of columns and columns definitions.
- Data type/precision, column name, column number, ...

□ Record Batch

- Internal data block that contains a certain # of rows.
- E.g, if N=1M rows with (Int32, Float64), this record-batch begins from 1M of Int32 array, then 1M of Float64 array follows.

□ Dictionary Batch

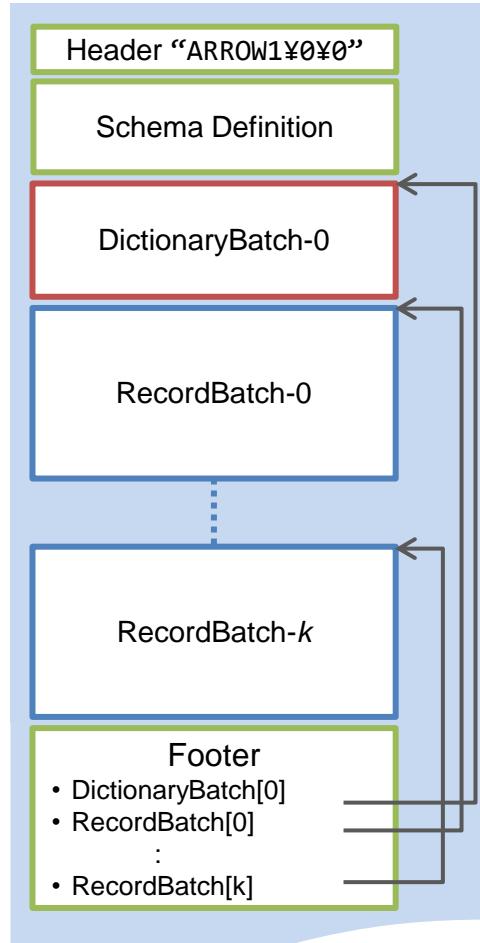
- Optional area for dictionary compression.
- E.g) 1 = "Tokyo", 2 = "Osaka", 3 = "Kyoto"
 - ➔ Int32 representation for frequently appearing words

□ Footer

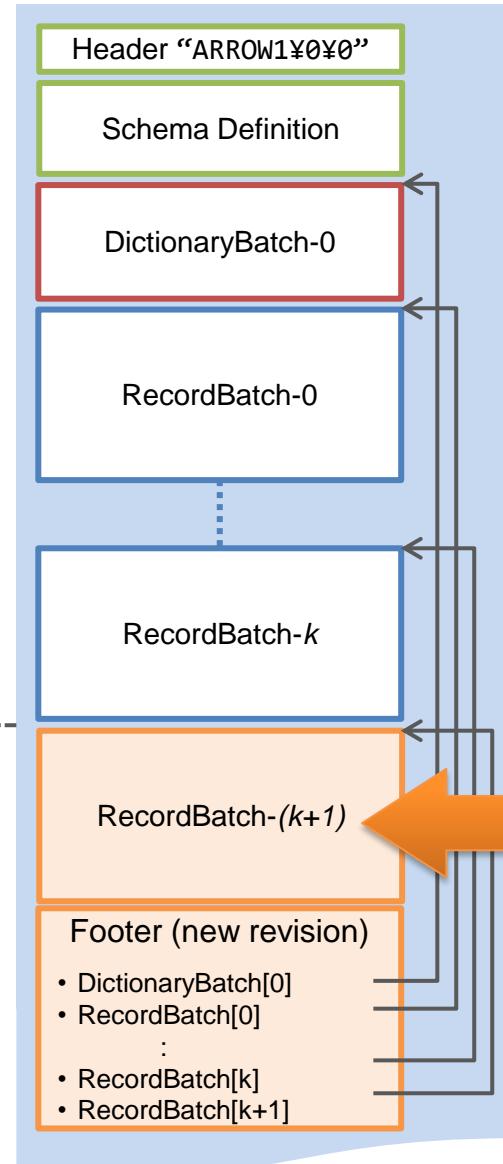
- Metadata - offset/length of RecordBatches and DictionaryBatches

Apache Arrow Internal (2/3) - writable Arrow_Fdw

Arrow file (before writes)

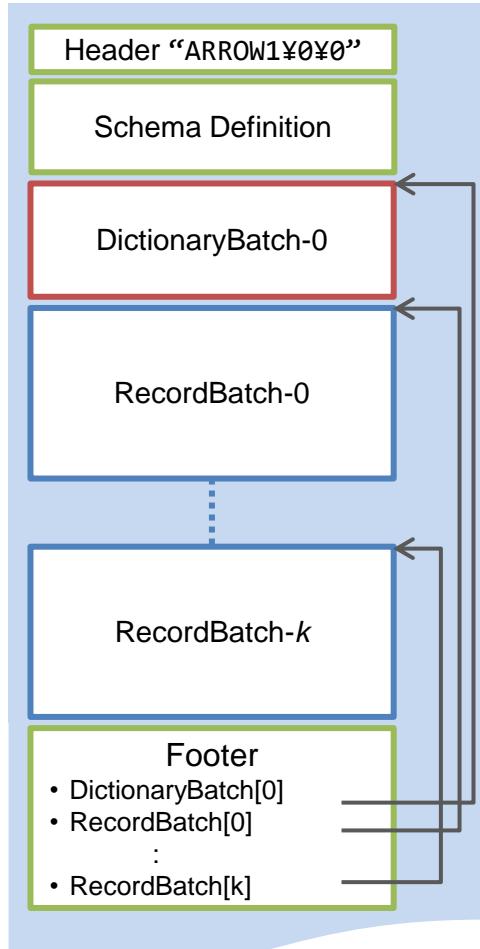


Arrow file (after writes)

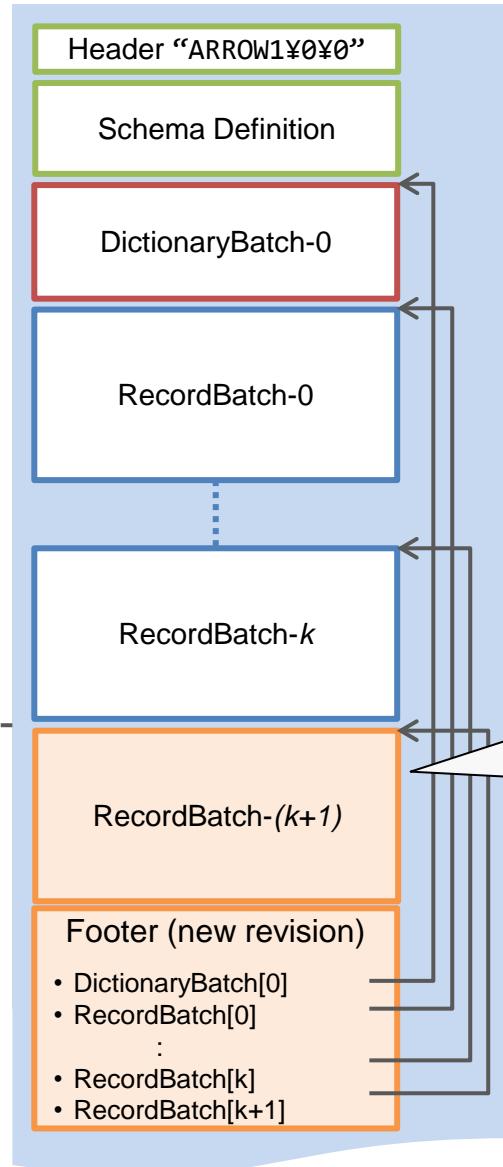


Apache Arrow Internal (3/3) - writable Arrow_Fdw

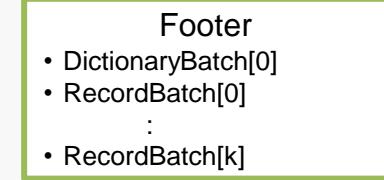
Arrow file (before writes)



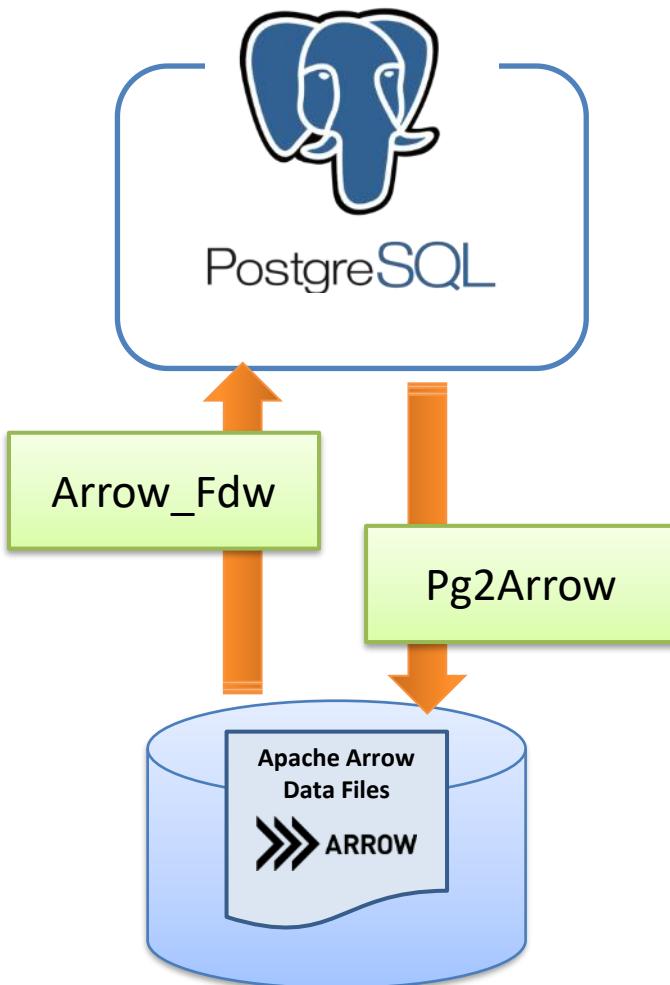
Arrow file (after writes)



Arrow_Fdw keeps the original Footer image until commit, for support of rollback.



Pg2Arrow / MySQL2Arrow enables to dump SQL results as Arrow file



```
$ pg2arrow --help
Usage:
  pg2arrow [OPTION] [database] [username]

General options:
  -d, --dbname=DBNAME    Database name to connect to
  -c, --command=COMMAND  SQL command to run
  -t, --table=TABLENAME  Table name to be dumped
    (-c and -t are exclusive, either of them must be given)
  -o, --output=FILENAME  result file in Apache Arrow format
    --append=FILENAME     result Apache Arrow file to be appended
    (--output and --append are exclusive. If neither of them
     are given, it creates a temporary file.)
```

Arrow format options:

```
-s, --segment-size=SIZE size of record batch for each
```

Connection options:

```
-h, --host=HOSTNAME    database server host
-p, --port=PORT         database server port
-u, --user=USERNAME    database user name
-w, --no-password      never prompt for password
-W, --password          force password prompt
```

Other options:

```
--dump=FILENAME        dump information of arrow file
--progress              shows progress of the job
--set=NAME:VALUE         config option to set before SQL execution
--help                  shows this message
```

✓ mysql2arrow also has almost equivalent capability