



# Parsing Postgres logs the non-pgBadger way



Kaarel Moppel, Freelance PostgreSQL Consultant  
pgConf.EU 2025, Riga

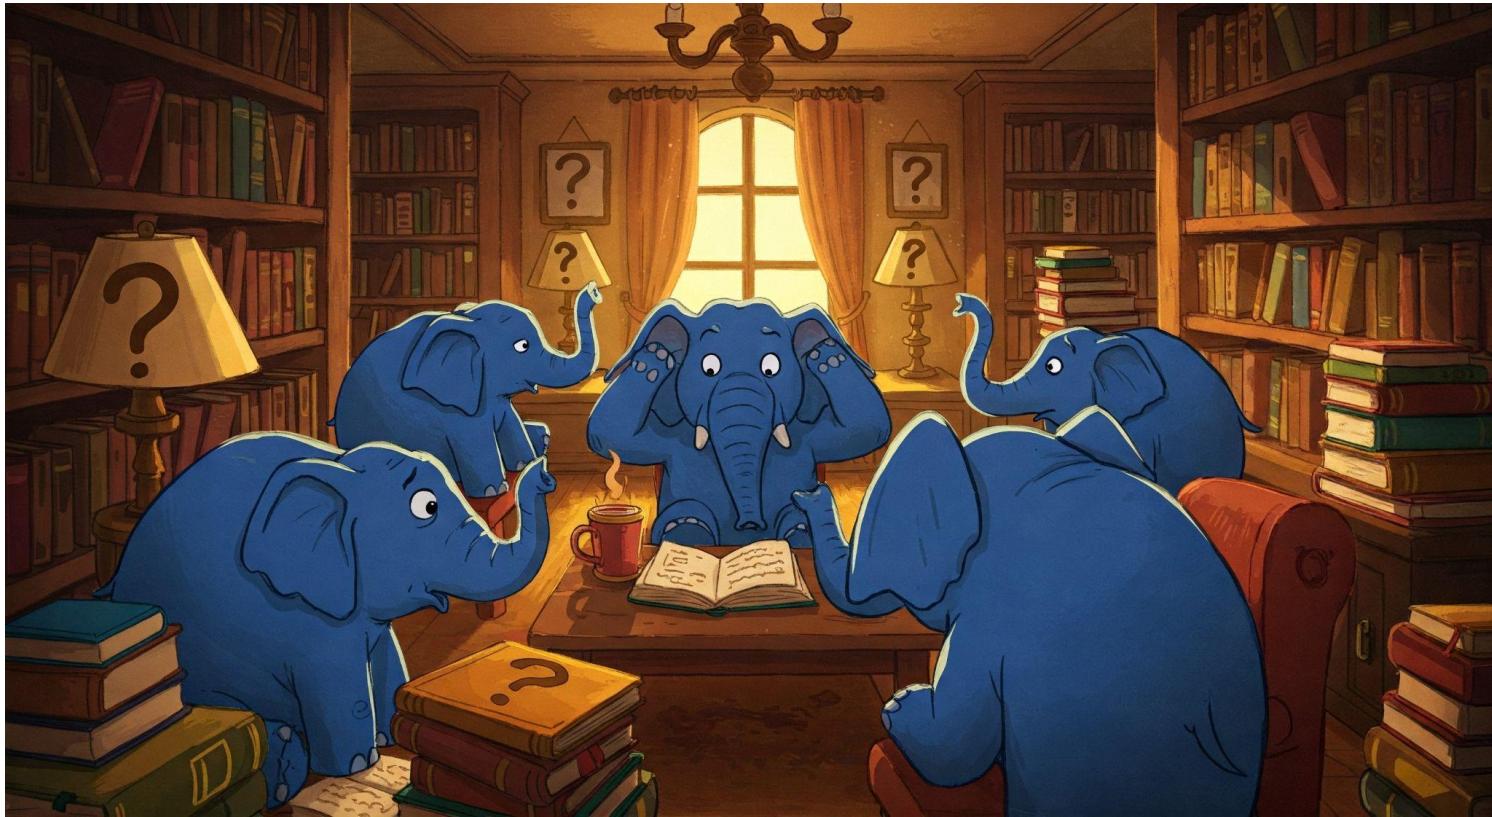
# \$ whoami

- Full-time “wrestling” with databases since 2007
- 20K+ hours in the Postgres ecosystem
  - Many hats along the way
  - Have developed somekind of a gut feeling on “Postgres”-y things if anything
- Freelance DBE / DBA / consultant
  - <https://kmoppel.github.io/> (Postgres Blog & Services)

# Agenda

- Value of Postgres server logs
- Log management approaches
- Common tools
- Cloud / SaaS tools
- AI-assisted tooling
- A fresh take - pgweasel

# Value of logs



# Why look at Postgres logs ?

The Postgres Stats subsystem doesn't cover everything by far!

- The obvious - get full details on PostgreSQL errors
- Slow queries with specific execution params
- Query plans (auto\_explain)
- Detect unauthenticated access / brute force attack attempts
  - Sadly possible by default with Postgres ...
- Background processing details (e.g. autovacuum, 3rd party exts.)
- Lock wait details
- Temp file details
- Backend terminations / OOM
- ...

# Log management approaches



# Log management approaches

- Ad-hoc searching/grepping
  - grep, pgBadger & co
  - Exposing logfiles directly via file\_fdw for SQL power
- Continuous analyzing
  - Loading into a DB (e.g. Postgres + ZFS)
  - Homegrown parsing (e.g. Cron+Python)
  - Filebeat / ELK (Elastic) pipeline
    - Sometimes via Syslog
  - Prometheus integration, e.g. [pgwatch](#) "server\_log\_event\_counts"
  - Generic framework (Nagios etc) Postgres plugins
  - Cloud tools e.g. Loggly or full APM suites
- Hooking into the Postgres log emitting mechanism
  - Tapping into "emit\_log\_hook", a la "[redislog](#)" (unmaintained sadly)
- ...

# Common tools



# Common tools - grep & co

The daily driver for many uses cases...

```
grep -E 'ERROR|FATAL' testdata/cloudsql.log  
egrep 'ERROR|FATAL' testdata/cloudsql.log  
egrep -c 'ERROR|FATAL' testdata/cloudsql.log
```

```
# Alternatives like "ripgrep" can be faster / nicer ...  
rg 'ERROR|FATAL' testdata/cloudsql.log
```

...

# Common tools - awk

For ad-hoc analysis it's "possible" to achieve below tasks (and DBA-s usually have something similar at their fingertips)...but not easy. LLMs can help a bit of course :)

- Find the N slowest queries
- Errors by type
- Connection attempts per user or host
- Average query duration
- Event counts by hour / minute
- ...

PS with CSV format life is somewhat nicer:

```
awk -F, '$12 == "ERROR" && $3 ~ /mydb/' postgresql.csv
```

# Common tools - awk - an example (not perfect)

```
gawk '{
    # time_parts[1] will contain "HH:MM:SS", time_parts[2] will contain the milliseconds
    split($2, time_parts, ":");

    timestamp_str = $1 " " time_parts[1];

    # Replace dashes and colons with spaces for mktime() format "YYYY MM DD HH MM SS".
    gsub(/[-:]/, " ", timestamp_str);

    # Convert the string to an epoch timestamp.
    epoch = mktime(timestamp_str);

    # Format the epoch time to an hourly key (e.g., "2025-05-21 10").
    hour_key = strftime("%Y-%m-%d %H", epoch);

    hourly_counts[hour_key]++;
}
END {
    print "--- Hourly Event Summary ---";
    for (hour in hourly_counts) {
        print hour ":00:00 | Events: " hourly_counts[hour];
    }
}' testdata/cloudsql.log
```

# Common tools - [pgBadger](#)

The #1 choice for many - a solid tool in principle! Main highlights:

- A HTML report with graphs by default
- Single-run and incremental modes
- Parallel processing (--jobs) support (not for CSV)
- Integrated remote access (SSH, http[s], [s]ftp)
- Ecosystem aware - RDS Cloudwatch & pgBouncer
- --exclude-query / --include-query regex object filtering
- Bind parameters “merging” for prepared statements
- JSON output
- --dump-raw-csv for loading / more structure
- ...

```
krl@e7420:~$ pgbadger --help | grep -E '^s+\-{1,2}' | wc -l  
106
```

# Common techniques - loading into Postgres

To gain the power of SQL! 💪 Works OOTB with CSV only...

```
CREATE FOREIGN TABLE pglog (
    log_time timestamp(3) with time zone,
    user_name text,
    database_name text,
    process_id integer,
    connection_from text,
    session_id text,
    session_line_num bigint,
    command_tag text,
    session_start_time timestamp with time zone,
    virtual_transaction_id text,
    transaction_id bigint,
    error_severity text,
    sql_state_code text,
    message text,
    detail text,
    hint text,
    internal_query text,
    internal_query_pos integer,
    context text,
    query text,
    query_pos integer,
    location text,
    application_name text,
    backend_type text,
    leader_pid integer,
    query_id bigint
) SERVER pglog
OPTIONS ( filename 'log/pglog.csv', format 'csv' );
```

The fields / format can be best viewed from “file\_fdw” documentation

<https://www.postgresql.org/docs/current/file-fdw.html>

PS - a reminder on pgbadger  
--dump-raw-csv mode again for default textual logs

# Cloud / SaaS tools



# Cloud / SaaS tools - no shortage

The “problem space” is not new actually - quite some tools out there!

- Top managed DBaaS providers have something built-in
  - Usually with some proprietary query language though :(
- Datadog
- AppDynamics
- Pganalyze
- Loggly
- BetterStack
- ...

**PS** Note that in many cases you need to adhere to their understanding of `log\_line\_prefix`!

# Tooling - AD 2025



# AI-assisted tooling ?

The big questions:

1. Can public LLM-s find the needle from a haystack?
2. Is it cost-effective?
3. Security ...

Based on my experiments:

- No point to throw everything at LLMs - need to pre-filter
  - Response time
  - \$\$ can get prohibitive - 1GB log file ~ \$1 with a fast model
  - Context window limitations / chunking
- Thus need to get hands dirty anyways - meaning direct usage loses in appeal...

# AI-assisted tooling ?

Demo - a simple 400kB / 4k line file with 10 problems:

<https://github.com/kmoppel/pgweasel/blob/main/testdata/cloudsql.log>

Result: Misses 2 simple “SELECT asd” ERROR’s + 3 FATAL’s with rarer protocol errors

Model: Gemini 2.5 Flash

The screenshot shows a dark-themed application window titled "cloudsql" with a "LOG" tab selected. A search bar at the top right contains the placeholder text "find all errors from attached postgresql log file". Below the search bar, a section titled "Analysis" is expanded, showing the following content:

The PostgreSQL log file, `cloudsql.log`, contains the following **ERROR** and **FATAL** entries:

**Errors (ERROR Level)**

- **Timestamp:** 2025-05-21 10:35:01.627 UTC
  - **Error:** `ERROR: relation "public.heartbeat" does not exist at character 13`
  - **User/DB:** `cloudsqladmin` connecting to `cloudsqladmin` from `127.0.0.1`
  - **Statement:** An `INSERT INTO public.heartbeat` statement that attempts to use `ON CONFLICT (id) DO UPDATE` failed because the `public.heartbeat` table was not found

**Fatal Errors (FATAL Level)**

- **Timestamp:** 2025-05-21 10:38:32.056 UTC
  - **Error:** `FATAL: password authentication failed for user "postgres"`
  - **User/DB:** `postgres` connecting to `postgres` from `91.129.106.131`

# AI-assisted tooling ?

But better already with  
autovacuum analysis!

As there are a gazillion of  
evolving models out there,  
maybe one day good enough



what was the longest autovacuum or autoanalyze dura

Analysis ▾

The longest duration for an **autovacuum** or **autoanalyze** operation in the log file was **2.15 s**.

This duration was for the **automatic analyze of table "postgres.public.pgbench\_accounts"**



Operation	Table	Elapsed Duration
Automatic vacuum	cloudsqladmin.public.heartbeat	0.50 s ⚡
Automatic analyze	cloudsqladmin.public.heartbeat	0.00 s ⚡
Automatic vacuum	postgres.public.pgbench_accounts	0.00 s ⚡
<b>Automatic analyze</b>	<b>postgres.public.pgbench_accounts</b>	<b>2.15 s</b> ⚡
Automatic analyze	postgres.public.pgbench_branches	0.00 s ⚡
Automatic analyze	postgres.public.pgbench_tellers	0.00 s ⚡
Automatic vacuum	cloudsqladmin.public.heartbeat	0.00 s ⚡
Automatic analyze	cloudsqladmin.public.heartbeat	0.00 s ⚡

Export to Sheets

# AI-assisted tooling ?

Some “specialized” progress is being made of course, e.g.:

<https://github.com/salesforce/logai>

...but it's harder than you think!



# Tooling - summary

I've personally tried out A LOT of tools...but mostly concluded:

- Not that they create as many problems as they solve, but every tool brings some limitations as usual
  - Which tend to surface only when past some first simpler milestones
- Sadly common pretty to hit bugs or performance / parallelism walls
- The tools keep changing...creating random backlog

**Consider if you really need all those features!**

Starting with something super-simple, but reliable, and fully under your control, is never a bad idea!

# pgweasel



# pgweasel - to complement pgBadger

A new project, aiming to gain community attention / action!

The tool tries to be different from pgBadger, focusing on the "Pareto" DBA flow of **"on the server" CLI usage** for humans :)

- Much simple to handle - command + subcommand approach
  - Less features - "less is more" philosophy
- Way faster! Currently in Golang
  - Rust re-write started (in "rust-rewrite" branch)
- A single binary - Perl can get prohibitive on containers
- Zero-config, always works (no --log-line-prefix)
- User-friendly - relative time inputs, command aliases, auto-detect logs from standard locations, automatic multi-core usage
- ...

# pgweasel - looking for contributors / co-owner!

- One person has quite a narrow view on things...
- Plus the typical side-project time issue
- Also...to make it work really good under all conditions a lot of server log examples would be needed!



In case you have some large real-life logfiles without security risks, please send me privately some S3 etc link or just open an Github issue!

1GB+ would be especially nice! 🙏

<https://github.com/kmoppel/pgweasel>

**Licence:** PostgreSQL license

All kinds of feedback, feature ideas, pull requests very much appreciated!  
Or just a  :)



DEMO time!